



**EECS 151/251A
Spring 2020
Digital Design and Integrated
Circuits**

Instructor:
John Wawrzynek

Lecture 11: RISC-V

Project Introduction

- You will design and optimize a RISC-V processor
- Phase 1: Design and demonstrate a processor
- Phase 2:
 - ASIC Lab – implement cache memory and generate complete chip layout
 - FPGA Lab – Add video display and graphics accelerator

Today discuss how to design the processor



What is RISC-V?

- Fifth generation of RISC design from UC Berkeley
- A high-quality, license-free, royalty-free RISC ISA specification
- Experiencing rapid uptake in both industry and academia
- Supported by growing shared software ecosystem
- Appropriate for all levels of computing system, from micro-controllers to supercomputers
 - 32-bit, 64-bit, and 128-bit variants (we're using 32-bit in class, textbook uses 64-bit)
- Standard maintained by non-profit RISC-V Foundation

<https://riscv.org/specifications/>

Foundation Members (60+)



Platinum:



Berkeley
Architecture
Research

bluespec

Google

IBM.

Mellanox®
TECHNOLOGIES



Microsoft

Rambus

Cryptography Research



SiFive

cortus

D R A P E R

SAMSUNG

Microsemi

QUALCOMM®



NVIDIA.

WD Western
Digital®



Gold, Silver, Auditors:

AMD

ANDES
TECHNOLOGY

Codasip

Esperanto
Technologies

IDT

INTRINSIX

MP Processor

Sur Technology

antmicro



Blockstream

GRAY
RESEARCH



ROA
LOGIC

Technion

ESPRESSIF

ETH zürich

LATTICE
SEMICONDUCTOR

Rumble

VectorBlox
embedded supercomputing

ICT

中国科学院
INSTITUTE OF COMPUTING TECHNOLOGIES

MEDIATEK

lowRISC

runtime.io

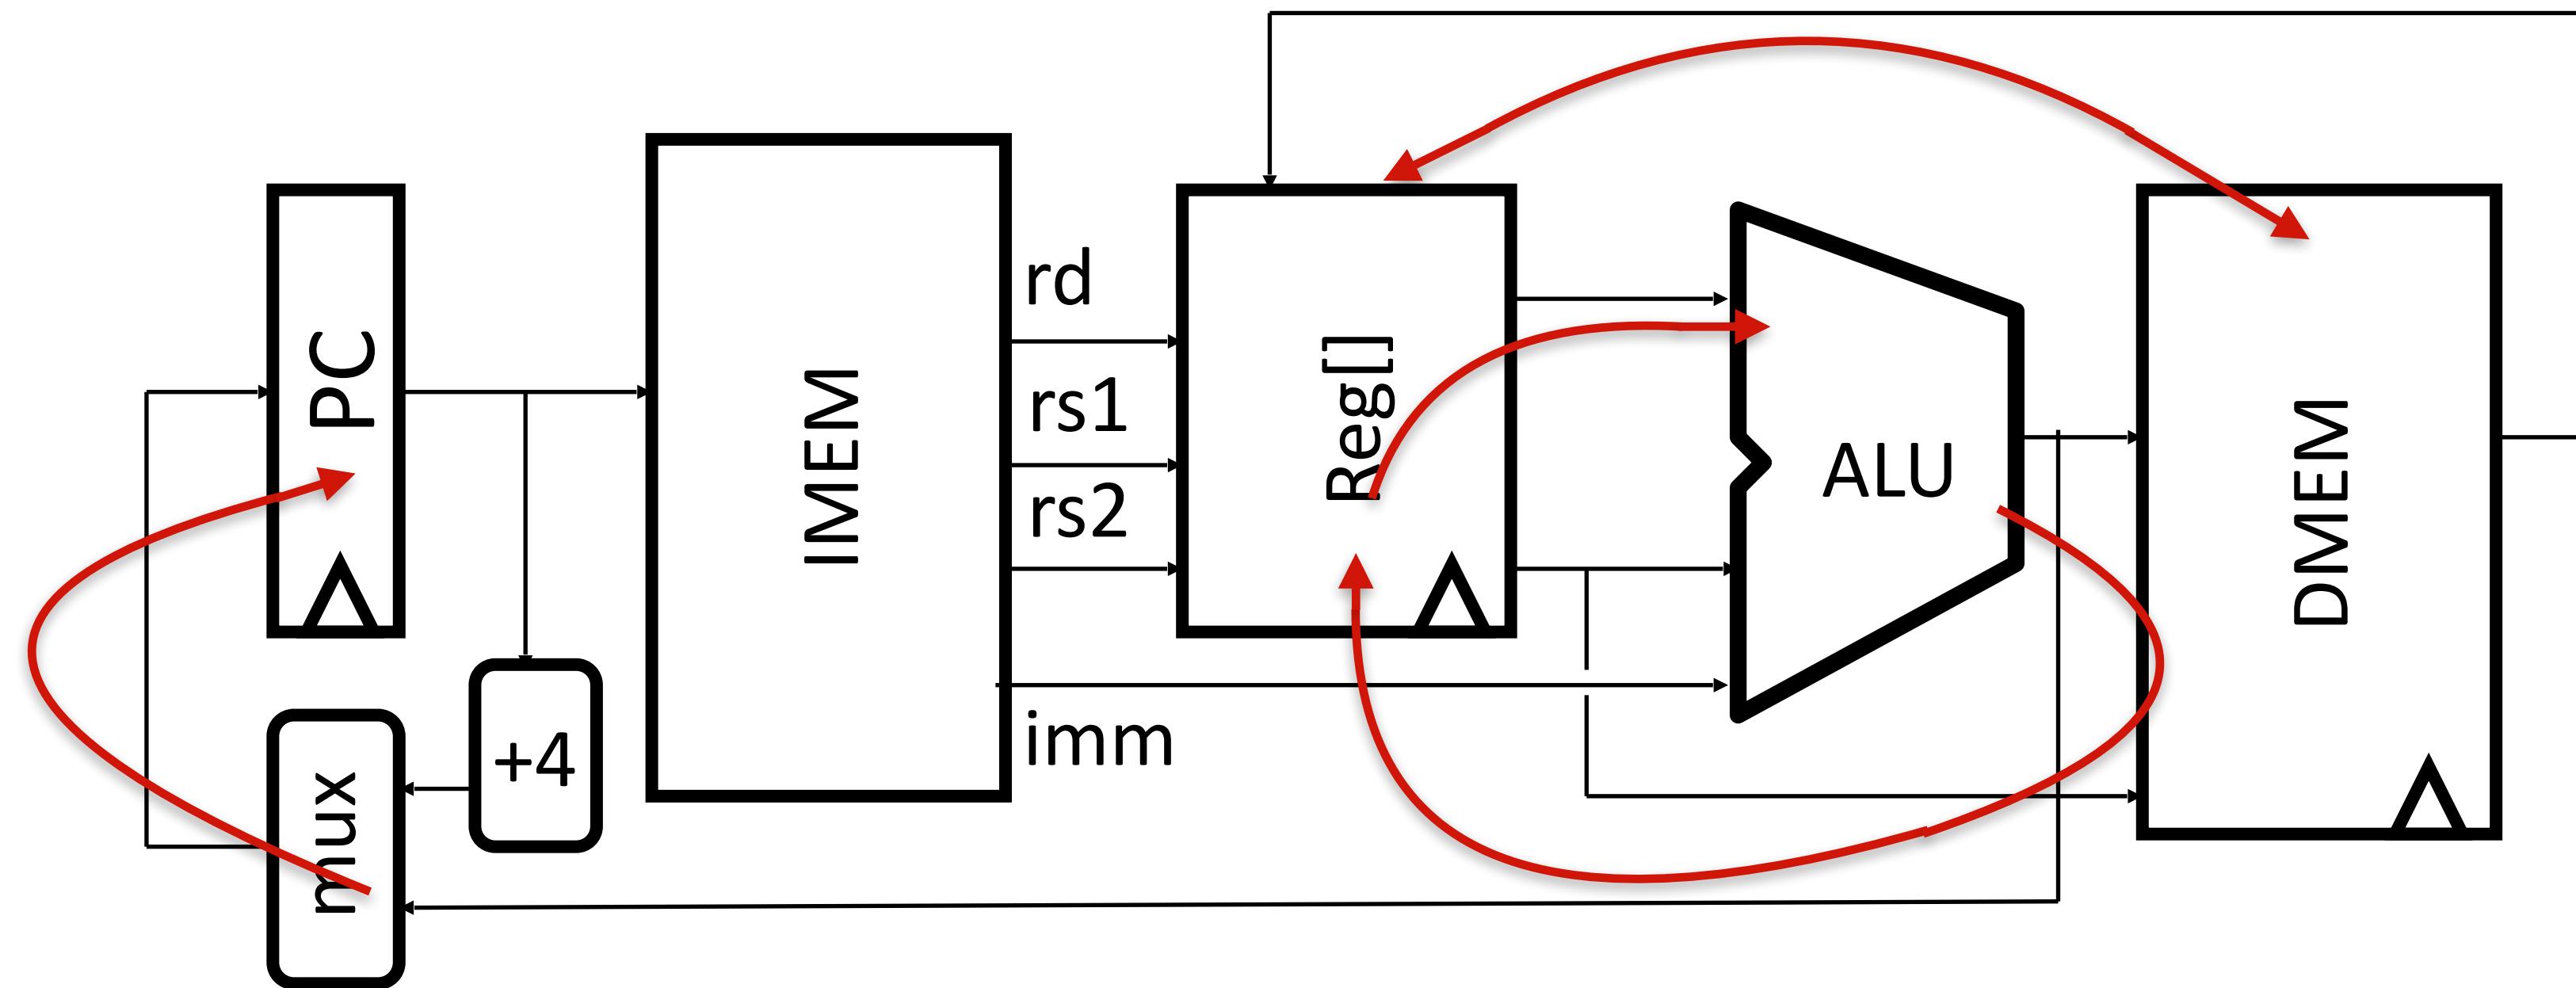
Instruction Set Architecture (ISA)

- Job of a CPU (*Central Processing Unit*, aka *Core*): execute *instructions*
- Instructions: CPU's primitives operations
 - Instructions performed one after another in sequence
 - Each instruction does a small amount of work (a tiny part of a larger program).
 - Each instruction has an *operation* applied to *operands*,
 - and might be used change the sequence of instruction.
- CPUs belong to “families,” each implementing its own set of instructions
- CPU’s particular set of instructions implements an *Instruction Set Architecture (ISA)*
 - Examples: ARM, Intel x86, MIPS, RISC-V, IBM/Motorola PowerPC (old Mac), Intel IA64, ...



If you need more info on processor organization.

RISC Processor Instructions in Brief



- *Compilers generate machine instructions to execute your programs in the following way:*
- Load/Store instructions move operands between main memory (cache hierarchy) and core register file.
- Register/Register instructions perform arithmetic and logical operations on register file values as operands and result returned to register file.
- Register/Immediate instructions perform arithmetic and logical operations on register file value and constants.
- Branch instructions are used for looping and if-than-else (data dependent operations).
- Jumps are used for function call and return.

Complete RV32I ISA

imm[31:12]			rd	0110111	LUI
imm[31:12]			rd	0010111	AUIPC
imm[20 10:1 11 19:12]			rd	1101111	JAL
imm[11:0]		rs1	000	rd	JALR
imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	BEQ
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	BNE
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	BLT
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	BGE
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	BLTU
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	BGEU
imm[11:0]		rs1	000	rd	LB
imm[11:0]		rs1	001	rd	LH
imm[11:0]		rs1	010	rd	LW
imm[11:0]		rs1	100	rd	LBU
imm[11:0]		rs1	101	rd	LHU
imm[11:5]	rs2	rs1	000	imm[4:0]	SB
imm[11:5]	rs2	rs1	001	imm[4:0]	SH
imm[11:5]	rs2	rs1	010	imm[4:0]	SW
imm[11:0]		rs1	000	rd	ADDI
imm[11:0]		rs1	010	rd	SLTI
imm[11:0]		rs1	011	rd	SLTIU
imm[11:0]		rs1	100	rd	XORI
imm[11:0]		rs1	110	rd	ORI
imm[11:0]		rs1	111	rd	ANDI
00000000		00010011			

0000000	shamt	rs1	001	rd	0010011	SLLI
0000000	shamt	rs1	101	rd	0010011	SRLI
0100000	shamt	rs1	101	rd	0010011	SRAI
0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND
0000	pred	succ	00000	000	00000	0001111
0000	0000	0000	00000	001	00000	0001111
000000000000			00000	000	00000	1110011
000000000001			00000	000	00000	1110011
csr	rs1	001	rd			1110011
csr	rs1	101	rd			1110011
csr	rs1	011	rd			1110011
csr	zimm	101	rd			1110011
csr	zimm	110	rd			1110011
csr	zimm	111	rd			1110011

Not in EECS151/251A

SLLI
SRLI
SRAI
ADD
SUB
SLL
SLT
SLTU
XOR
SRL
SRA
OR
AND
FENCE
FENCE.I
ECALL
EBREAK *CSRRW
CSRRS
CSRRC
CSRRWI
CSRRI
CSRCI

* implemented in the ASIC project

Summary of RISC-V Instruction Formats

Binary encoding of machine instructions. Note the common fields.

31	30	25 24	21	20	19	15 14	12 11	8	7	6	0	
												R-type
												I-type
												S-type
												B-type
												U-type
												J-type

“State” Required by RV32I ISA

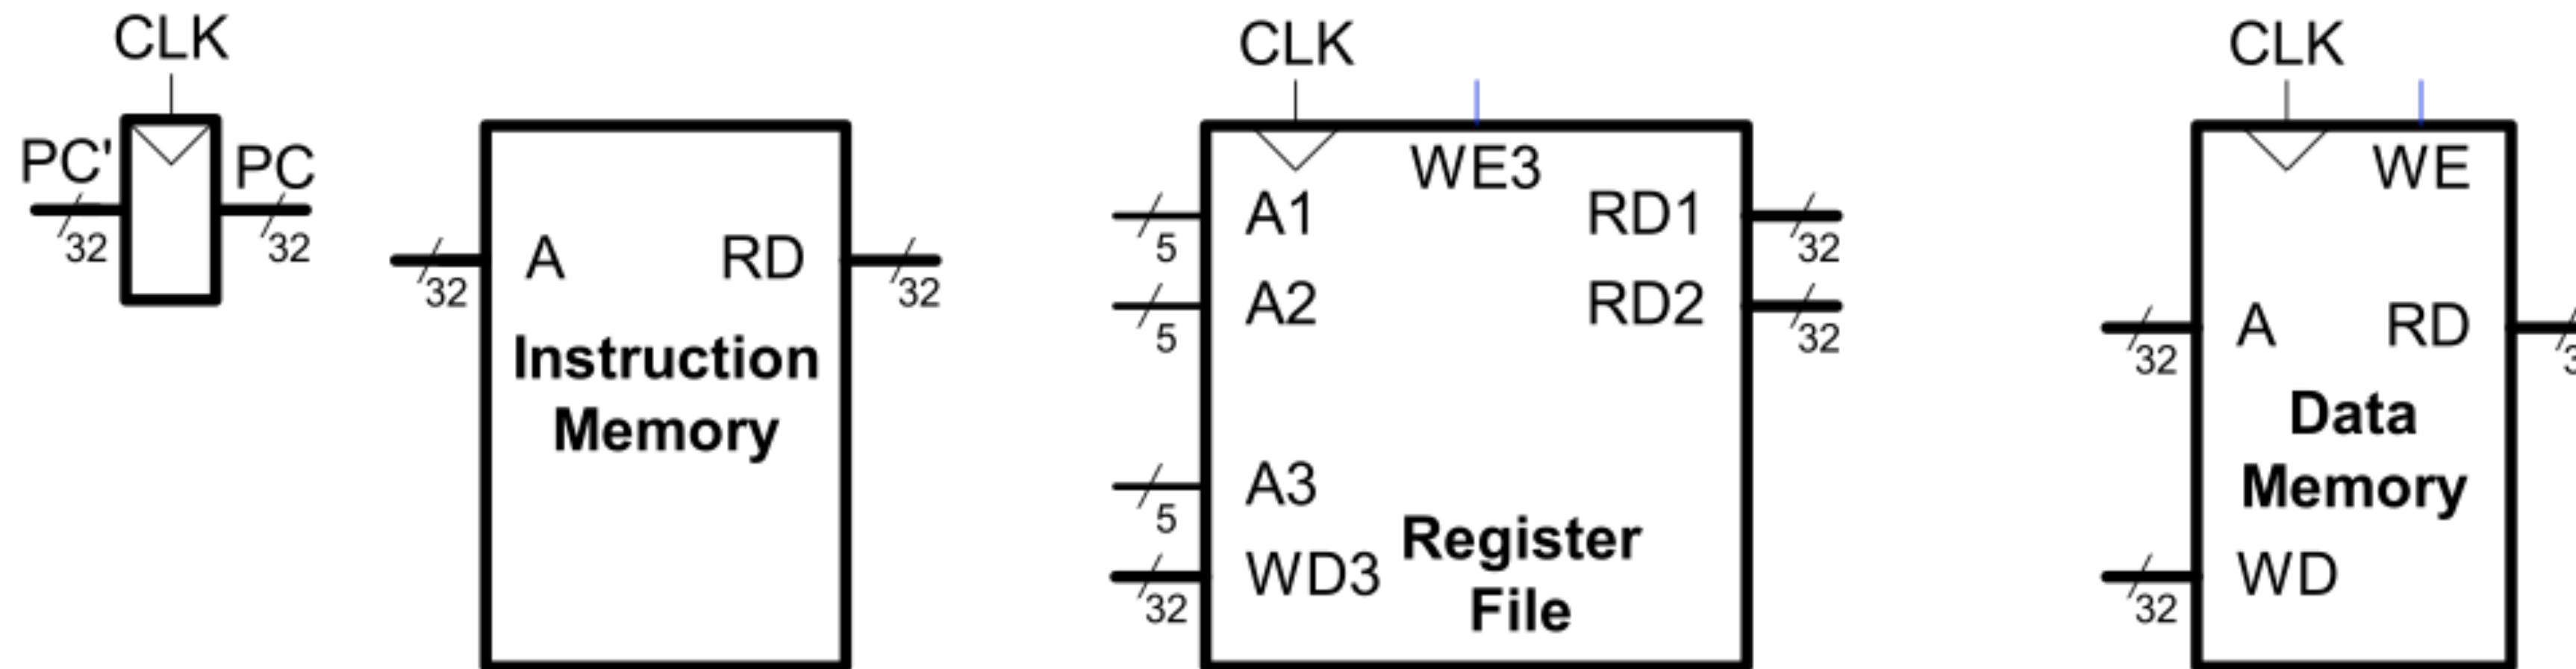
Each instruction reads and updates this state during execution:

- Registers (**x**0 .. **x**31)
 - Register file (or *regfile*) **Reg** holds 32 registers x 32 bits/register:
Reg[0] .. **Reg**[31]
 - First register read specified by *rs1* field in instruction
 - Second register read specified by *rs2* field in instruction
 - Write register (destination) specified by *rd* field in instruction
 - **x**0 is always 0 (writes to **Reg**[0] are ignored)
- Program Counter (PC)
 - Holds address of current instruction
- Memory (MEM)
 - Holds both instructions & data, in one 32-bit byte-addressed memory space
 - We’ll use separate memories for instructions (**IMEM**) and data (**DMEM**)
 - *Later we’ll replace these with instruction and data caches*
 - Instructions are read (*fetched*) from instruction memory (assume **IMEM** read-only)
 - Load/store instructions access data memory

RISC-V State Elements

- State encodes everything about the execution status of a processor:

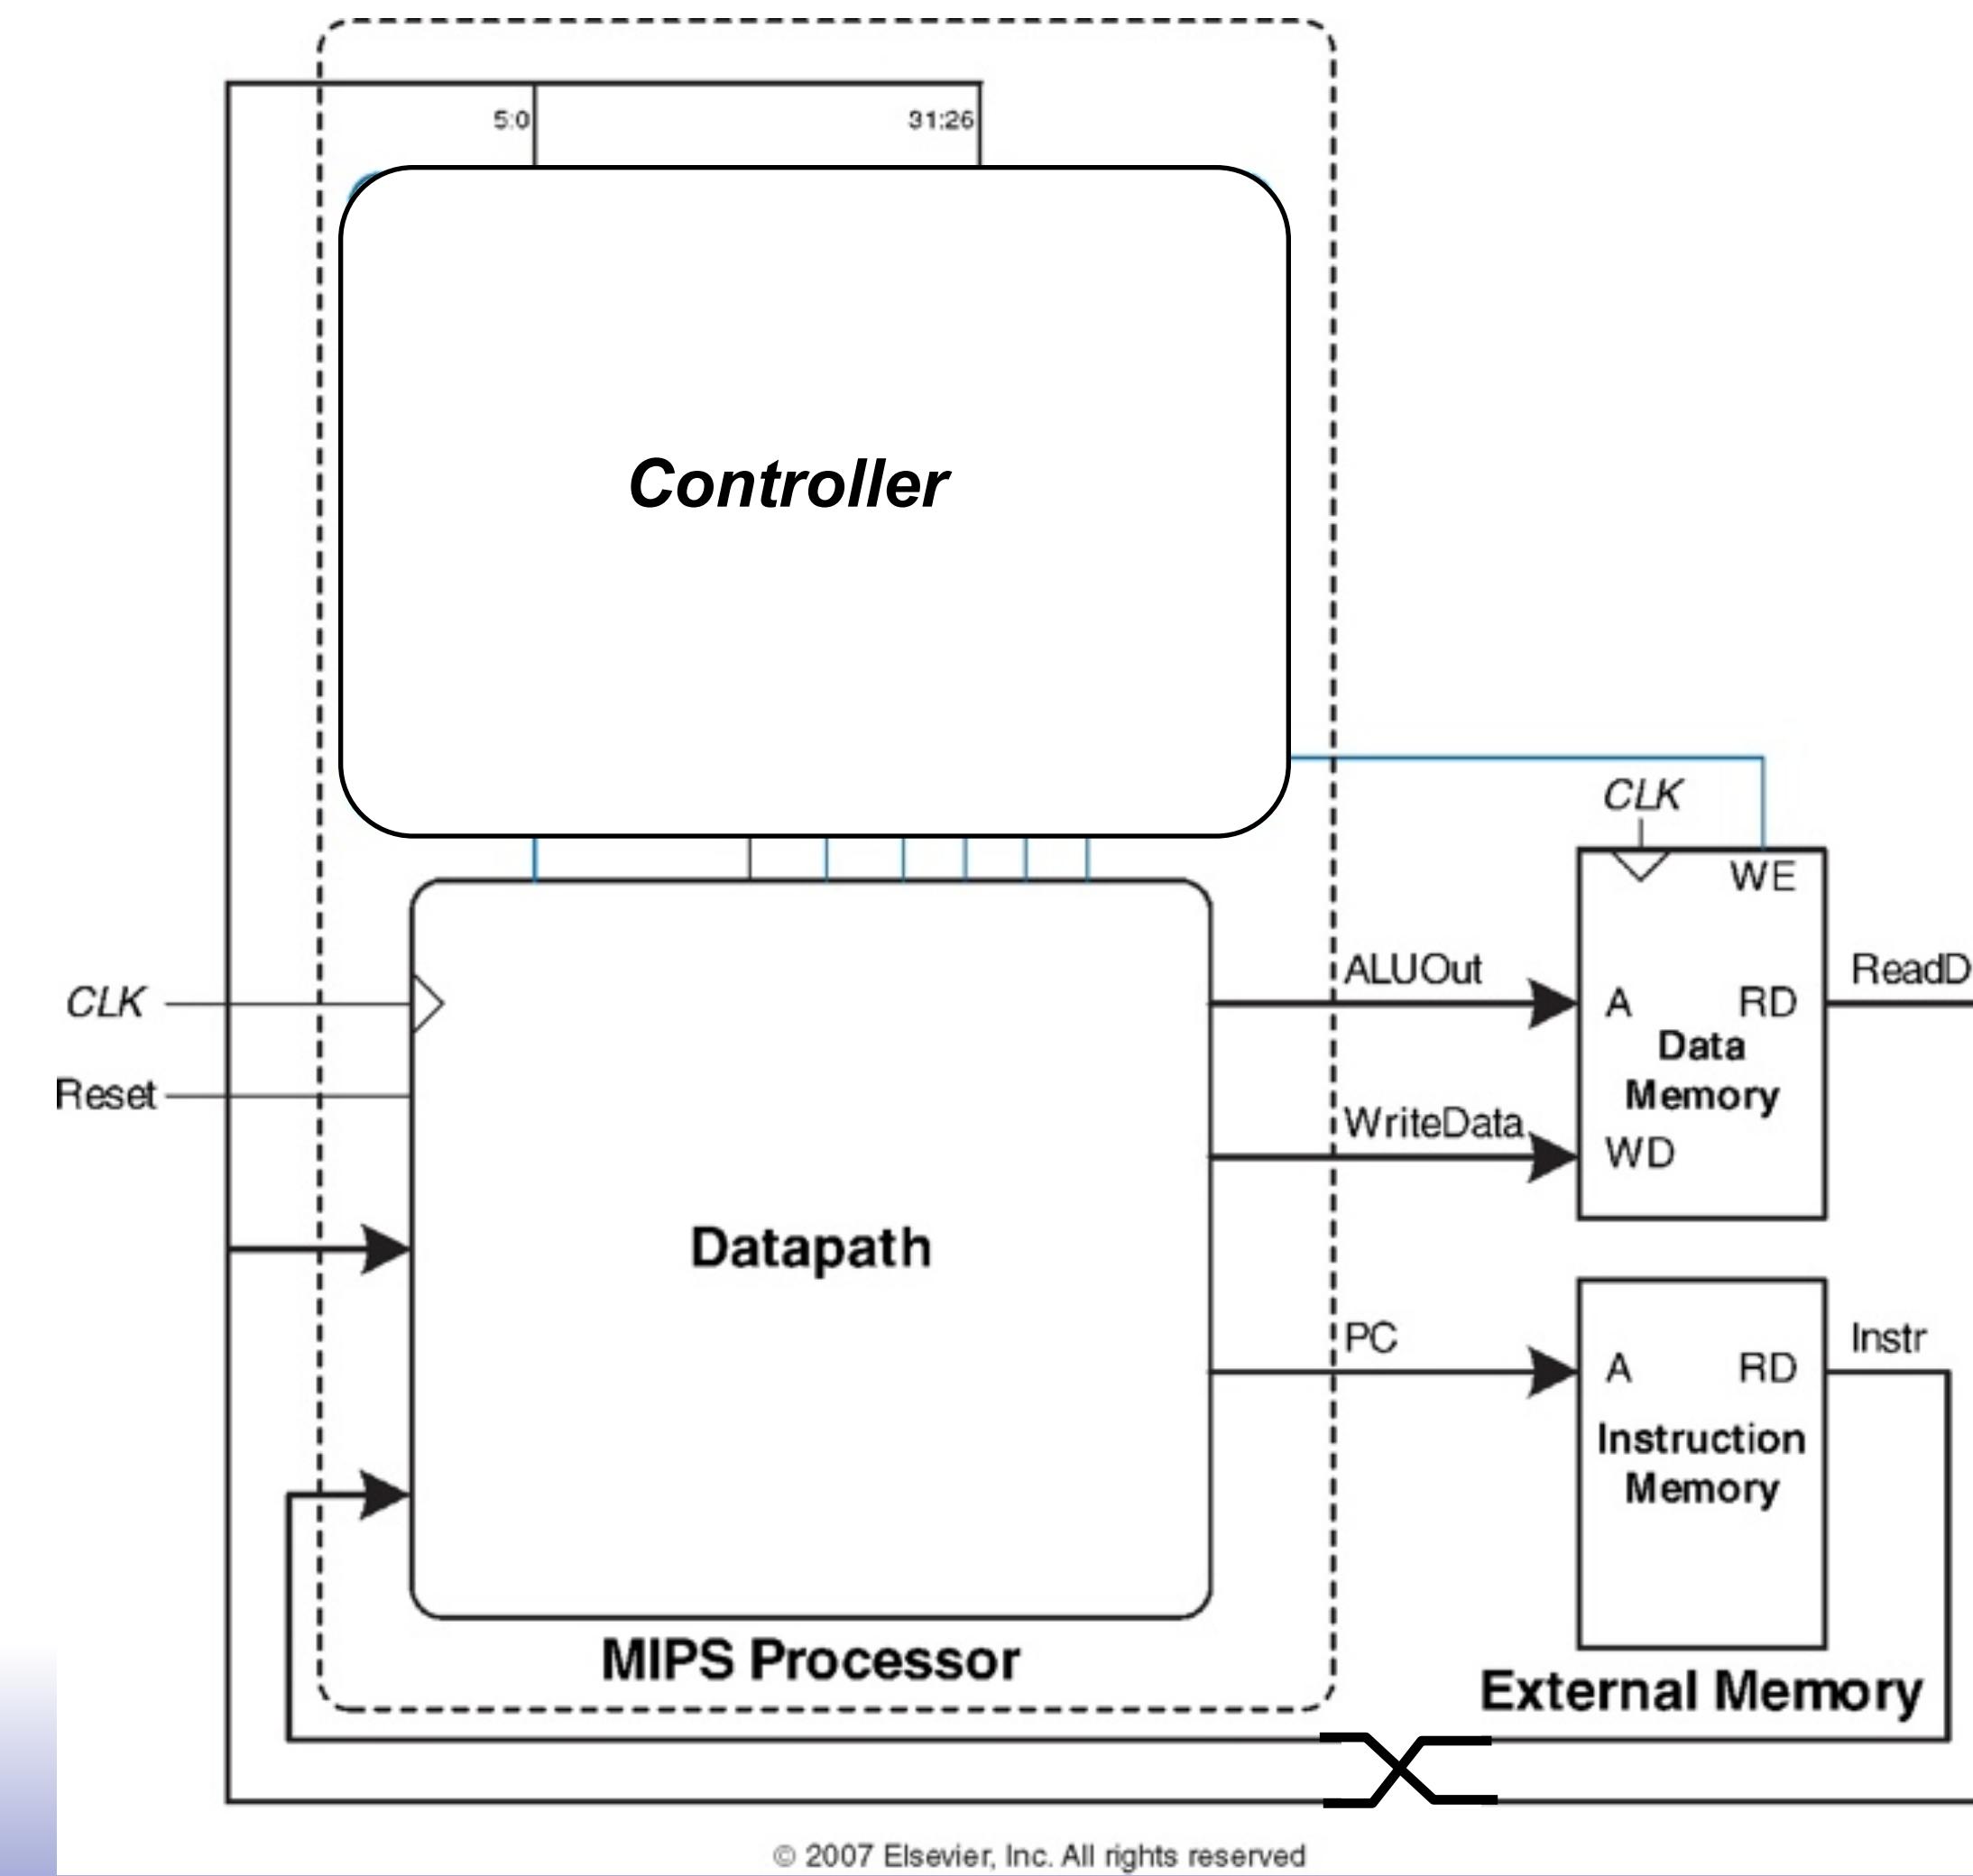
- PC register
- 32 registers
- Memory



Note: for these state elements, clock is used for write but not for read (asynchronous read, synchronous write).

RISC-V Microarchitecture Organization

Datapath + Controller + External Memory



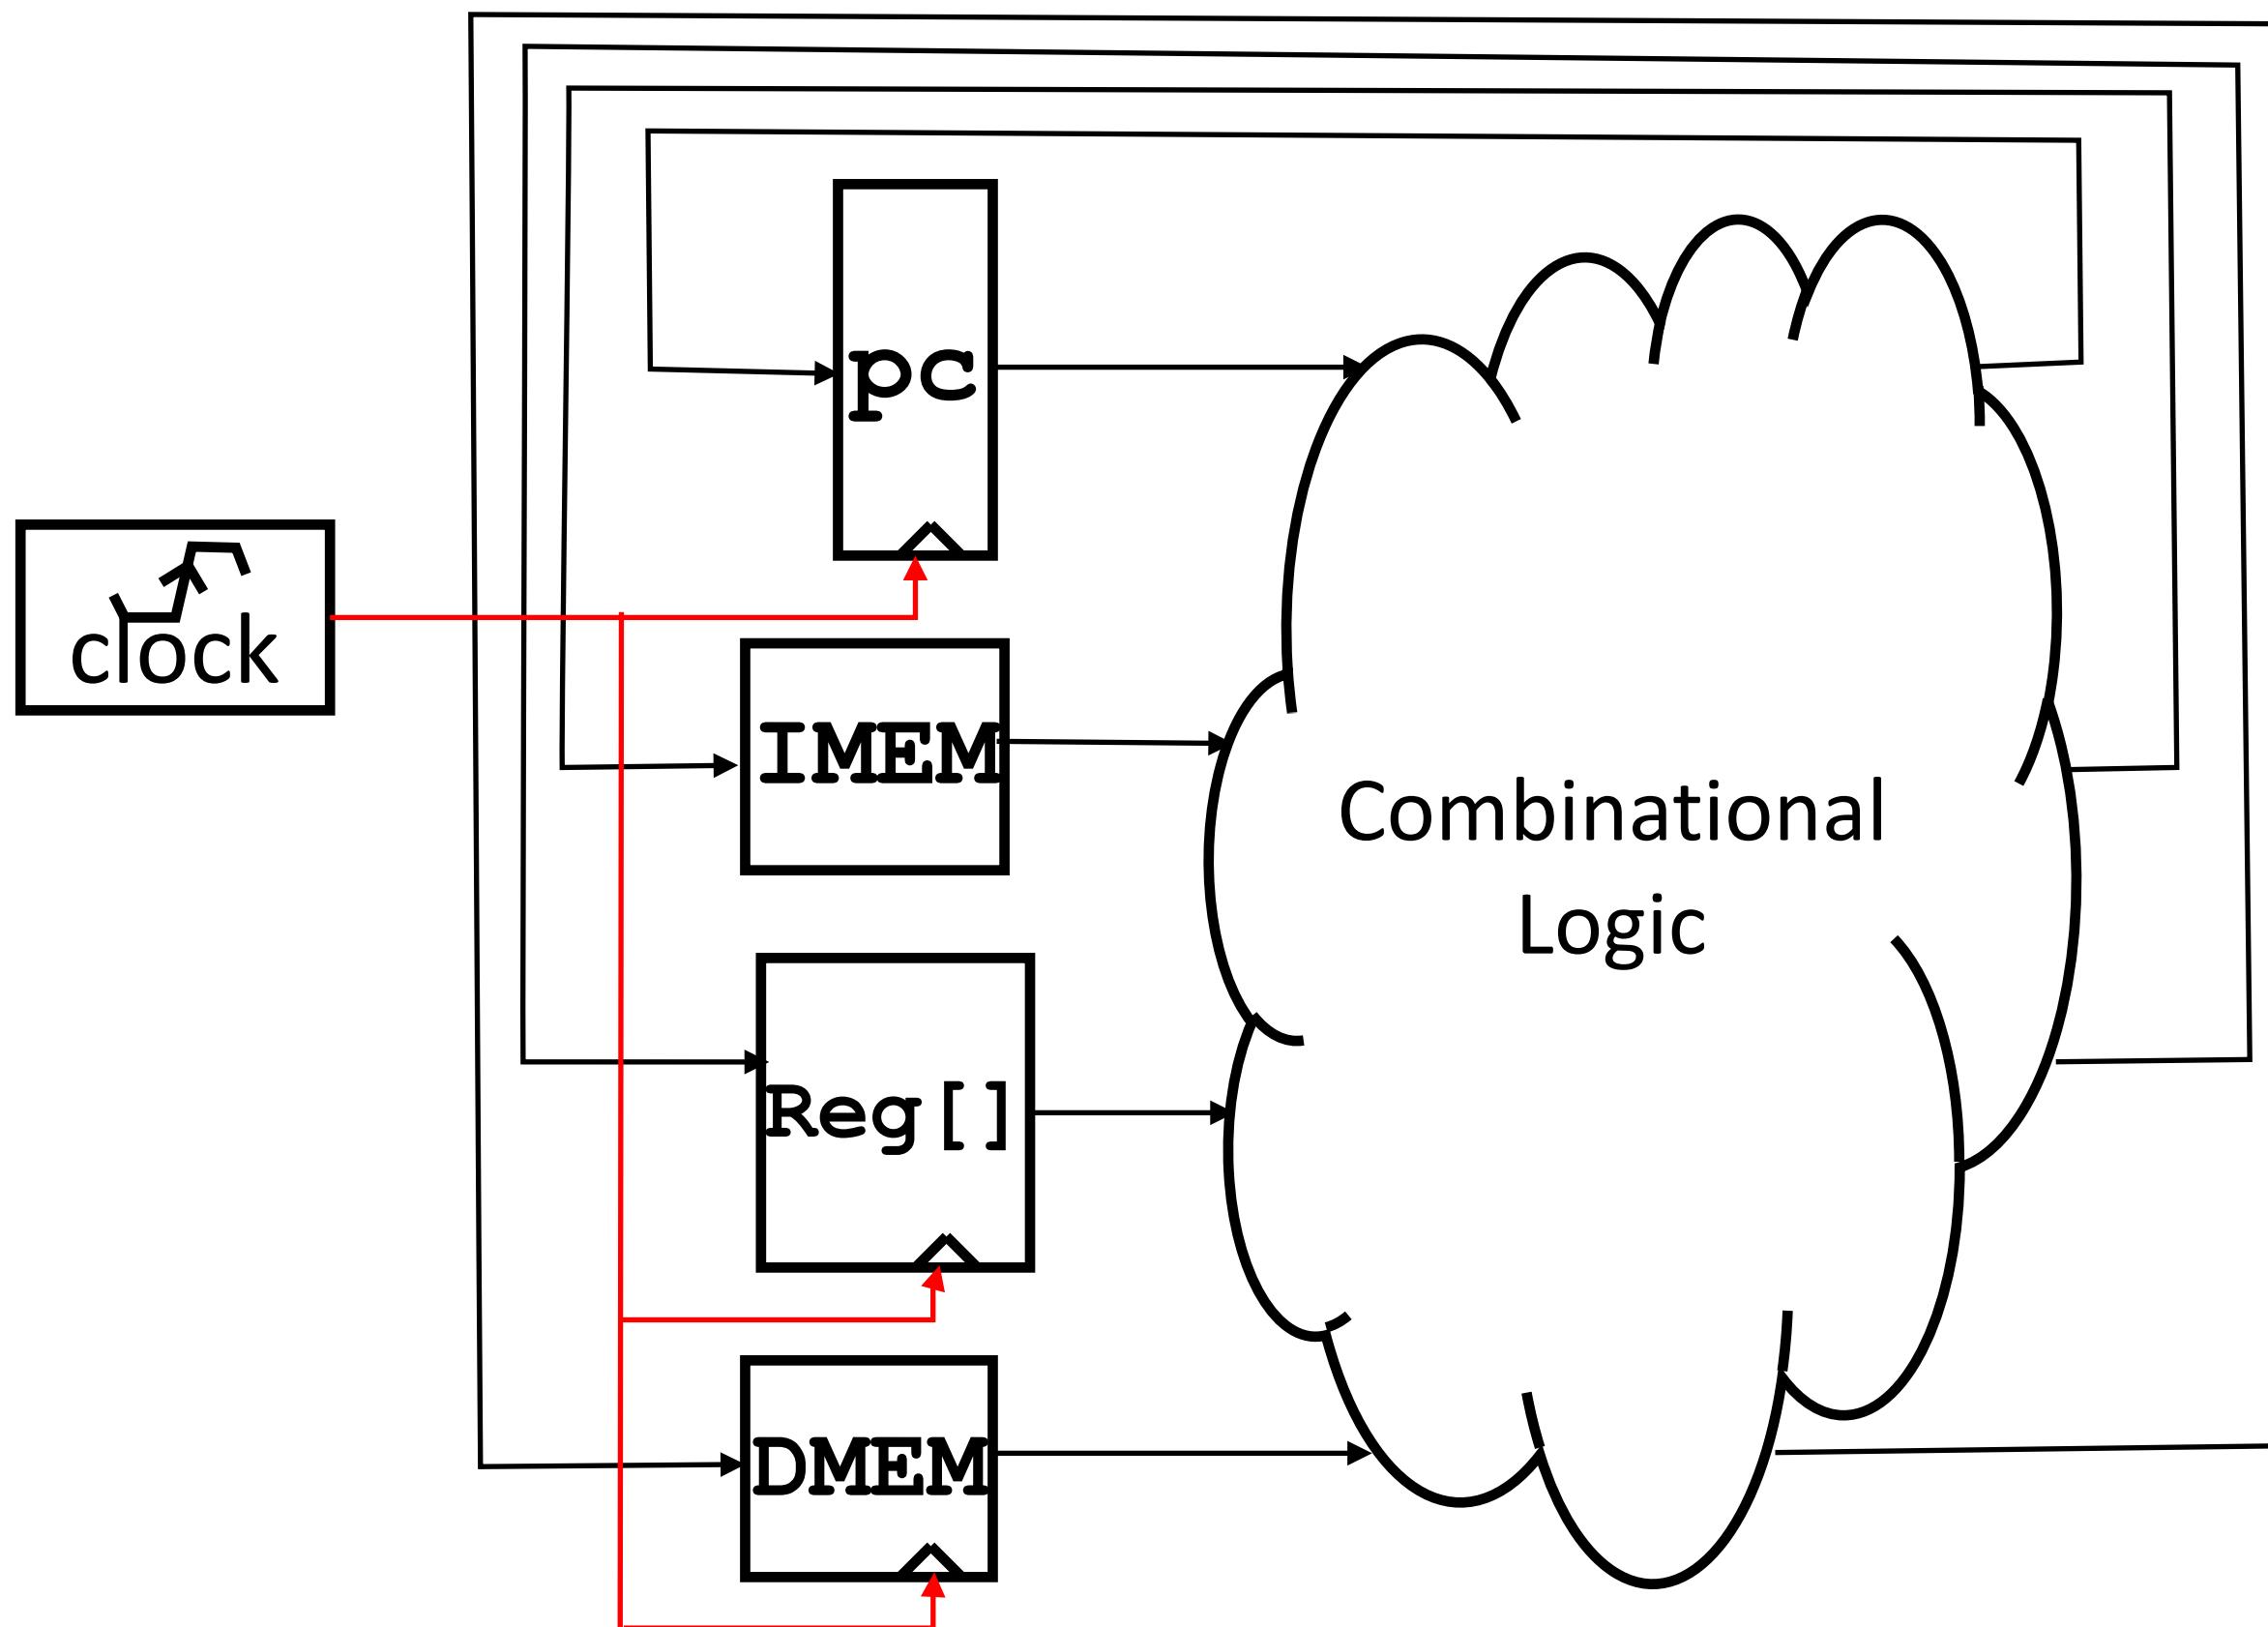
Microarchitecture

Multiple implementations for a single instruction set architecture:

- *Single-cycle*
 - Each instruction executes in a single clock cycle.
- *Multicycle*
 - Each instruction is broken up into a series of shorter steps with one step per clock cycle.
- *Pipelined (variant on “multicycle”)*
 - Each instruction is broken up into a series of steps with one step per clock cycle
 - Multiple instructions execute at once by overlapping in time.
- *Superscalar*
 - Multiple functional units to execute multiple instructions at the same time
- *Out of order...*
 - Instructions are reordered by the hardware

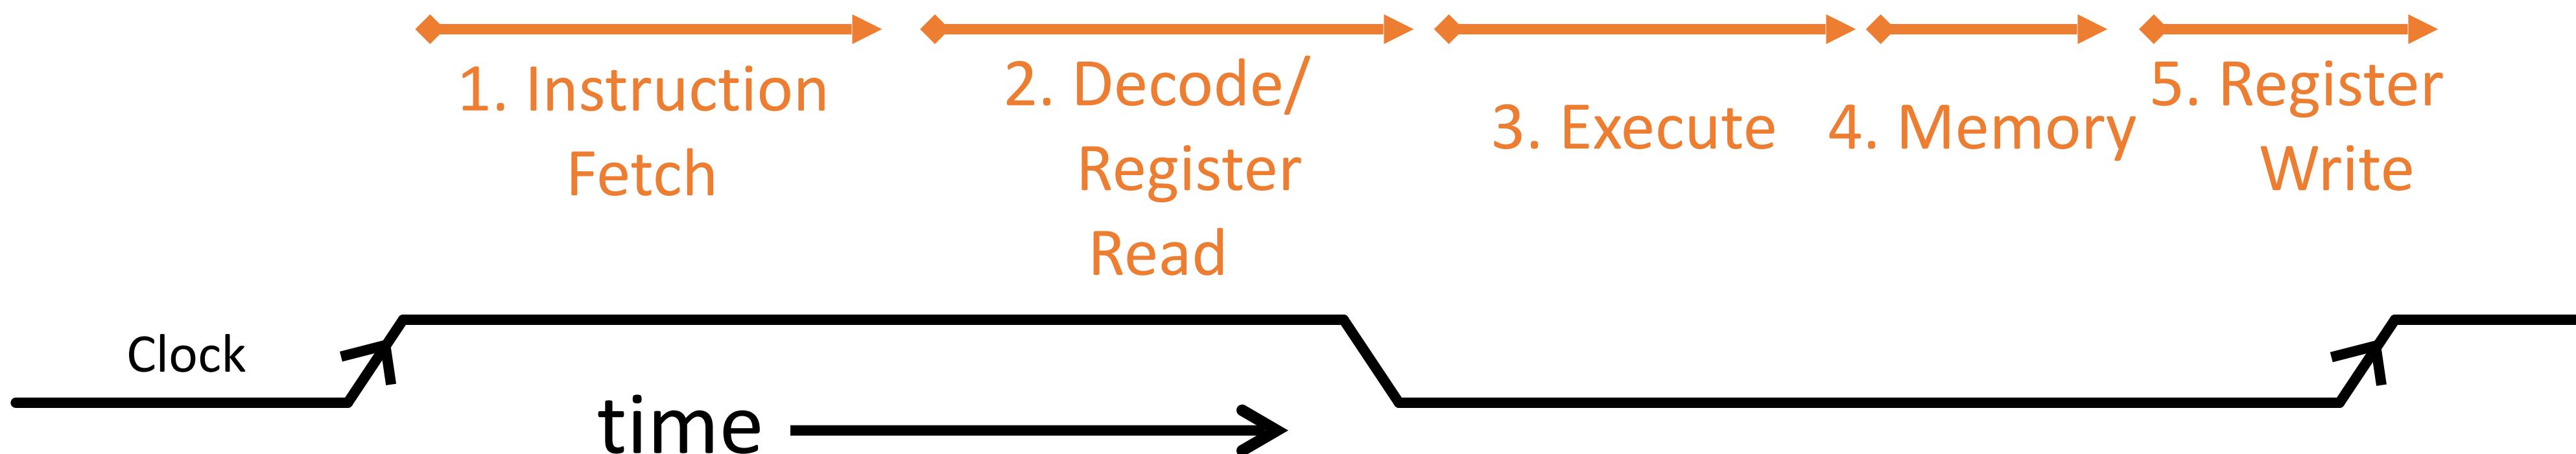
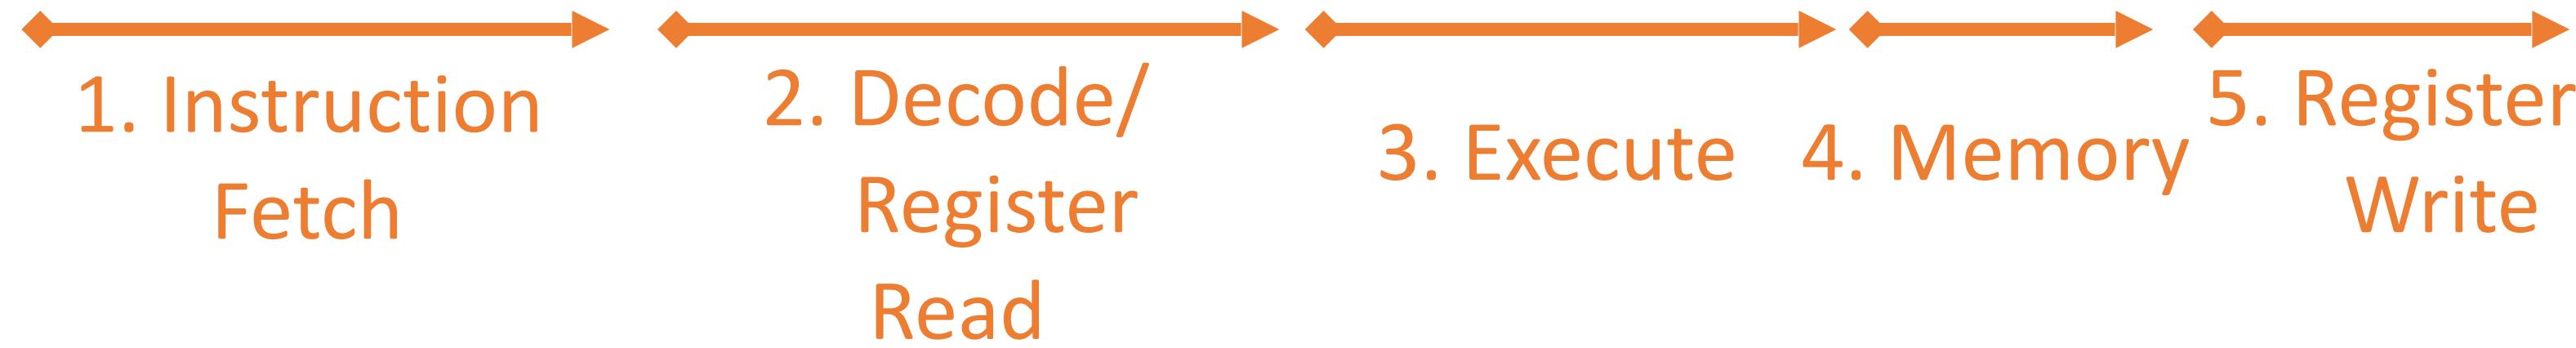
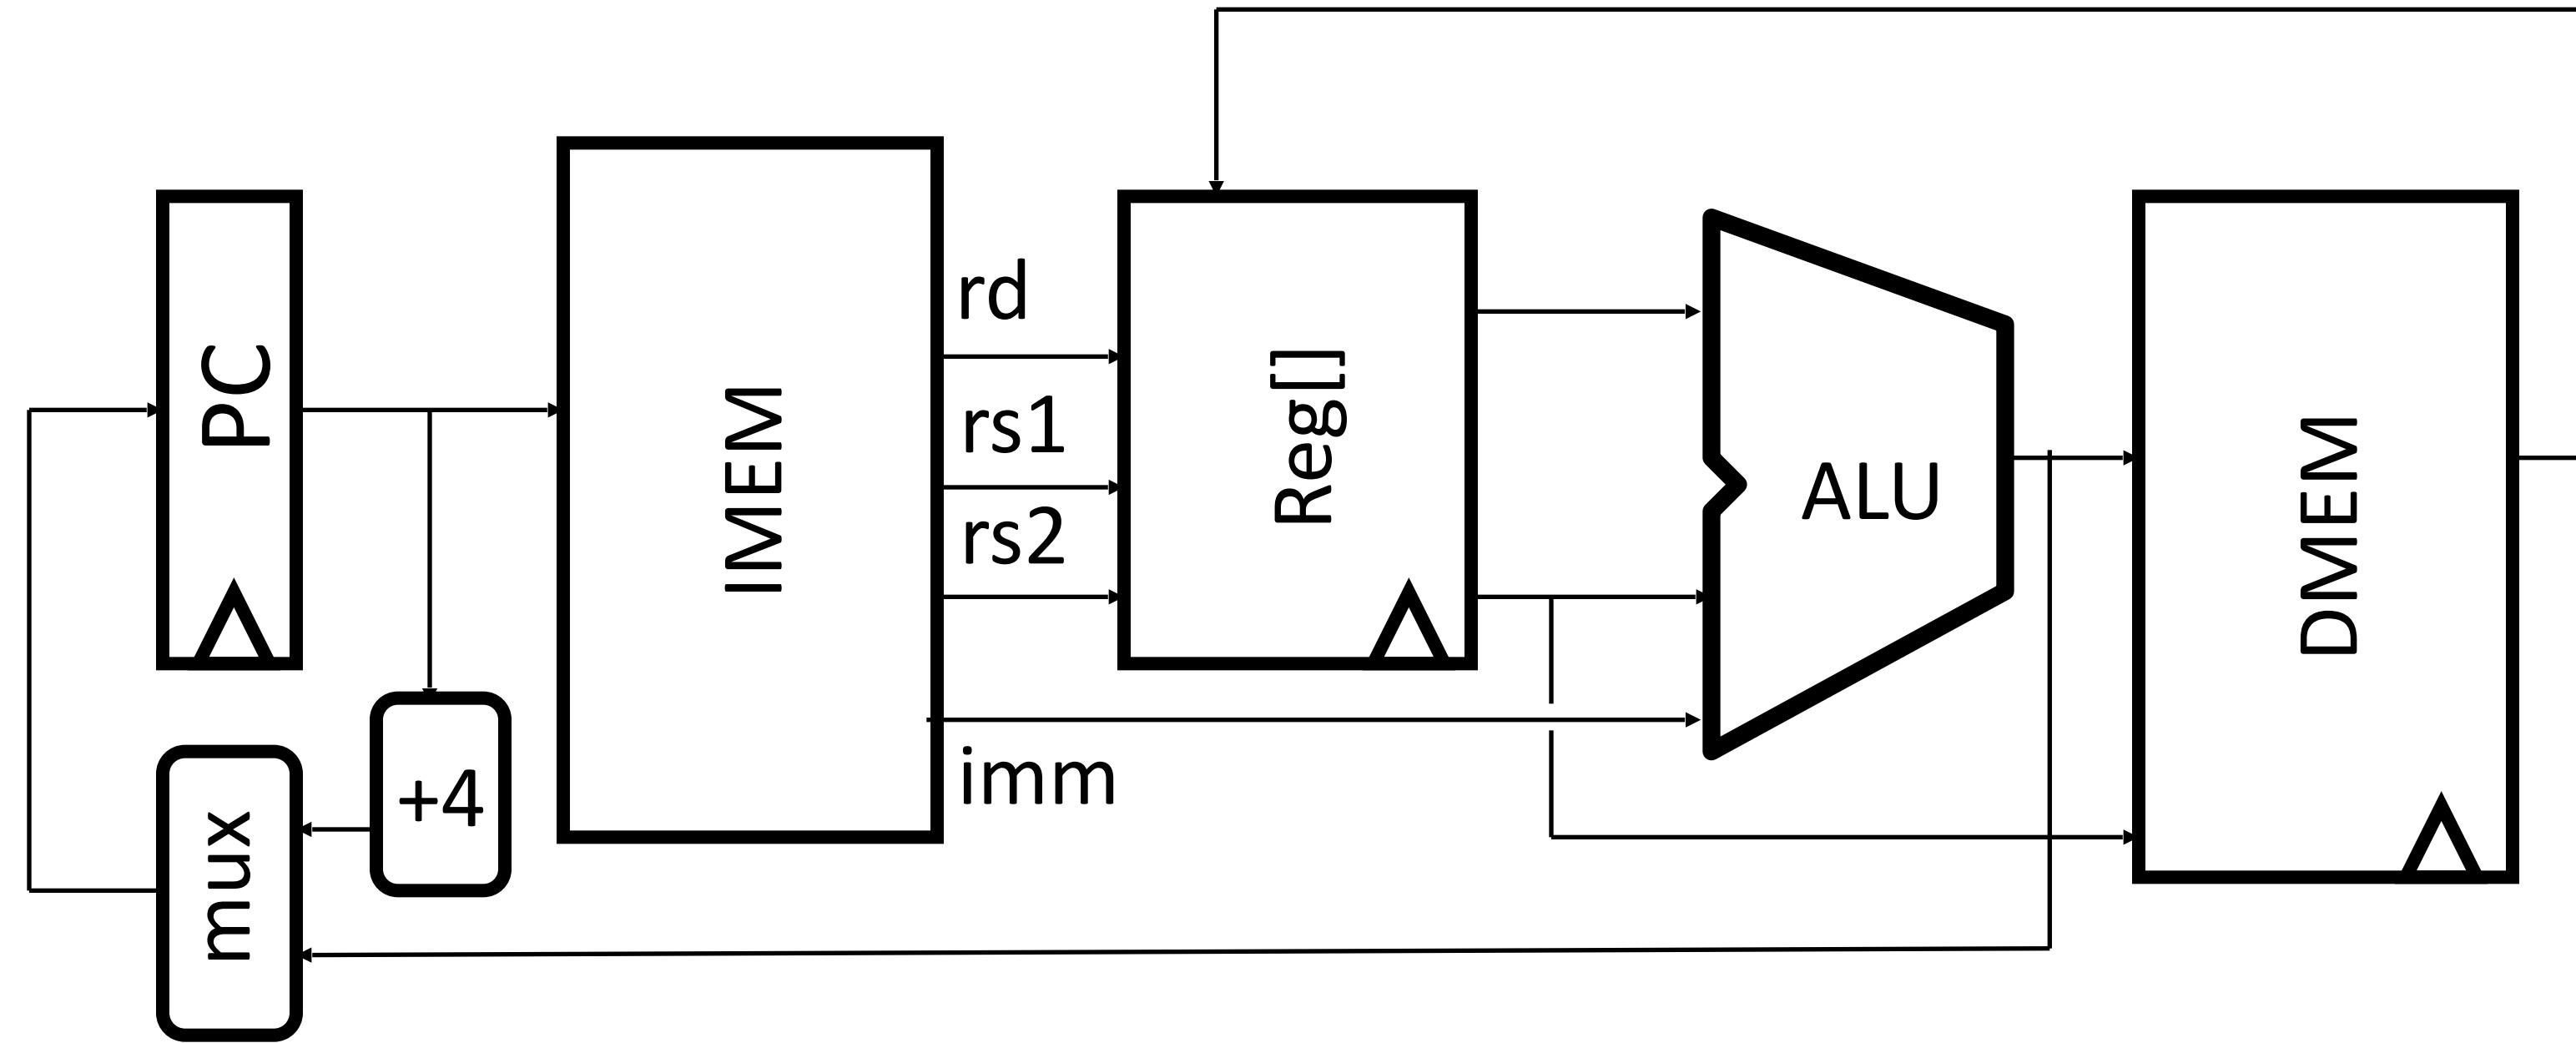
First Design: One-Instruction-Per-Cycle RISC-V Machine

On every tick of the clock, the computer executes one instruction



1. Current state outputs drive the inputs to the combinational logic, whose outputs settle at the values of the state before the next clock edge
2. At the rising clock edge, all the state elements are updated with the combinational logic outputs, and execution moves to the next clock cycle (next instruction)
13

Basic Phases of Instruction Execution



Implementing the add instruction

0000000	rs2	rs1	000	rd	0110011	ADD
---------	-----	-----	-----	----	---------	-----

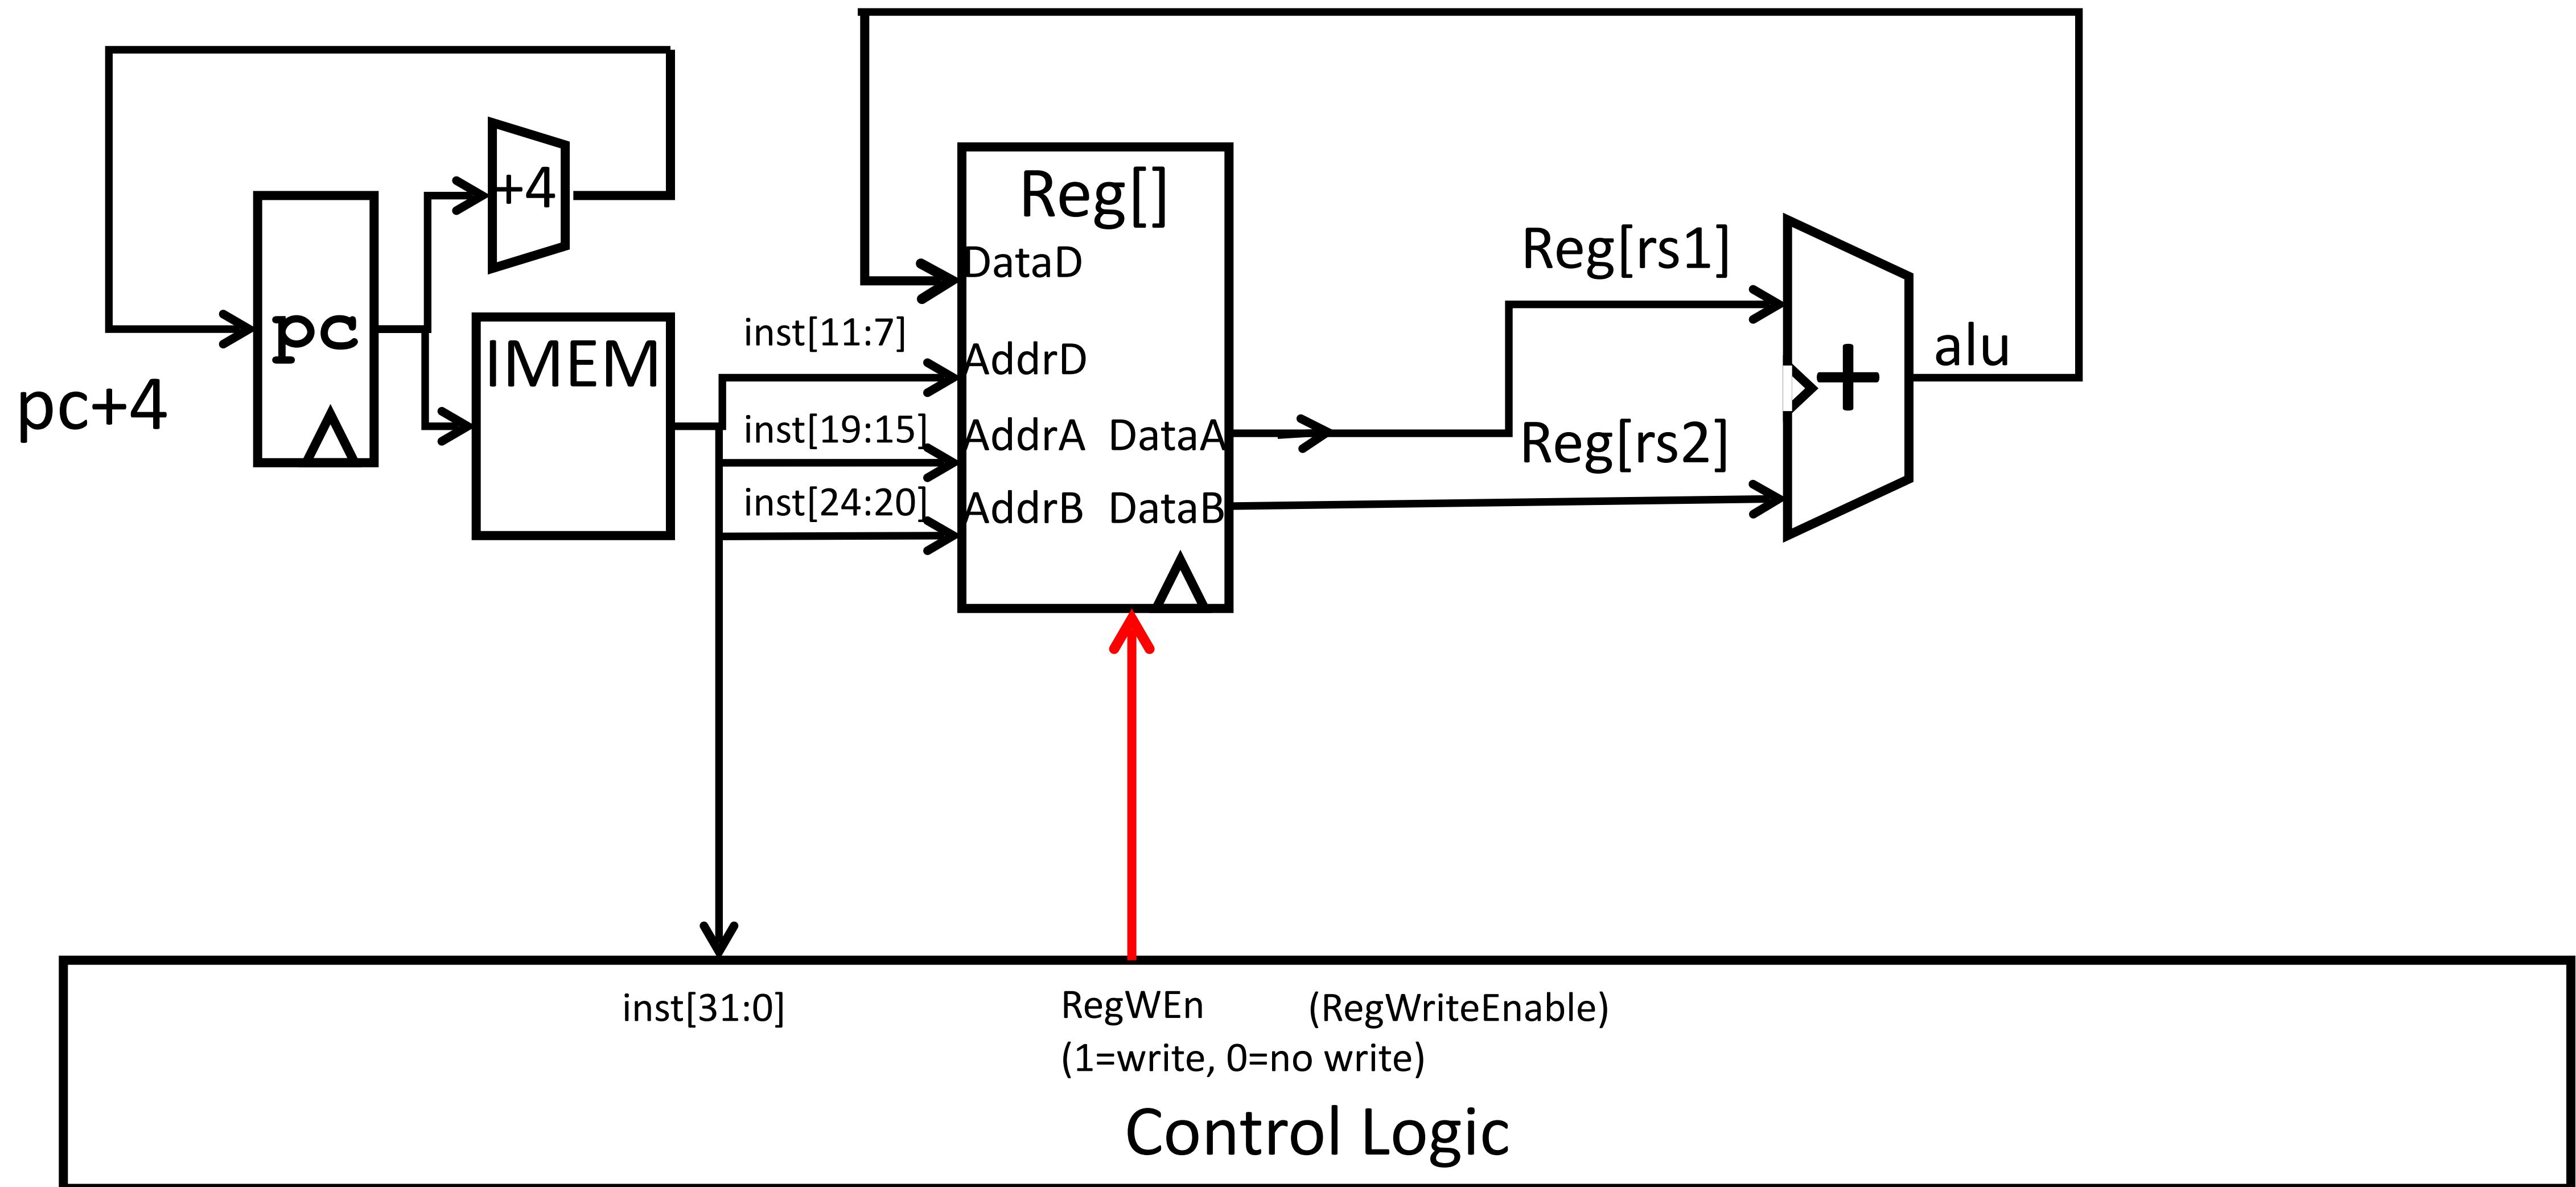
add rd, rs1, rs2

- Instruction makes two changes to machine's state:

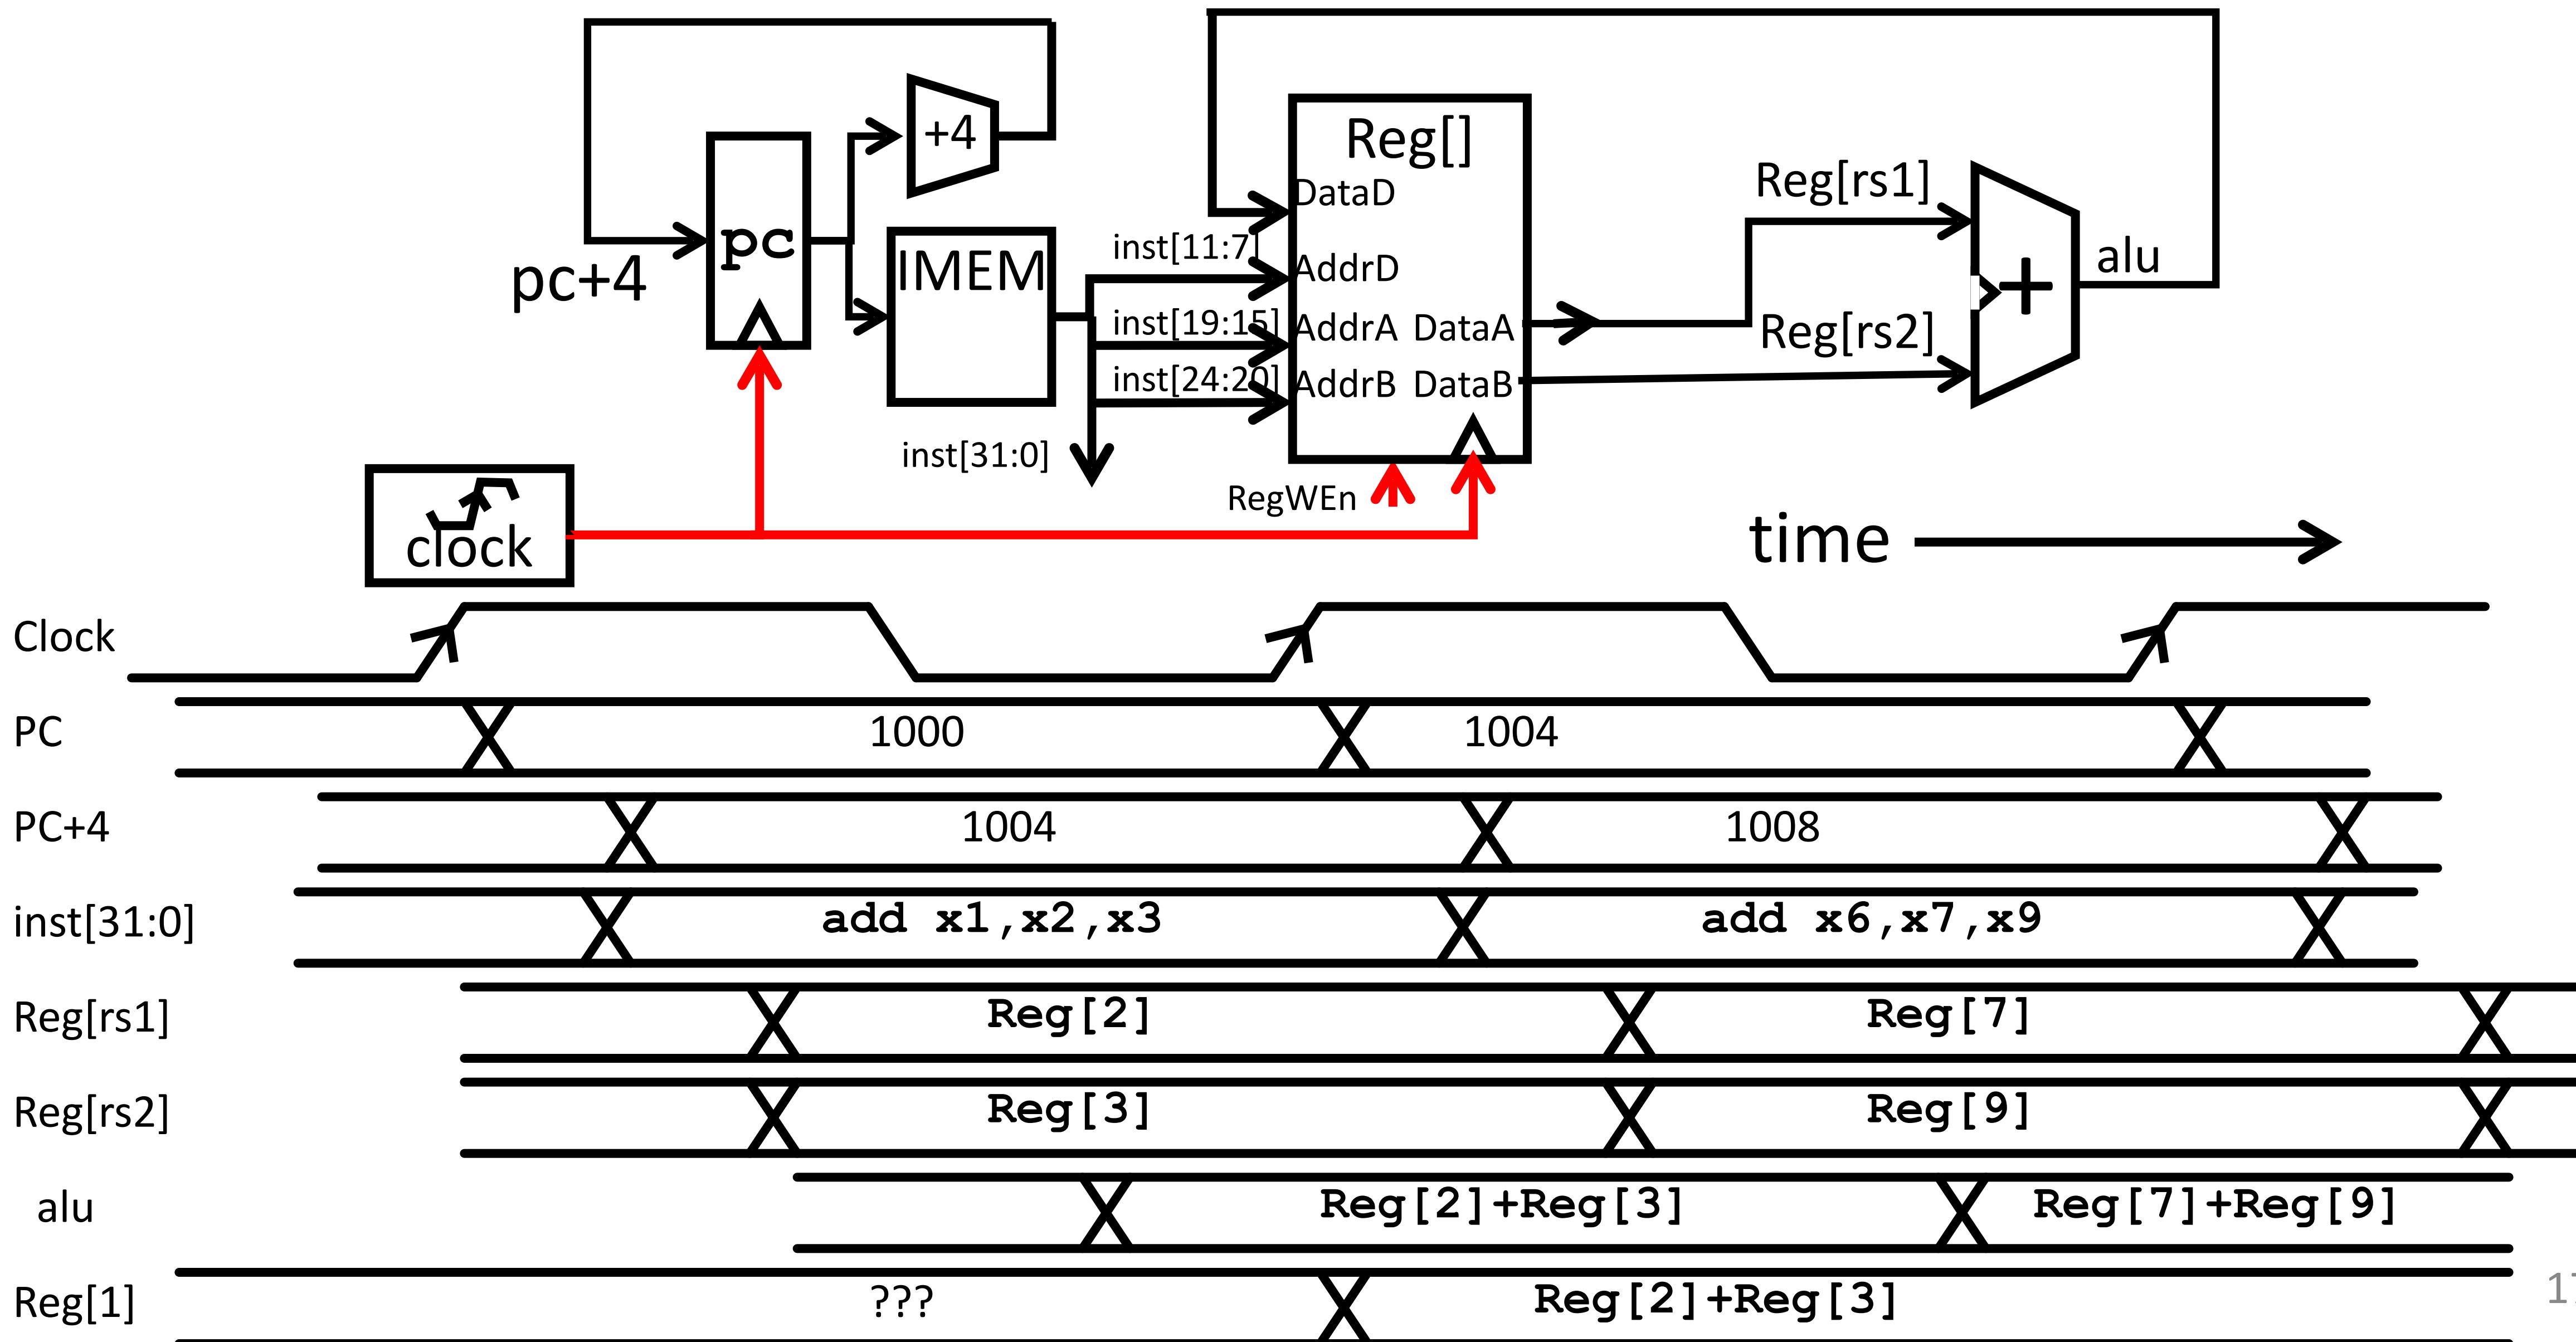
$$\text{Reg}[rd] = \text{Reg}[rs1] + \text{Reg}[rs2]$$

$$PC = PC + 4$$

Datapath for add



Timing Diagram for add



Implementing the **sub** instruction

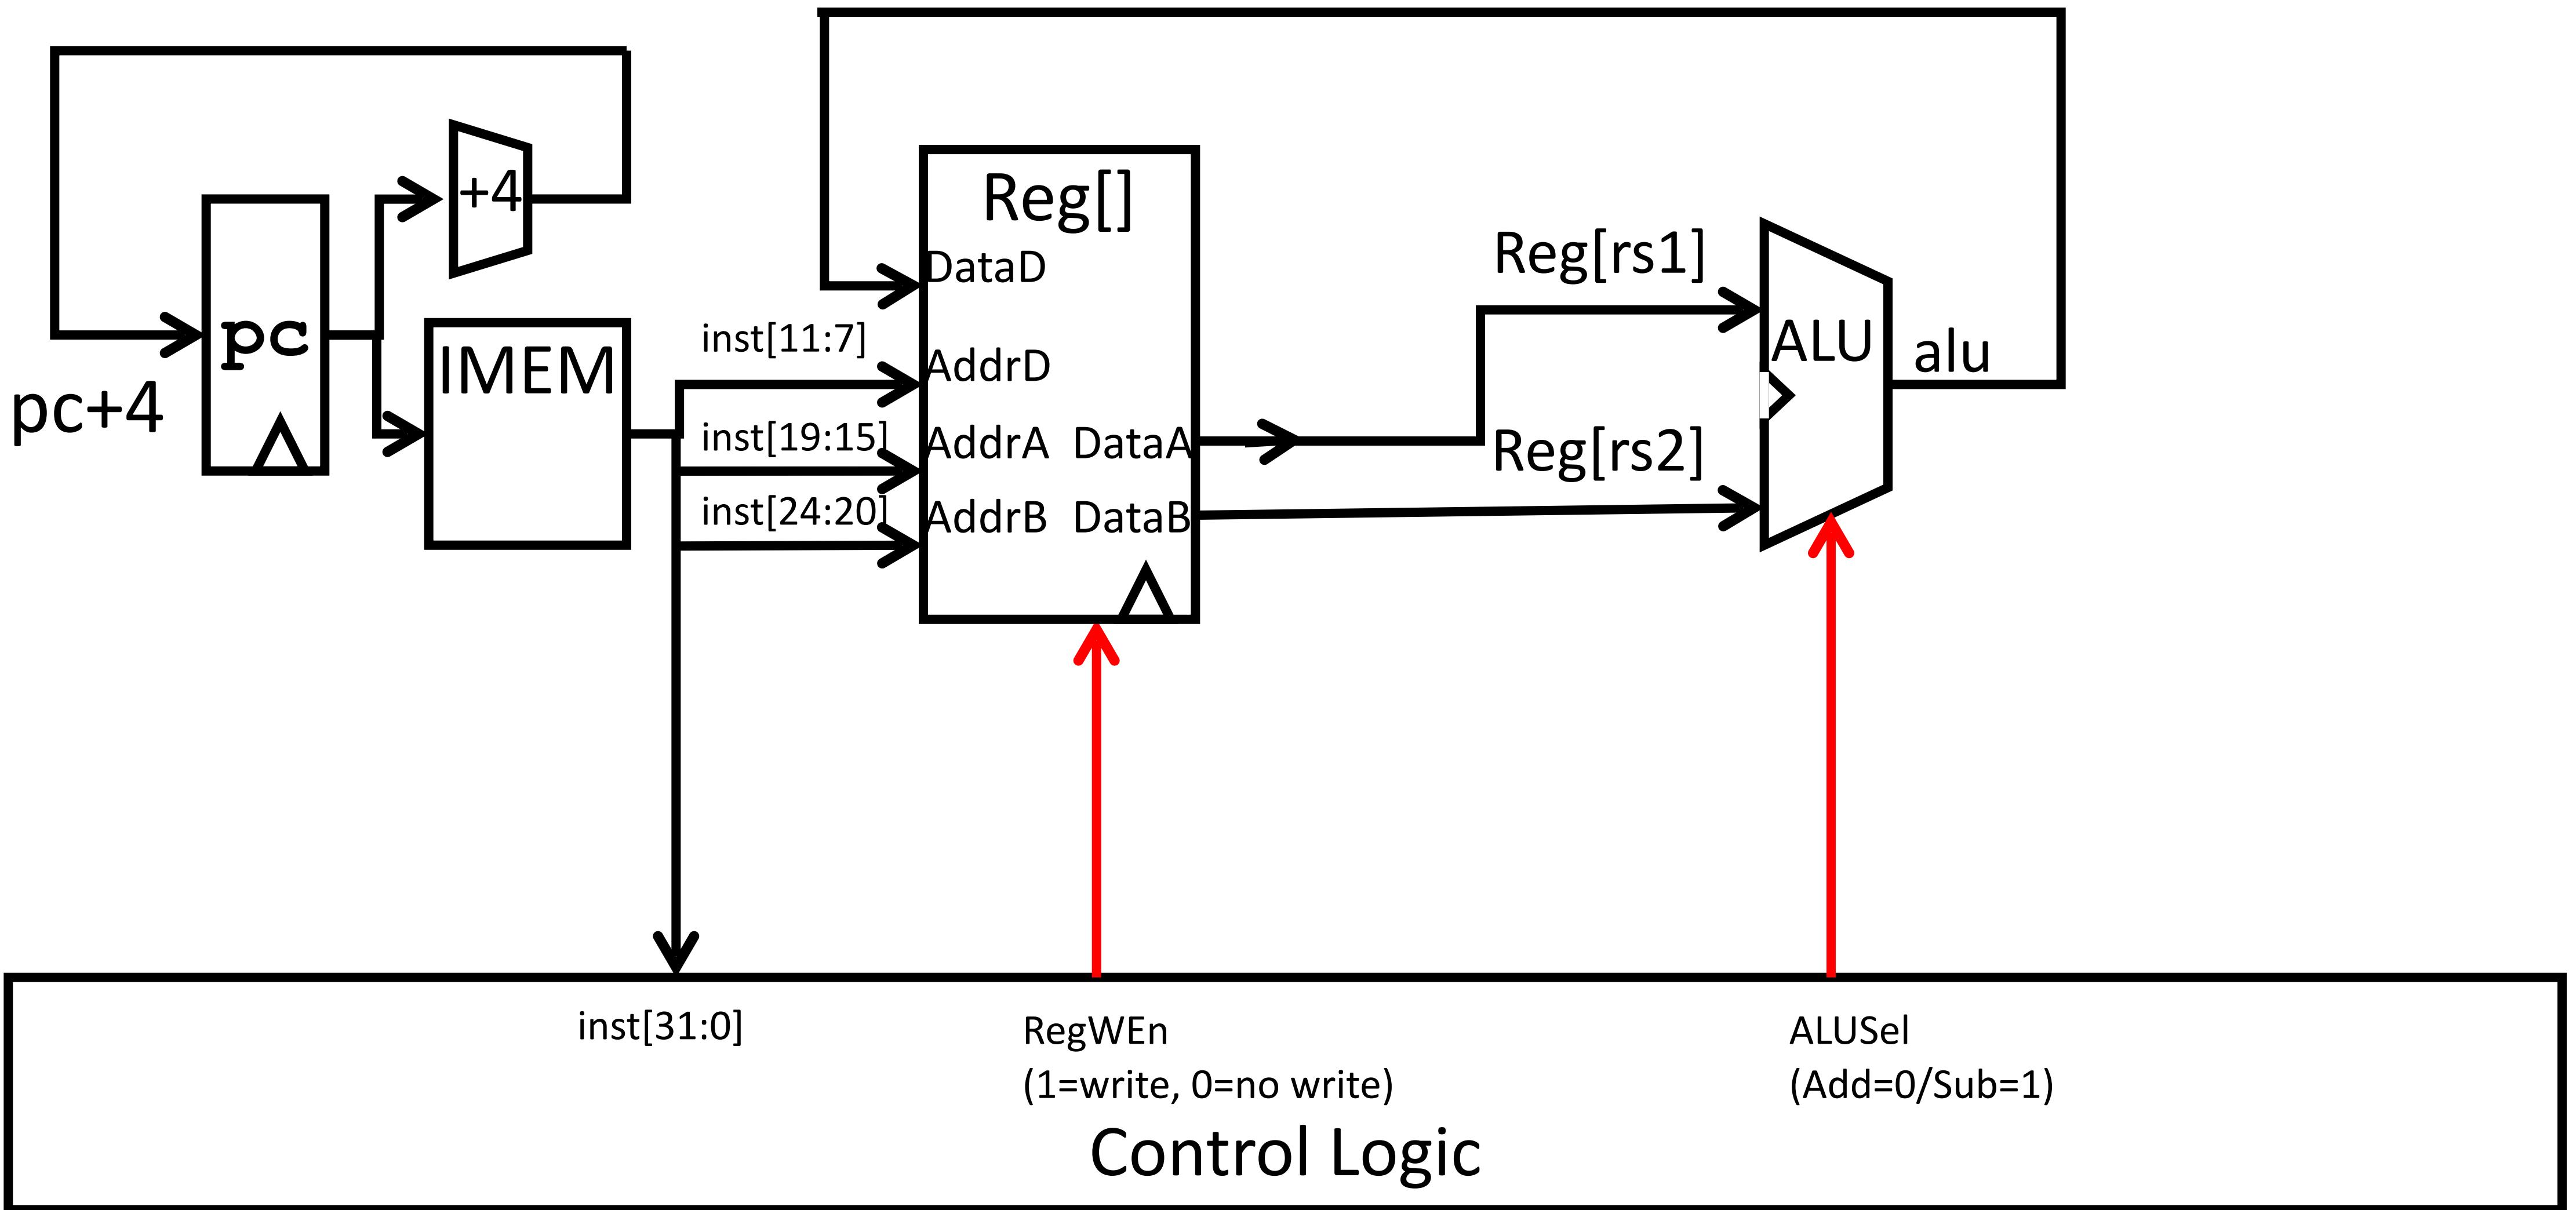
0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB

sub rd, rs1, rs2

$$\text{Reg}[rd] = \text{Reg}[rs1] - \text{Reg}[rs2]$$

- Almost the same as add, except now have to subtract operands instead of adding them
- `inst[30]` selects between add and subtract

Datapath for add/sub



Implementing other R-Format instructions

0000000	rs2	rs1	000	rd	0110011	ADD
0100000	rs2	rs1	000	rd	0110011	SUB
0000000	rs2	rs1	001	rd	0110011	SLL
0000000	rs2	rs1	010	rd	0110011	SLT
0000000	rs2	rs1	011	rd	0110011	SLTU
0000000	rs2	rs1	100	rd	0110011	XOR
0000000	rs2	rs1	101	rd	0110011	SRL
0100000	rs2	rs1	101	rd	0110011	SRA
0000000	rs2	rs1	110	rd	0110011	OR
0000000	rs2	rs1	111	rd	0110011	AND

- All implemented by decoding funct3 and funct7 fields and selecting appropriate ALU function

Implementing the addi instruction

- RISC-V Assembly Instruction: *Uses the “I-type” instruction format*

addi rd, rs1, integer

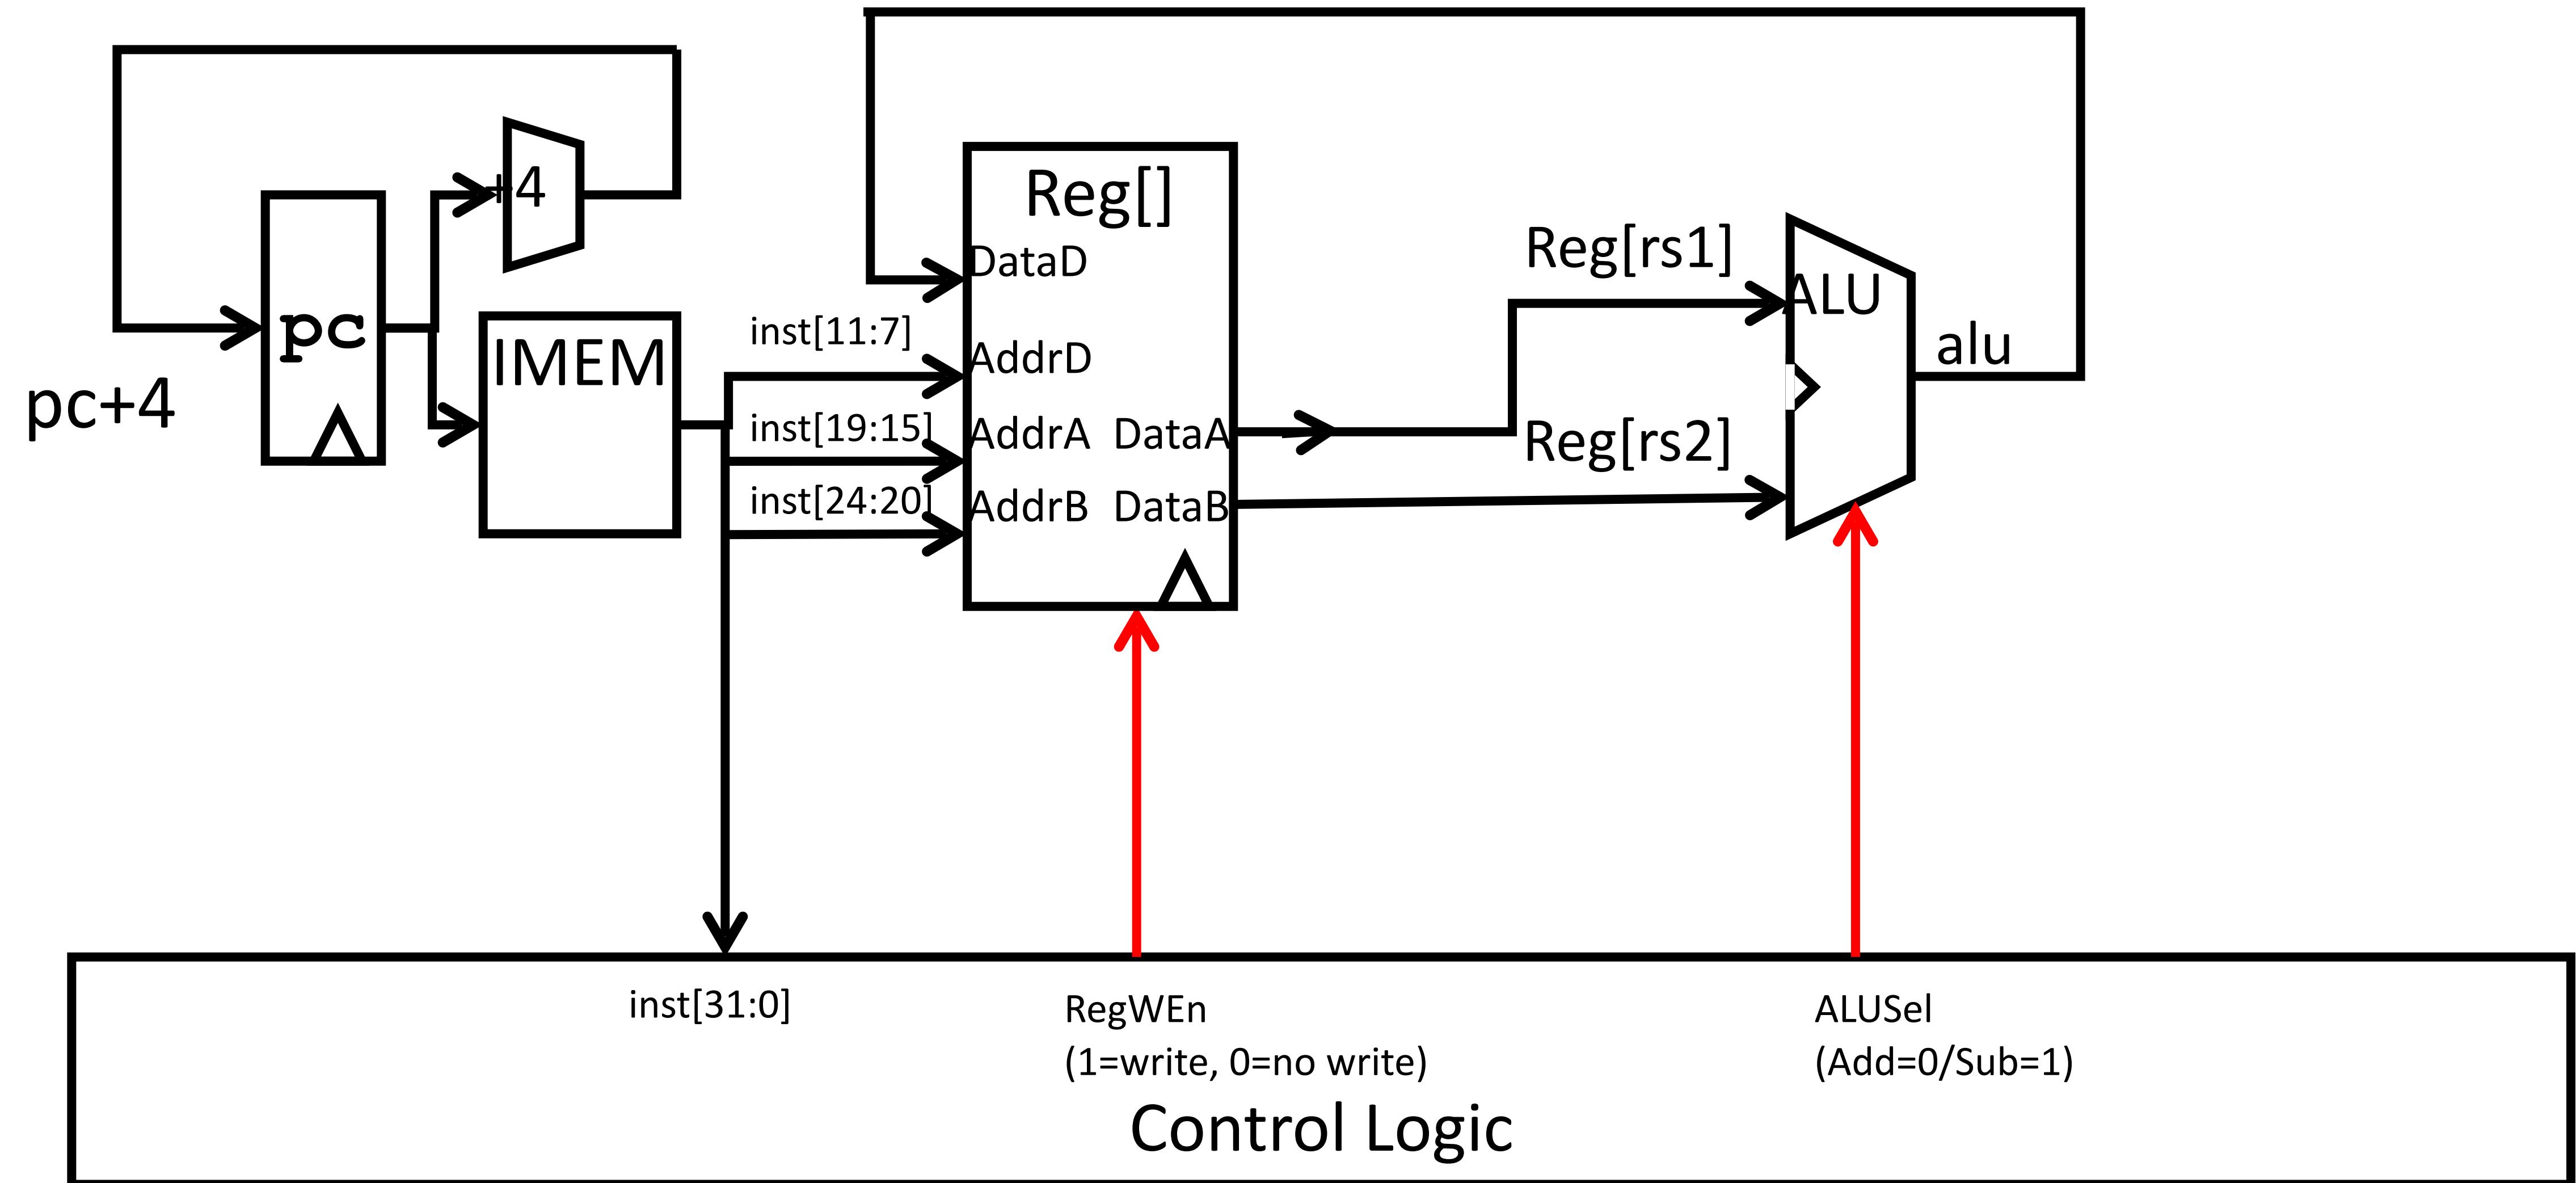
$Reg[rd] = Reg[rs1] + \text{sign_extend}(\text{immediate})$

example:

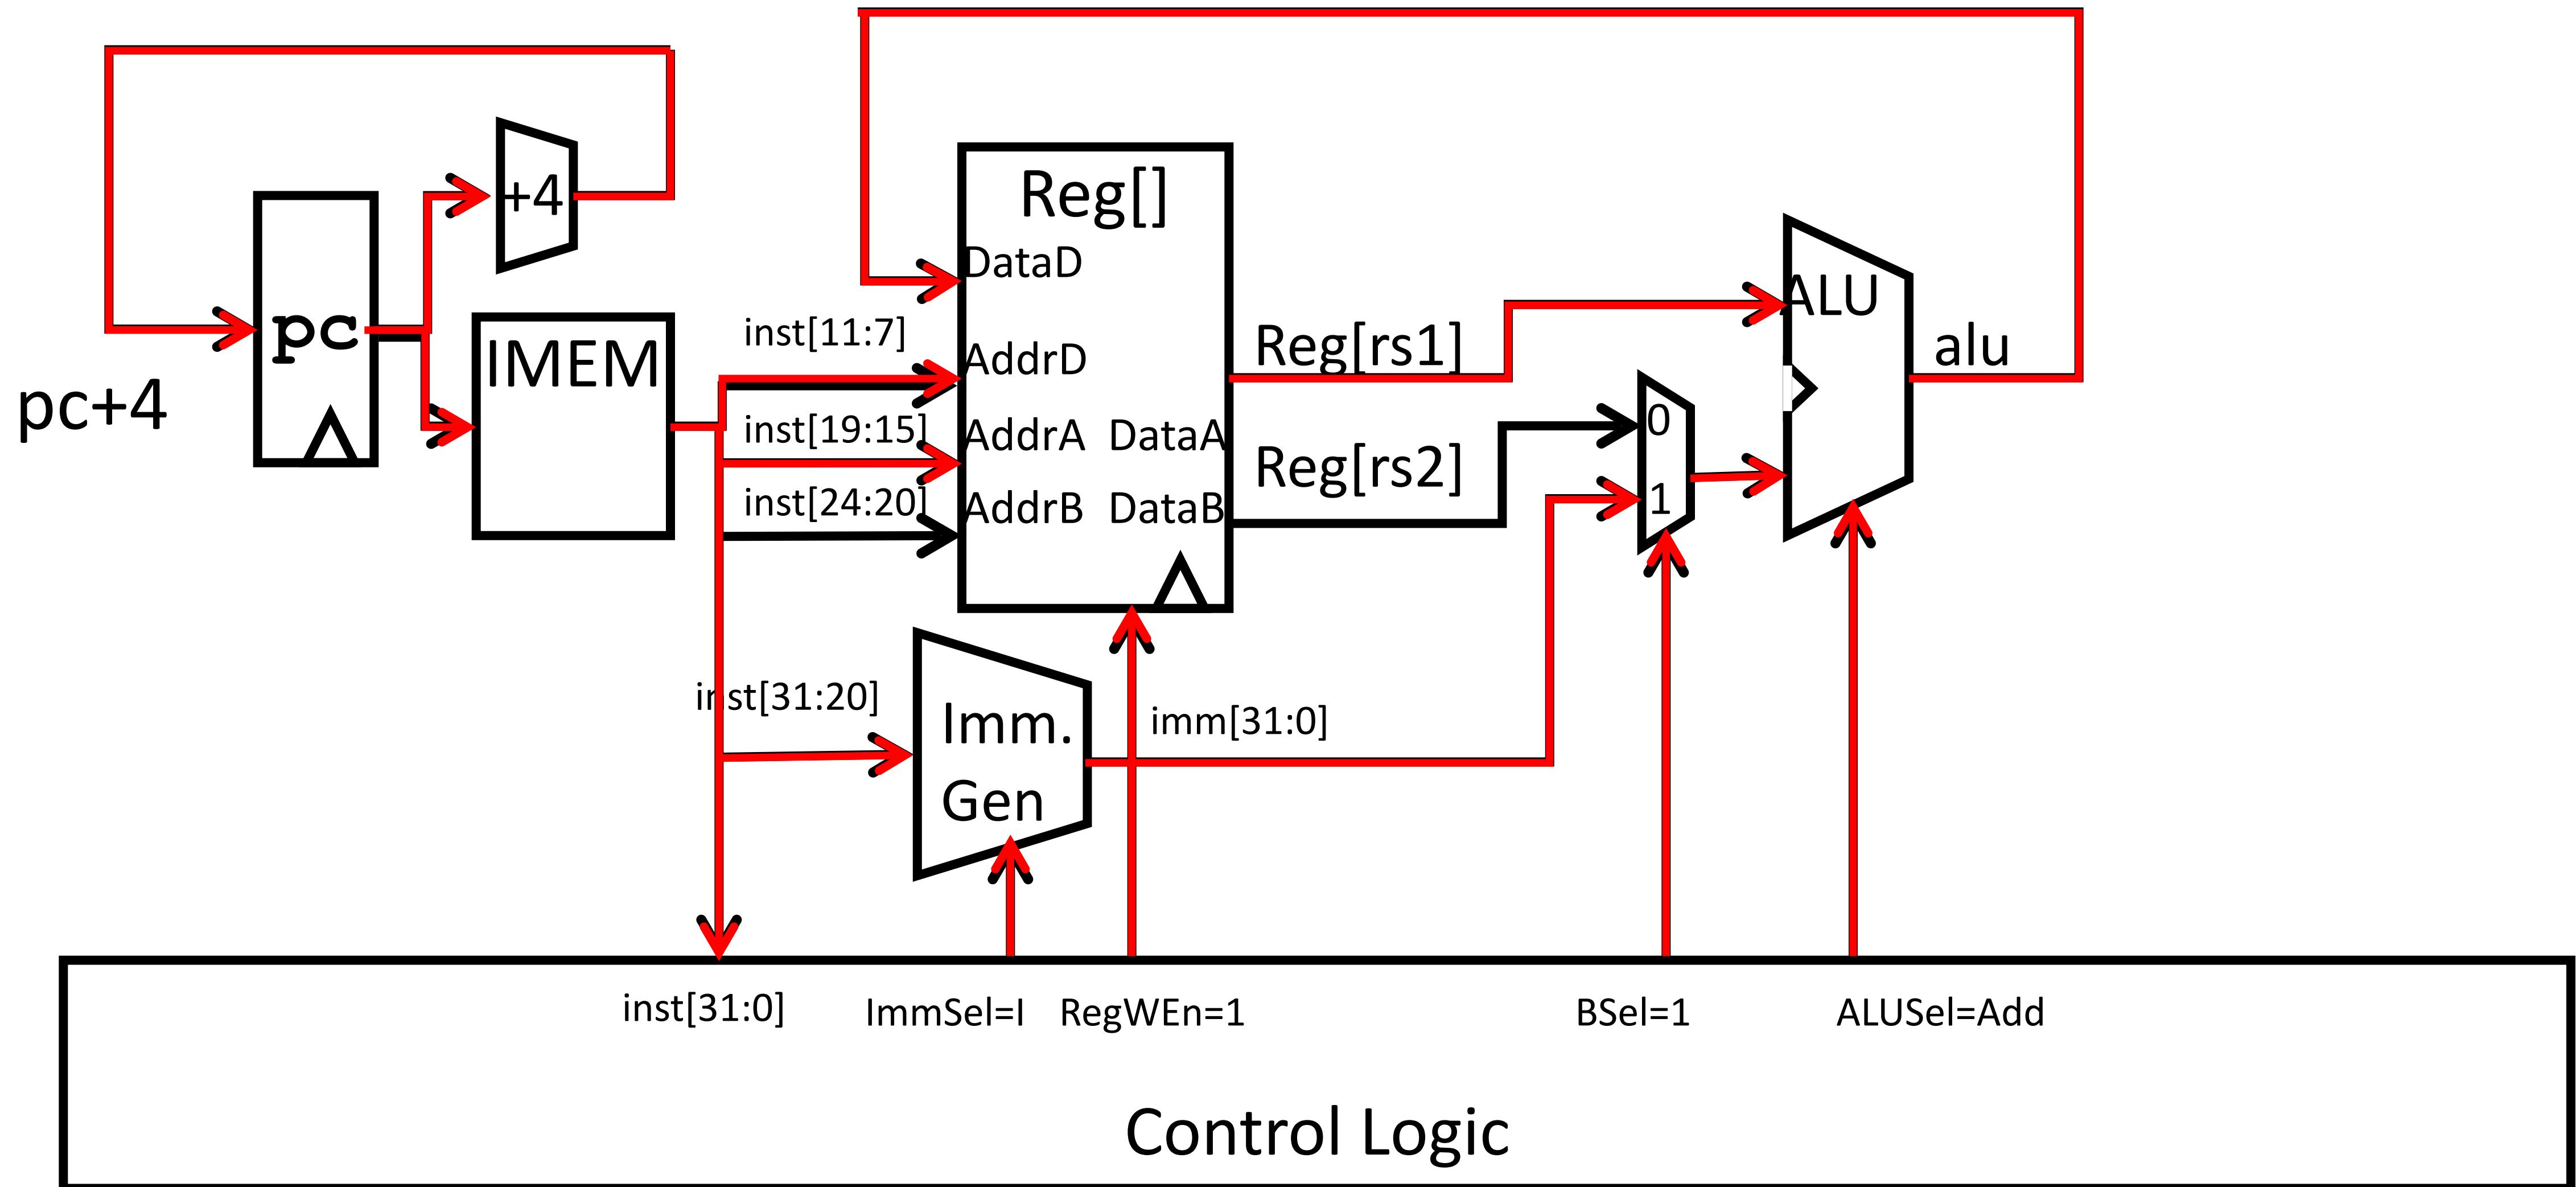
addi x15, x1, -50

31	20 19	15 14	12 11	7 6	0
imm[11:0]	rs1	funct3	rd	opcode	
12	5	3	5	7	0
111111001110	00001	000	01111	0010011	
imm=-50		rs1=1	ADD	rd=15	OP-Imm

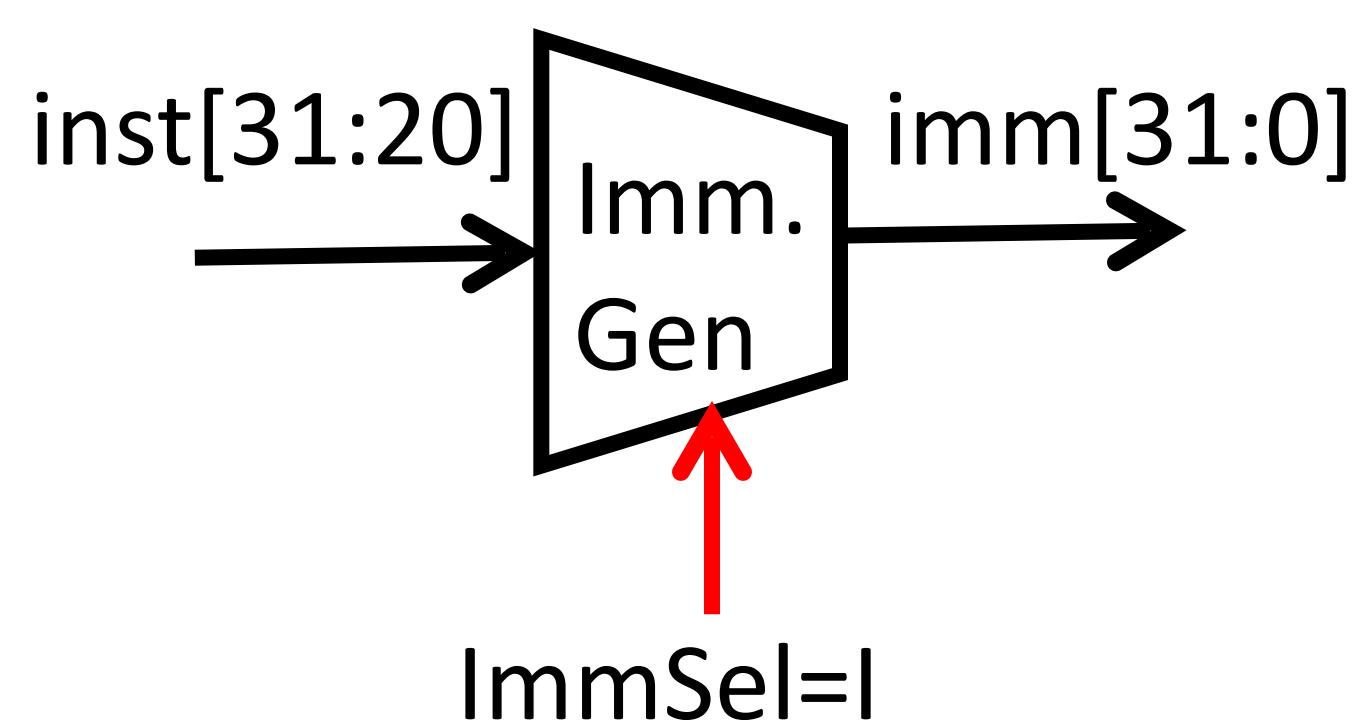
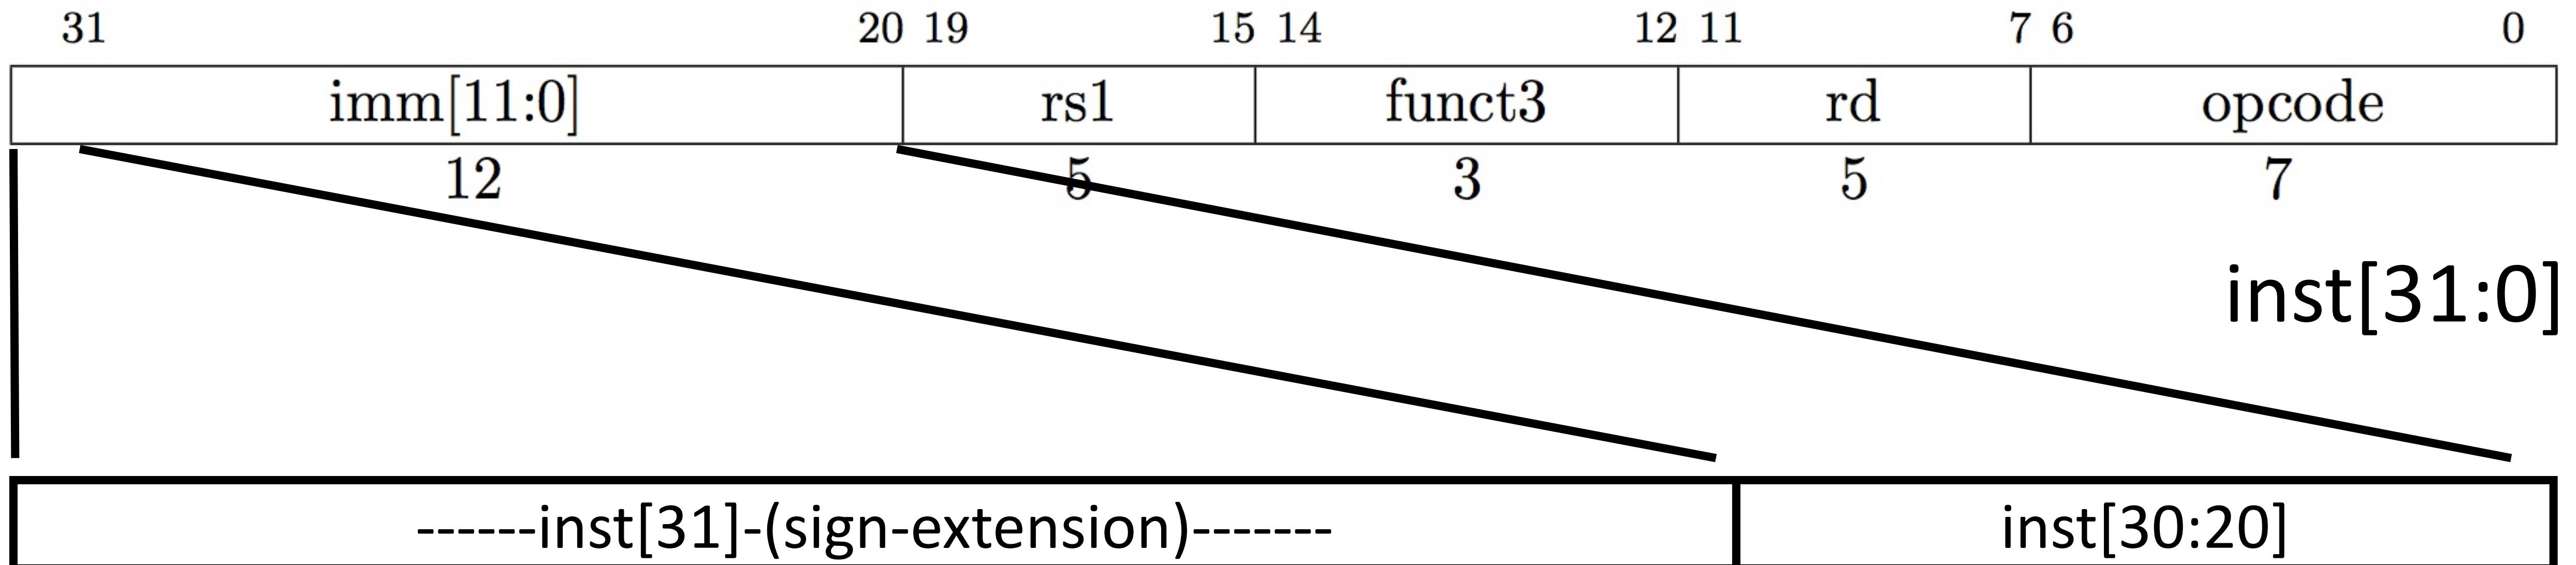
Review: Datapath for add/sub



Adding addi to datapath

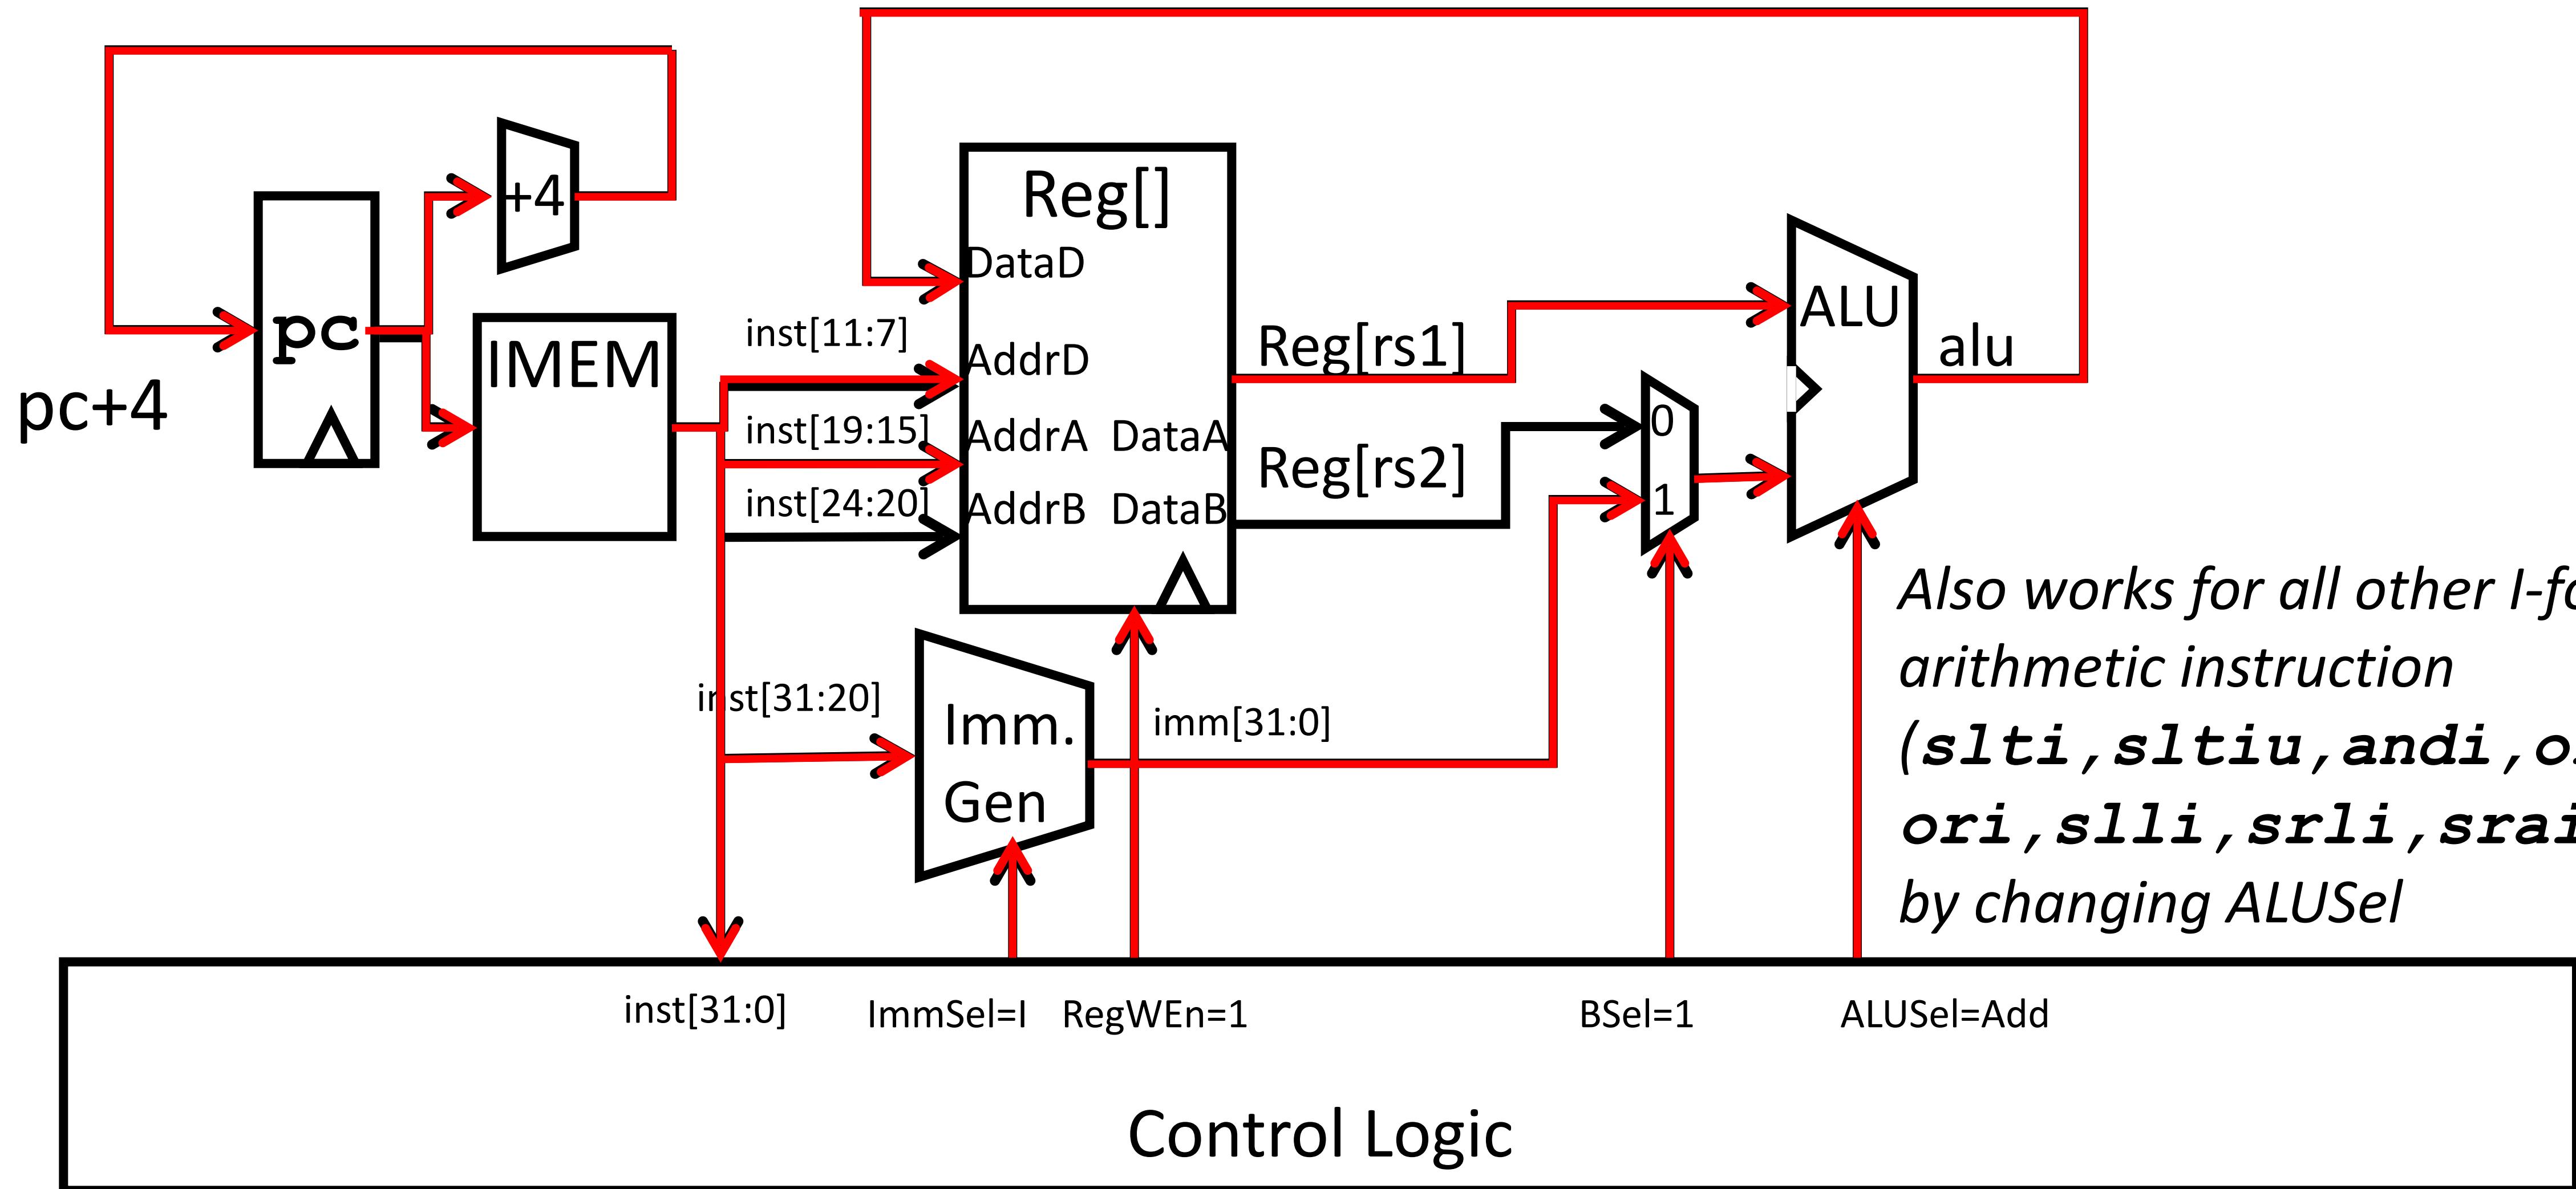


I-type Format immediates



- High 12 bits of instruction (inst[31:20]) copied to low 12 bits of immediate (imm[11:0])
- Immediate is sign-extended by copying value of inst[31] to fill the upper 20 bits of the immediate value (imm[31:12])

Adding addi to datapath



Implementing Load Word instruction

- RISC-V Assembly Instruction: *Also uses the “I-type” instruction format*

lw rd, integer(rs1)

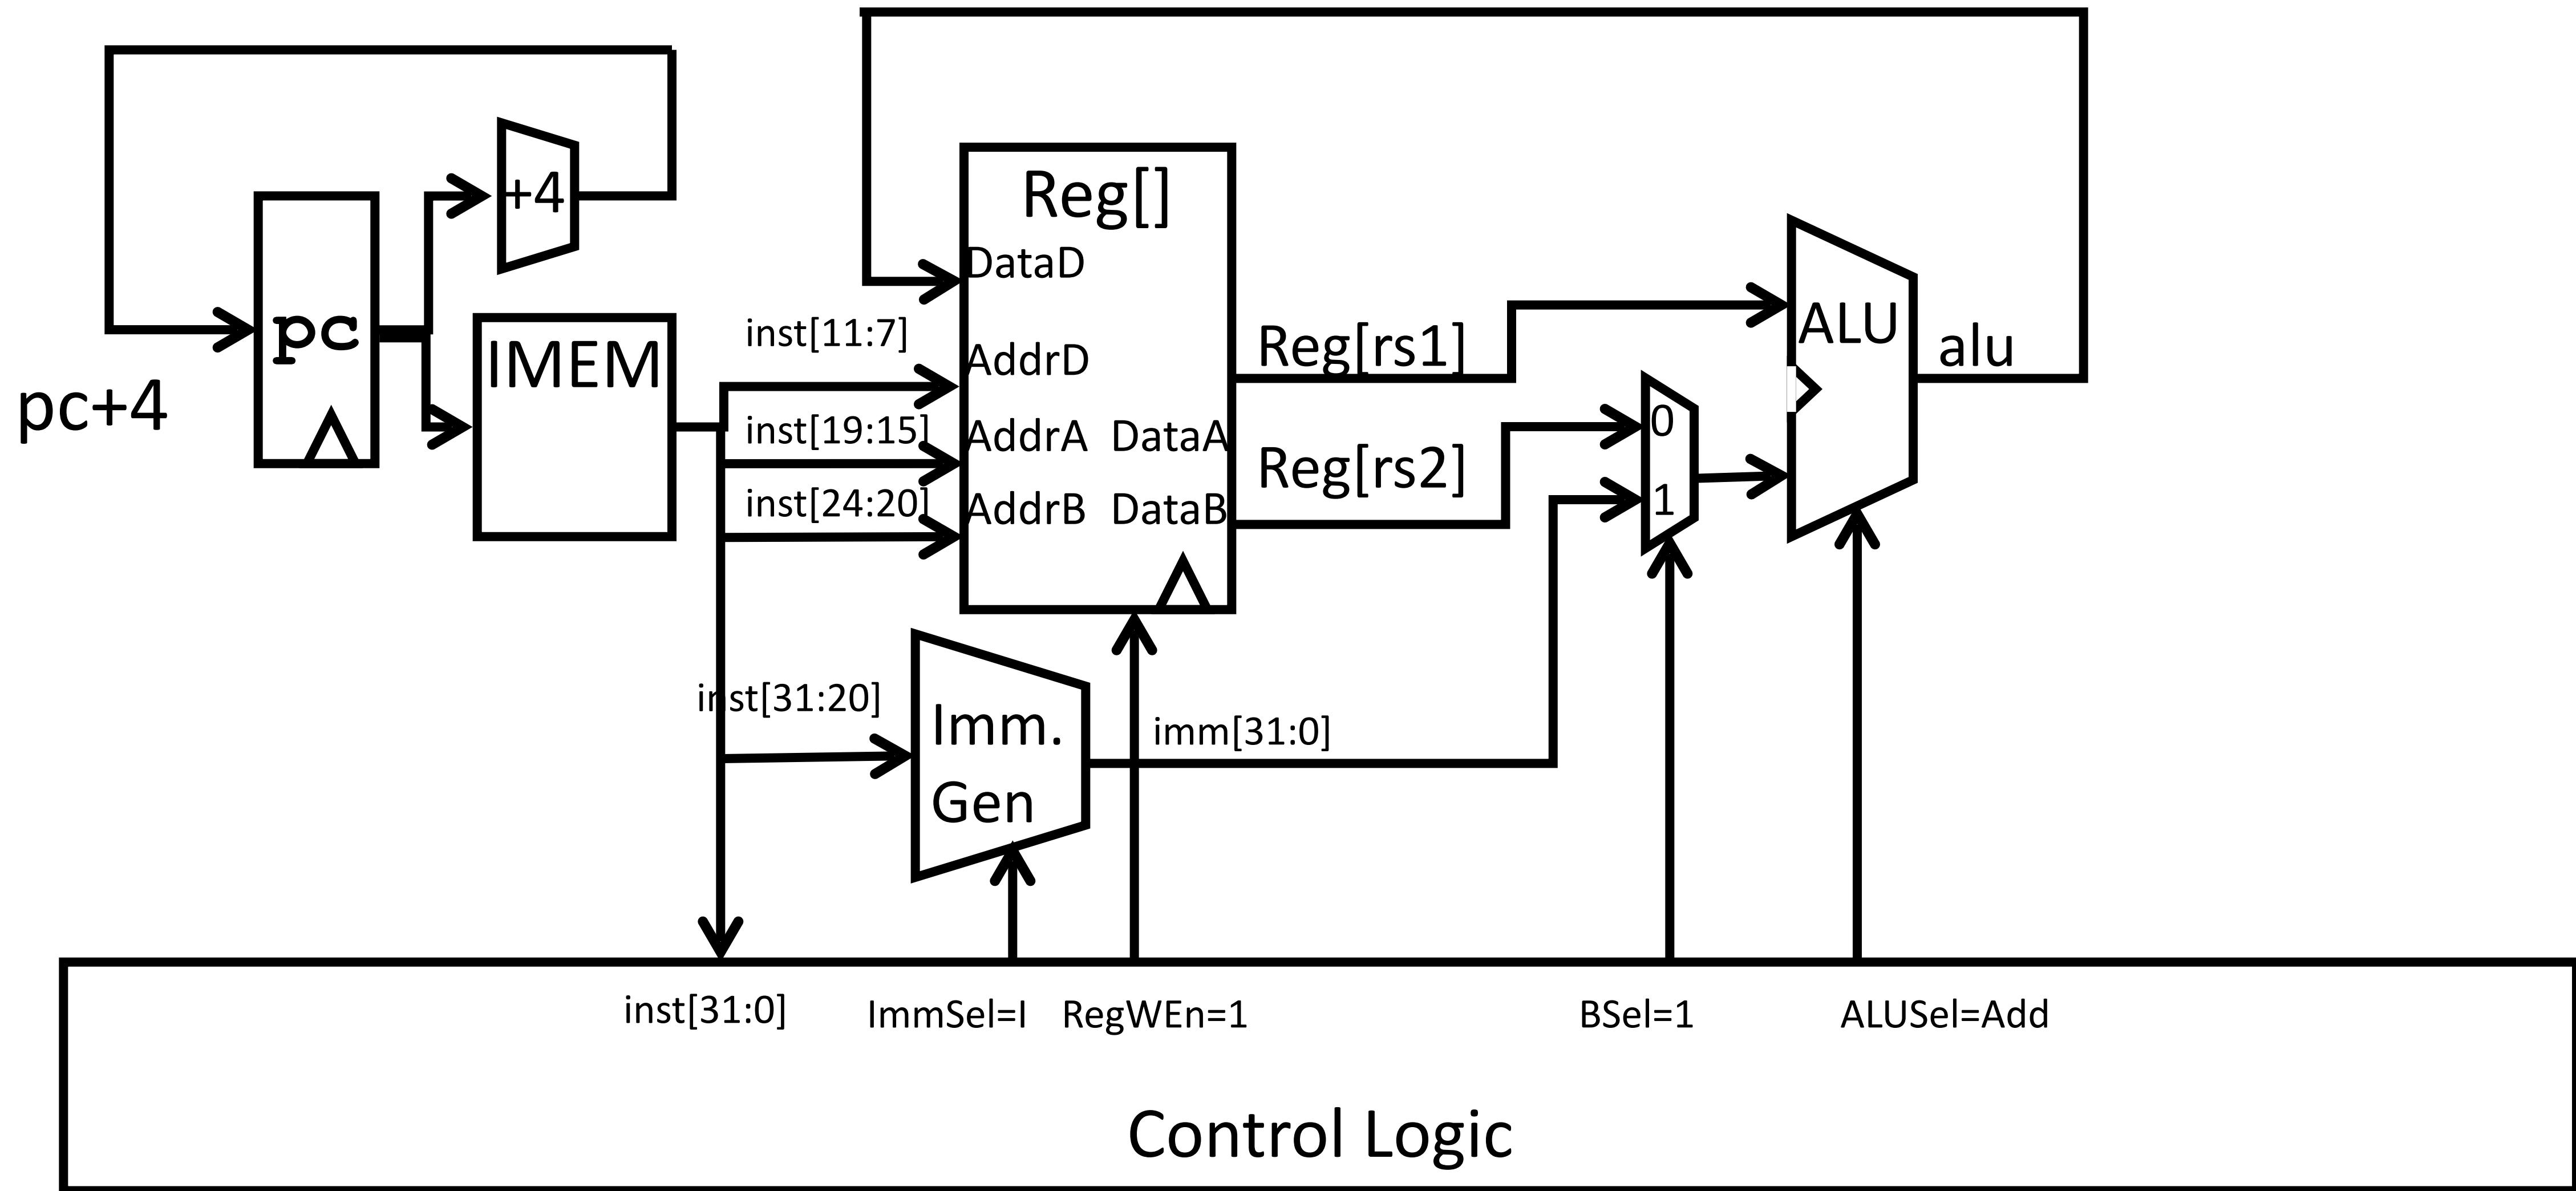
$Reg[rd] = DMEM[Reg[rs1]] + sign_extend(immediate)$

example:

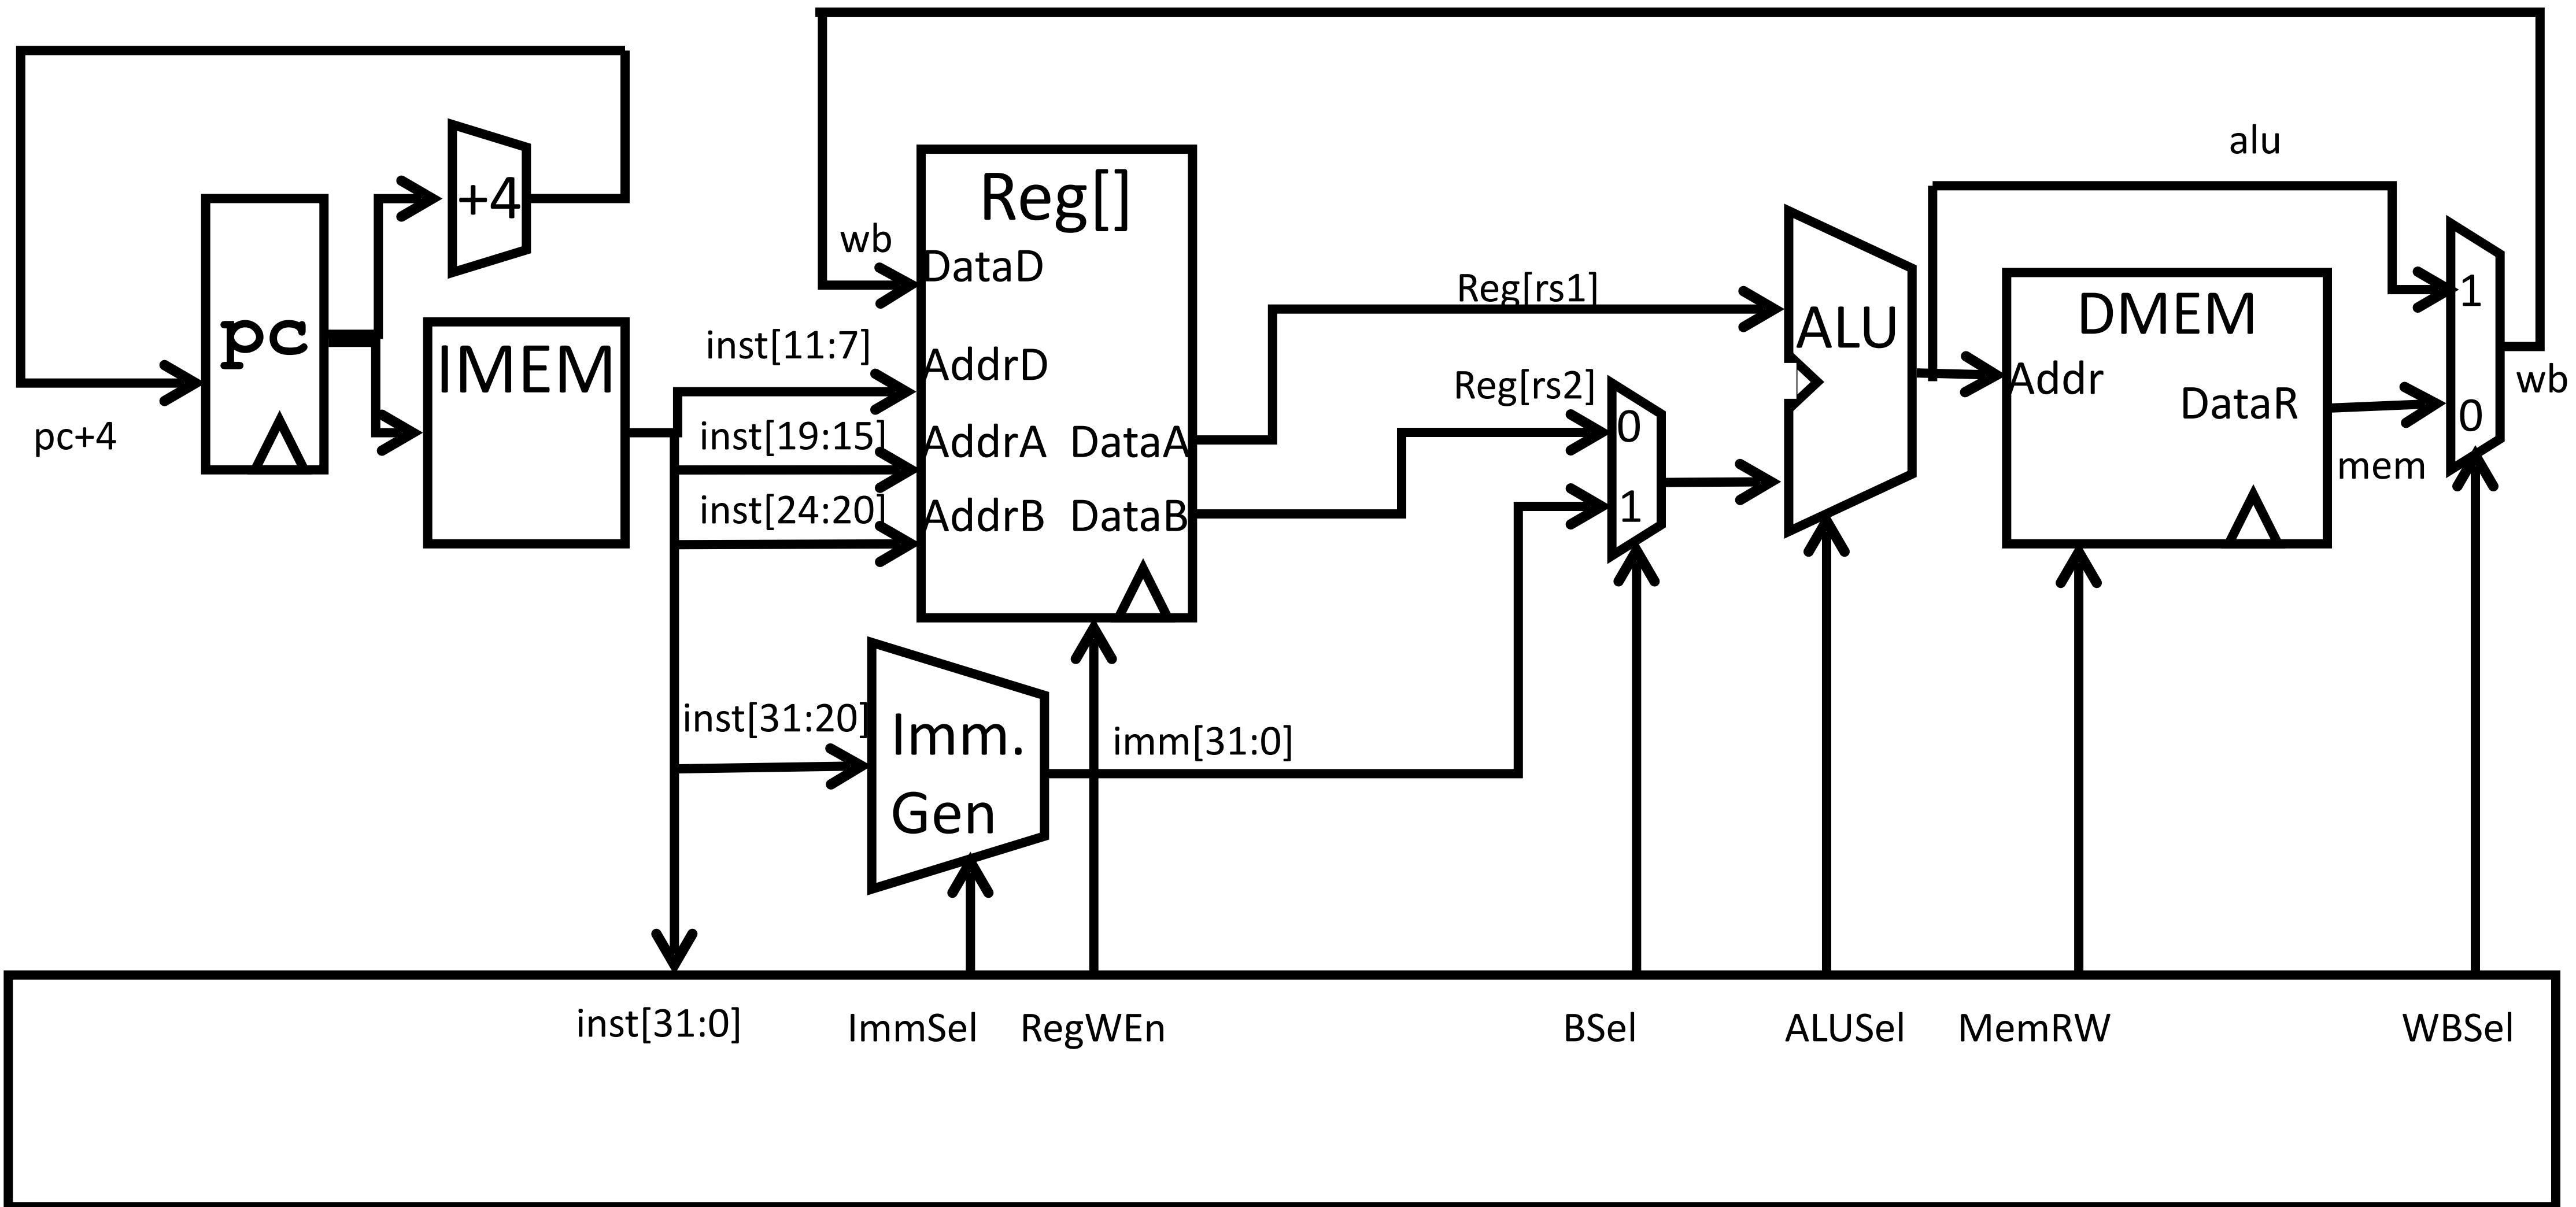
addi x14,8(x2)

31	20 19	15 14	12 11	7 6	0
imm[11:0]	rs1	funct3	rd	opcode	
12	5	3	5	7	0
000000001000	00010	010	01110	0000011	
imm=+8		rs1=2	LW	rd=14	LOAD

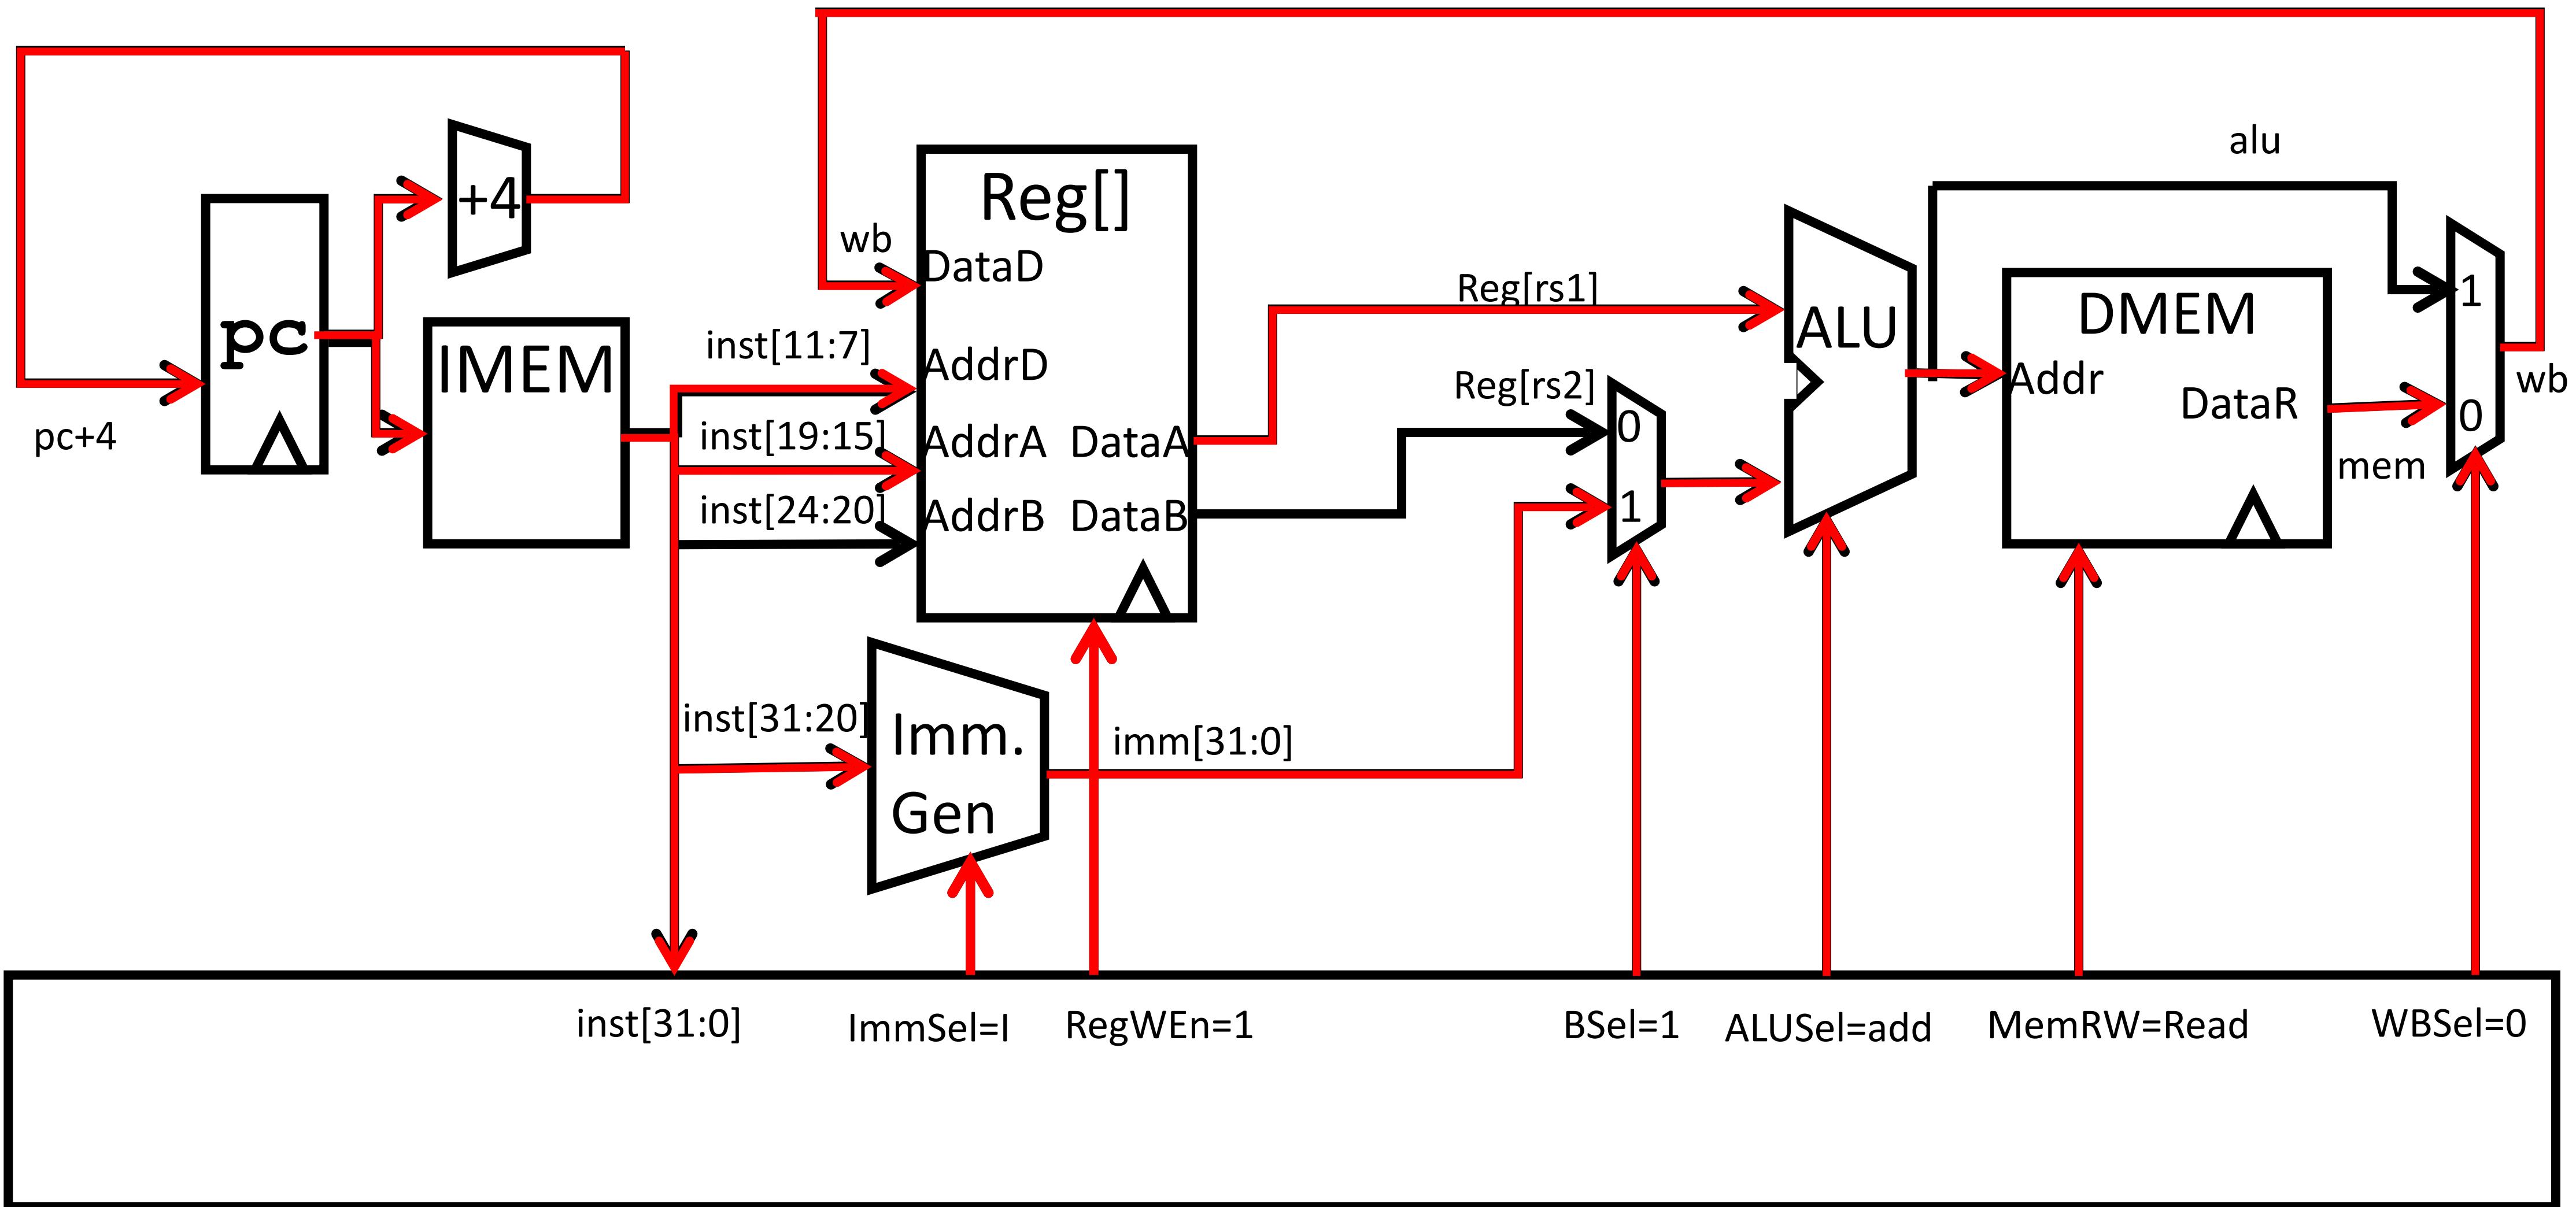
Review: Adding addi to datapath



Adding 1w to datapath



Adding 1w to datapath



All RV32 Load Instructions

imm[11:0]	rs1	000	rd	0000011	LB
imm[11:0]	rs1	001	rd	0000011	LH
imm[11:0]	rs1	010	rd	0000011	LW
imm[11:0]	rs1	100	rd	0000011	LBU
imm[11:0]	rs1	101	rd	0000011	LHU

↑ funct3 field encodes size and signedness of load data

- Supporting the narrower loads requires additional circuits to extract the correct byte/halfword from the value loaded from memory, and sign- or zero-extend the result to 32 bits before writing back to register file.

Implementing Store Word instruction

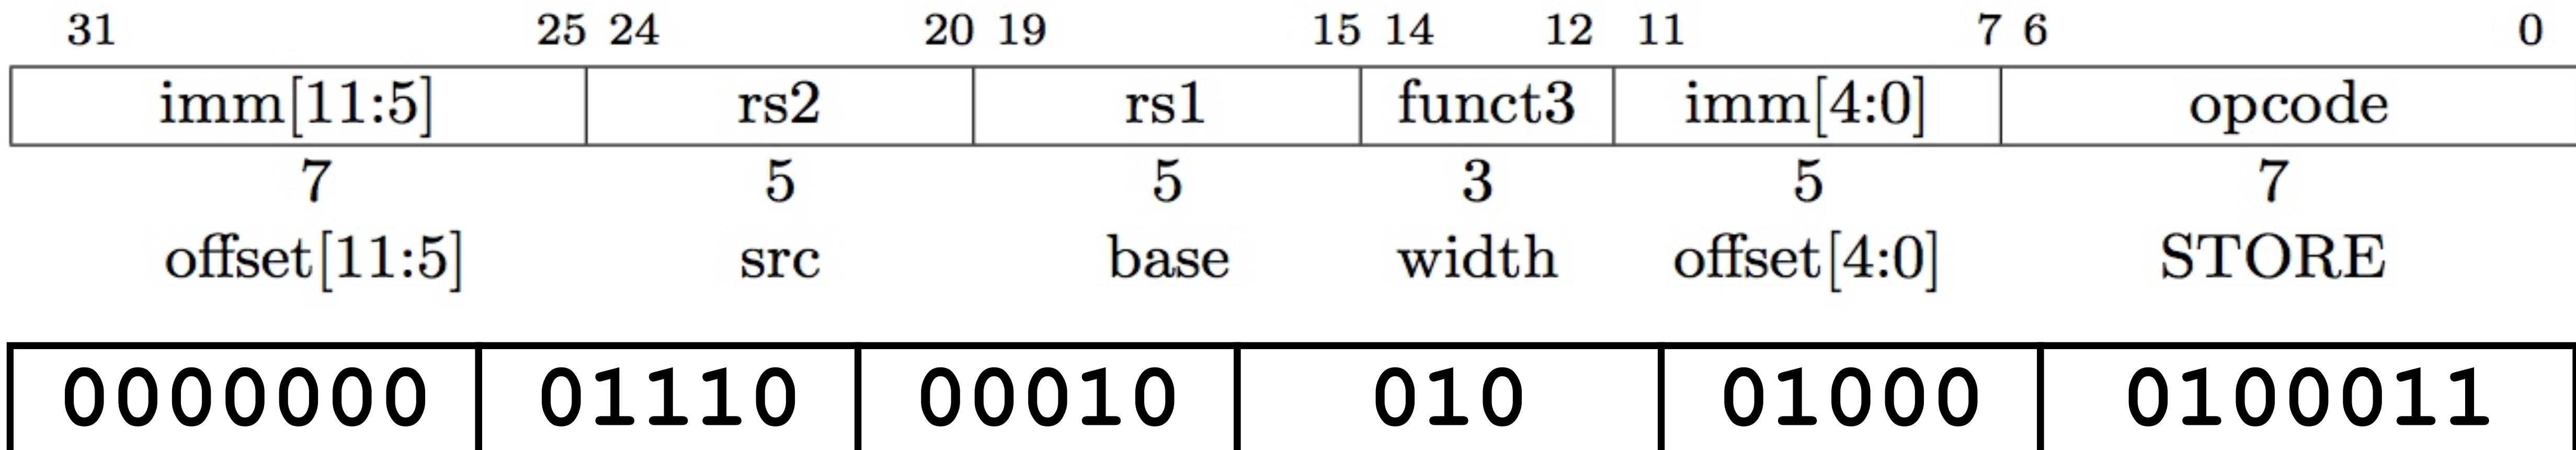
- RISC-V Assembly Instruction: *Uses the “S-type” instruction format*

```
sw rs2, integer(rs1
```

$$DDEM[Reg[rs1] + sign\ extend(immediate)] = Reg[rs2]$$

example:

sw x14, 8(x2)



offset[11:5] rs2=14 rs1=2 SW offset[4:0] STORE

= 0

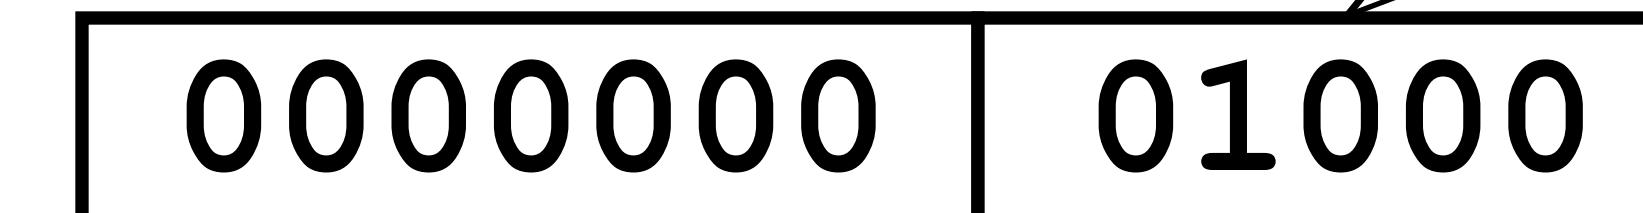
rs1=2

SW

offset[4:0]

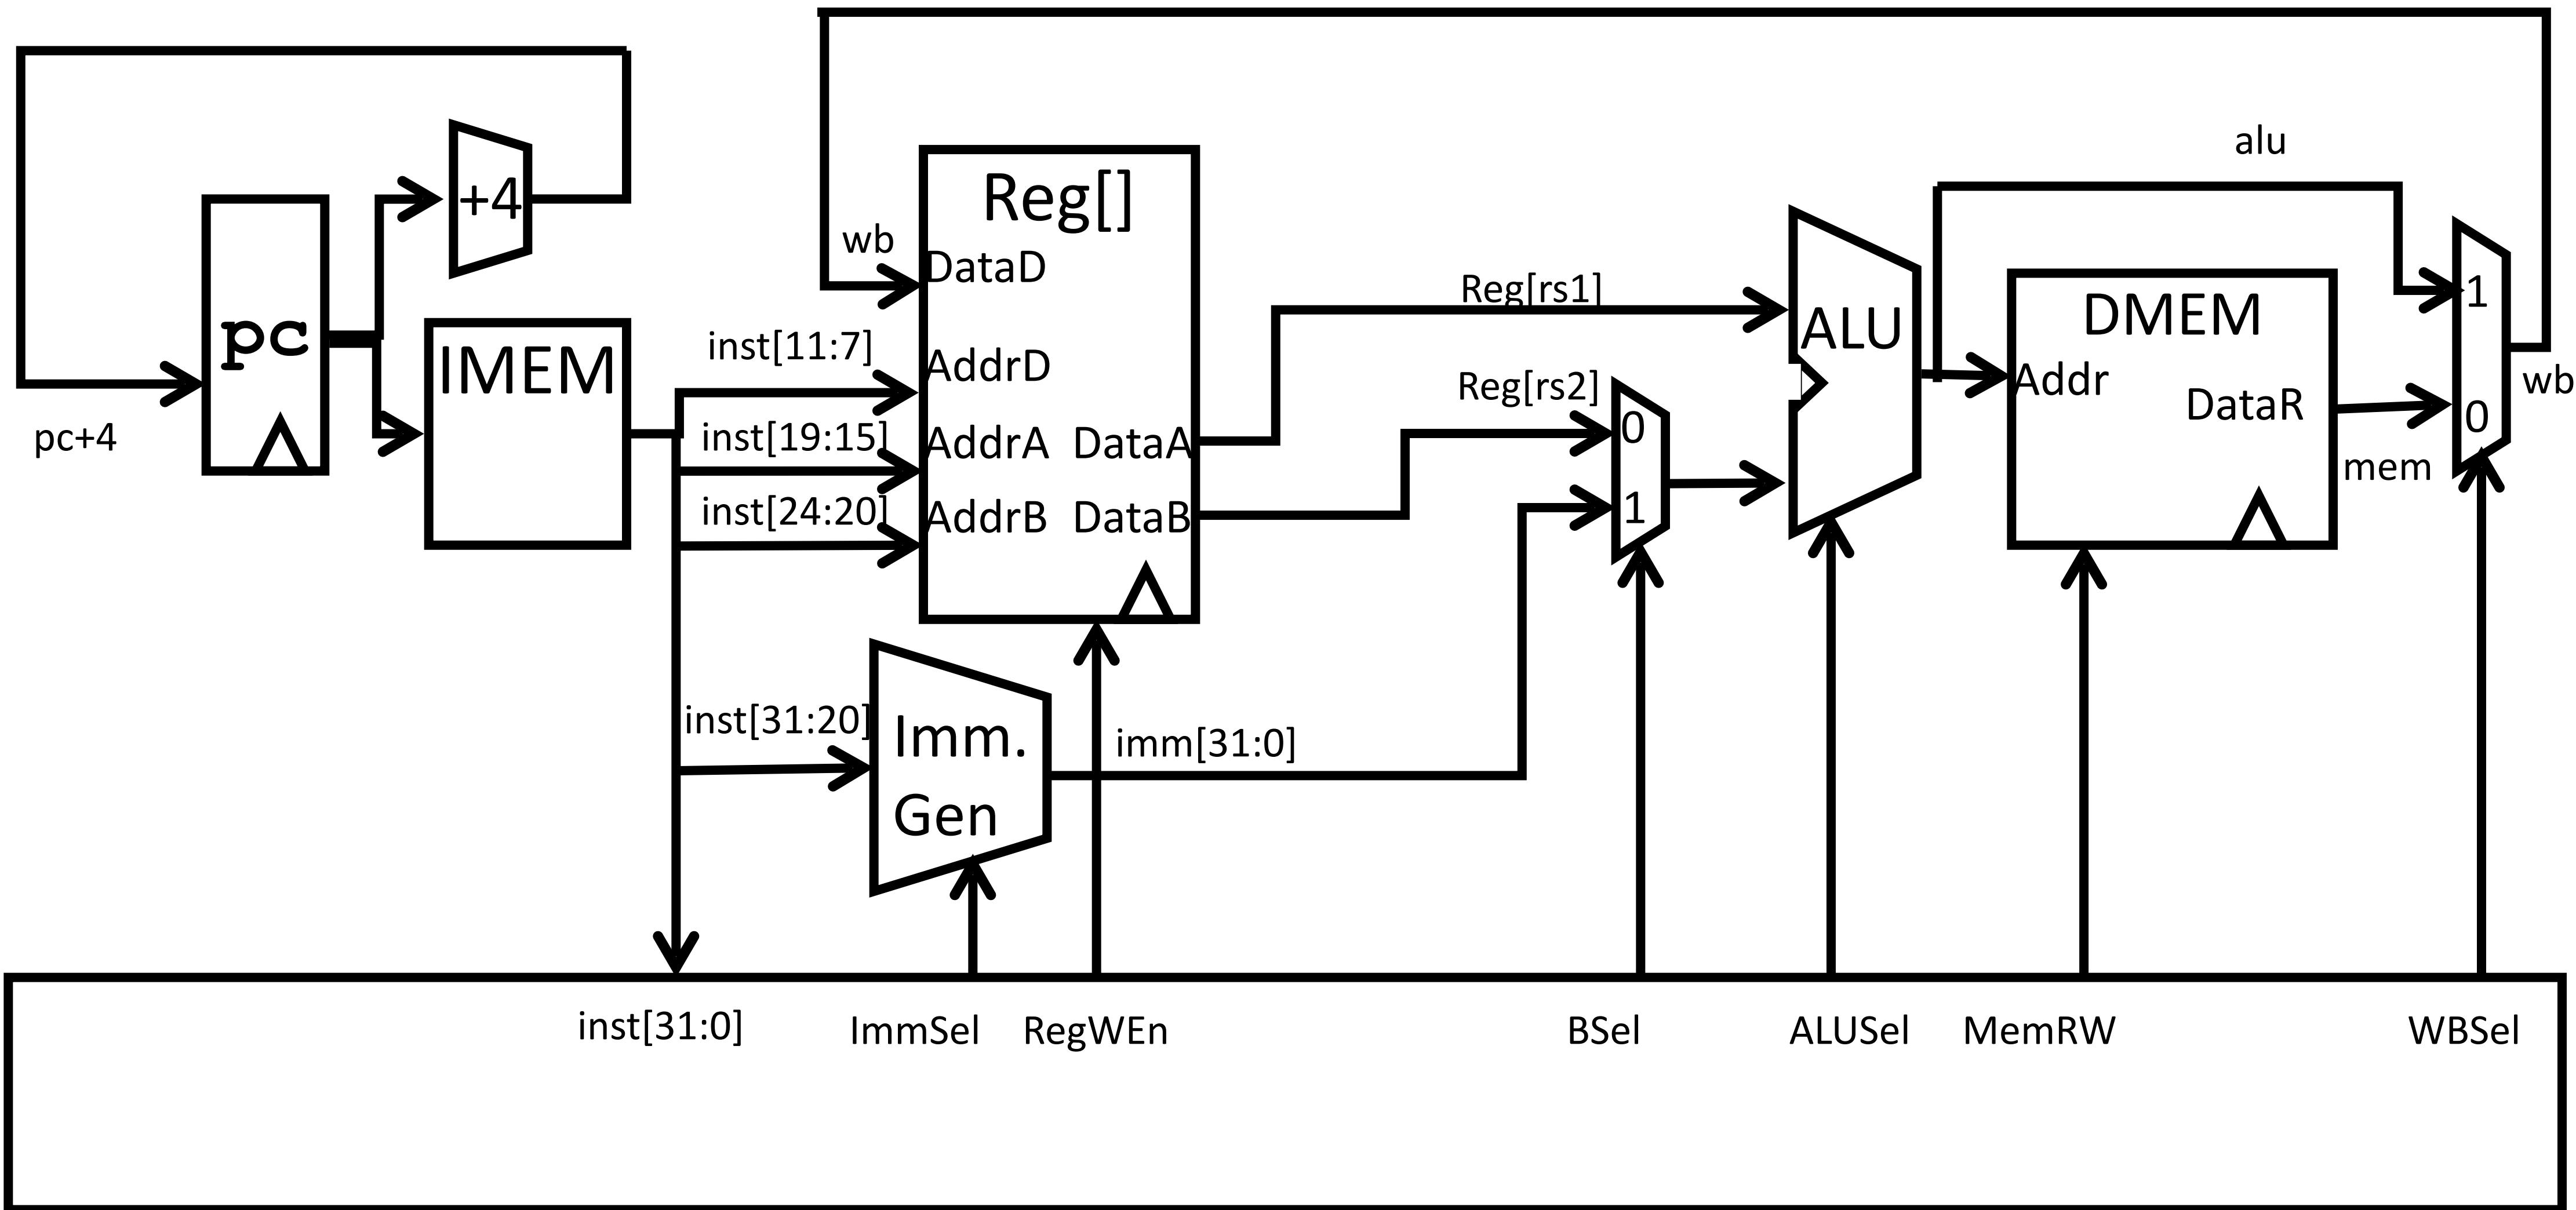
STORE

= 8

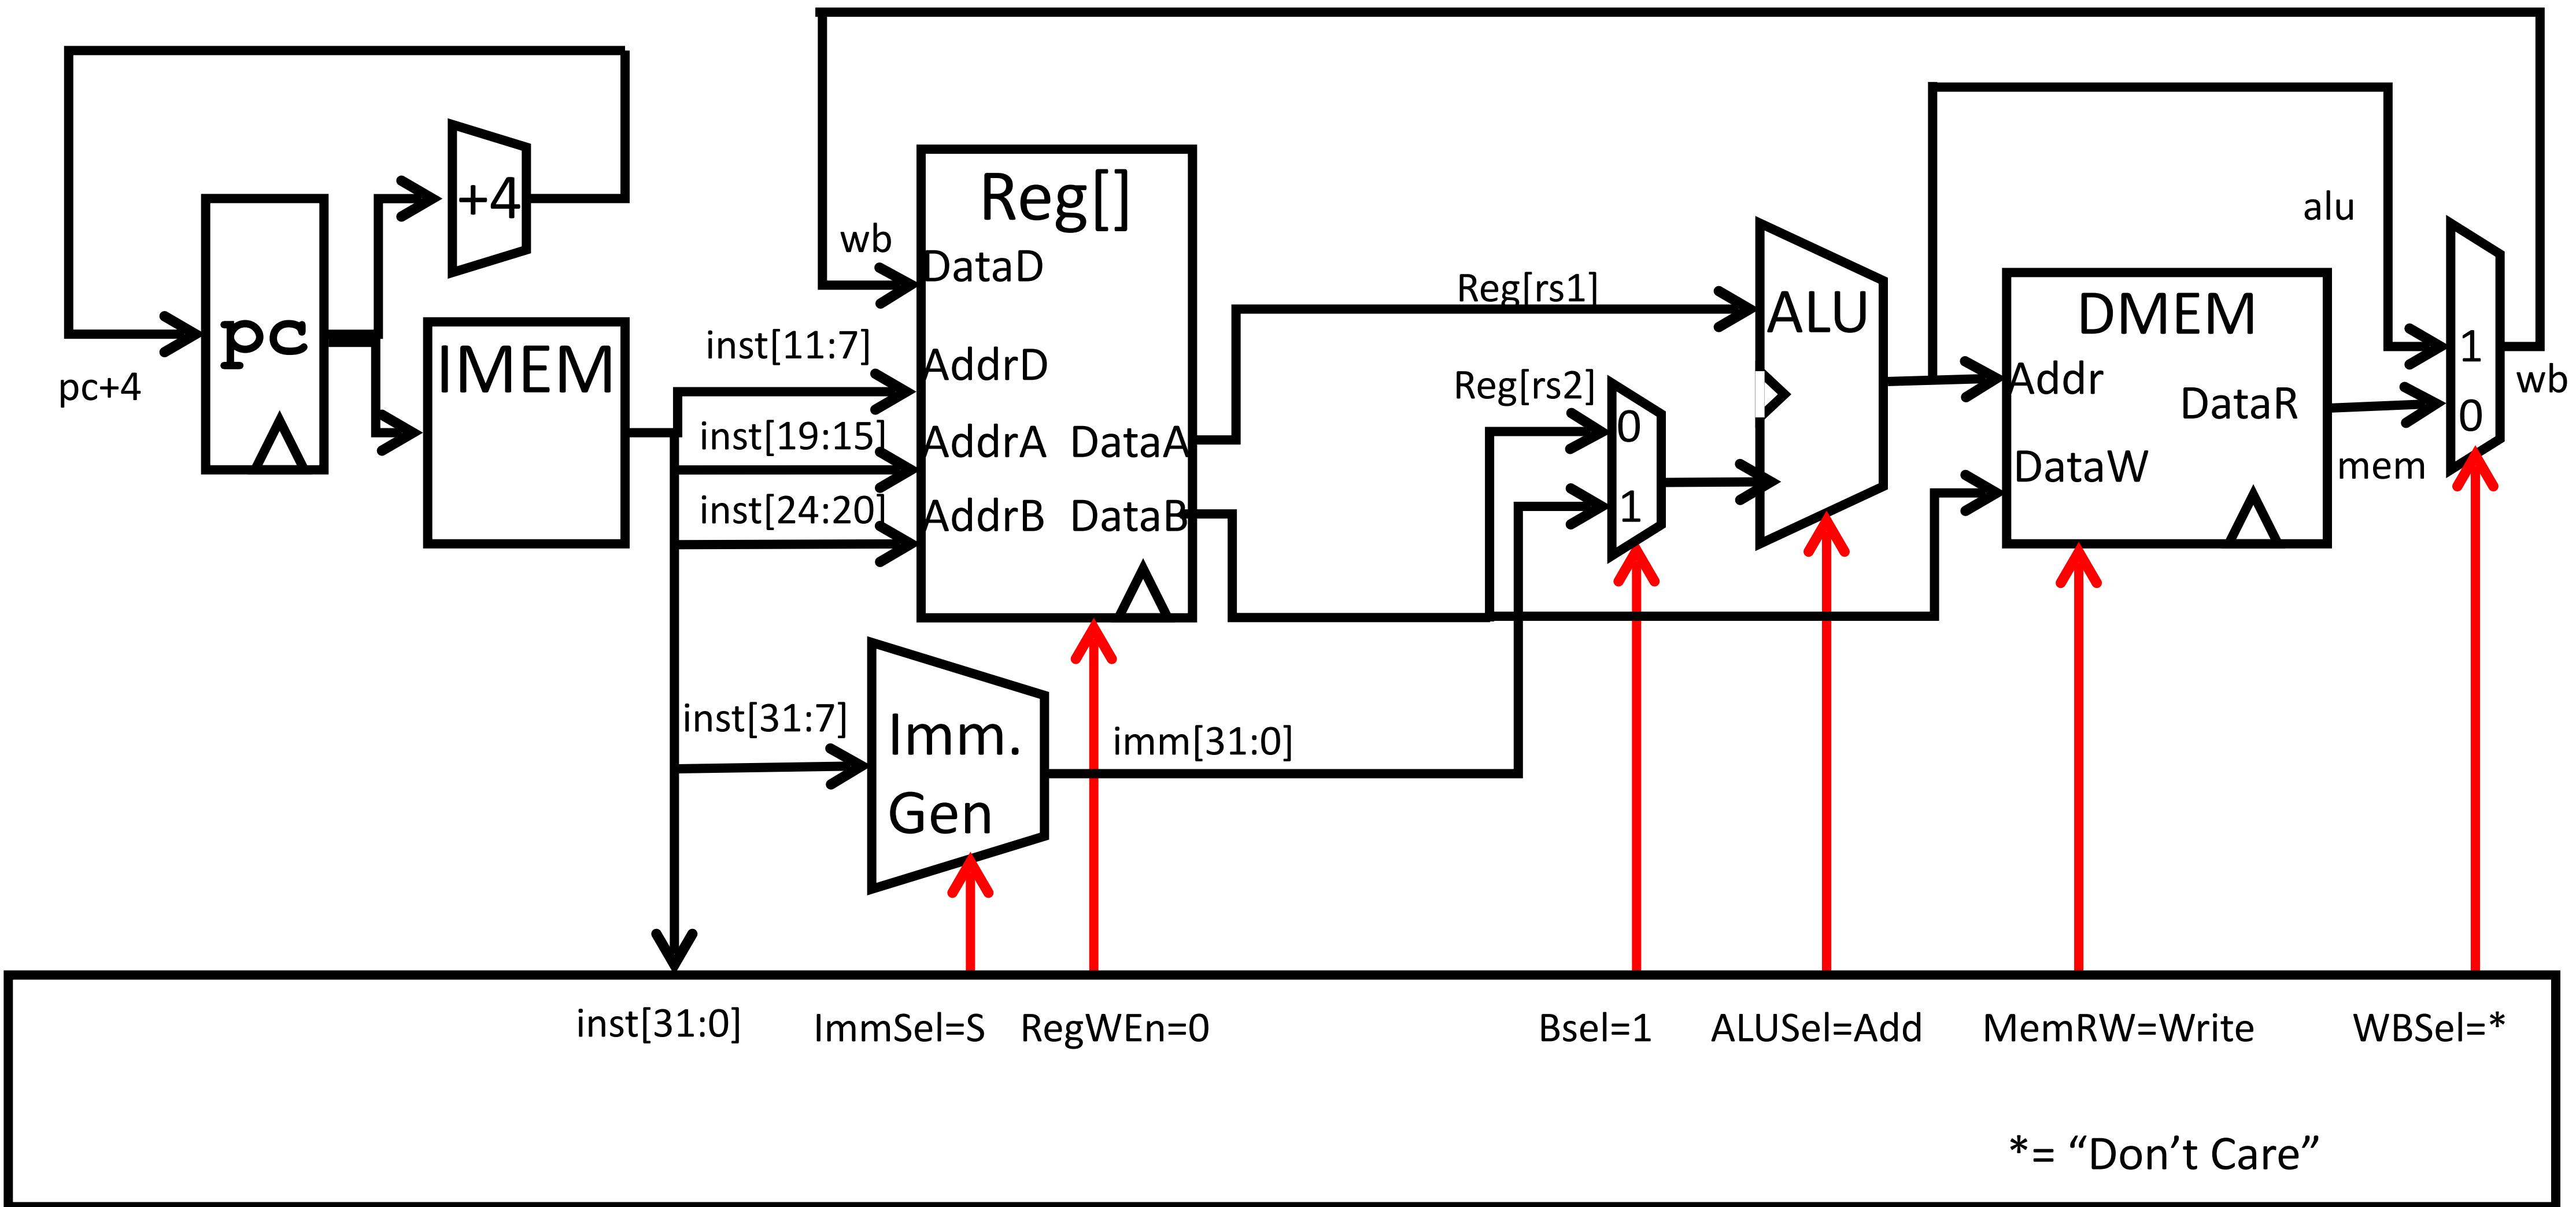


combined 12-bit offset = 8

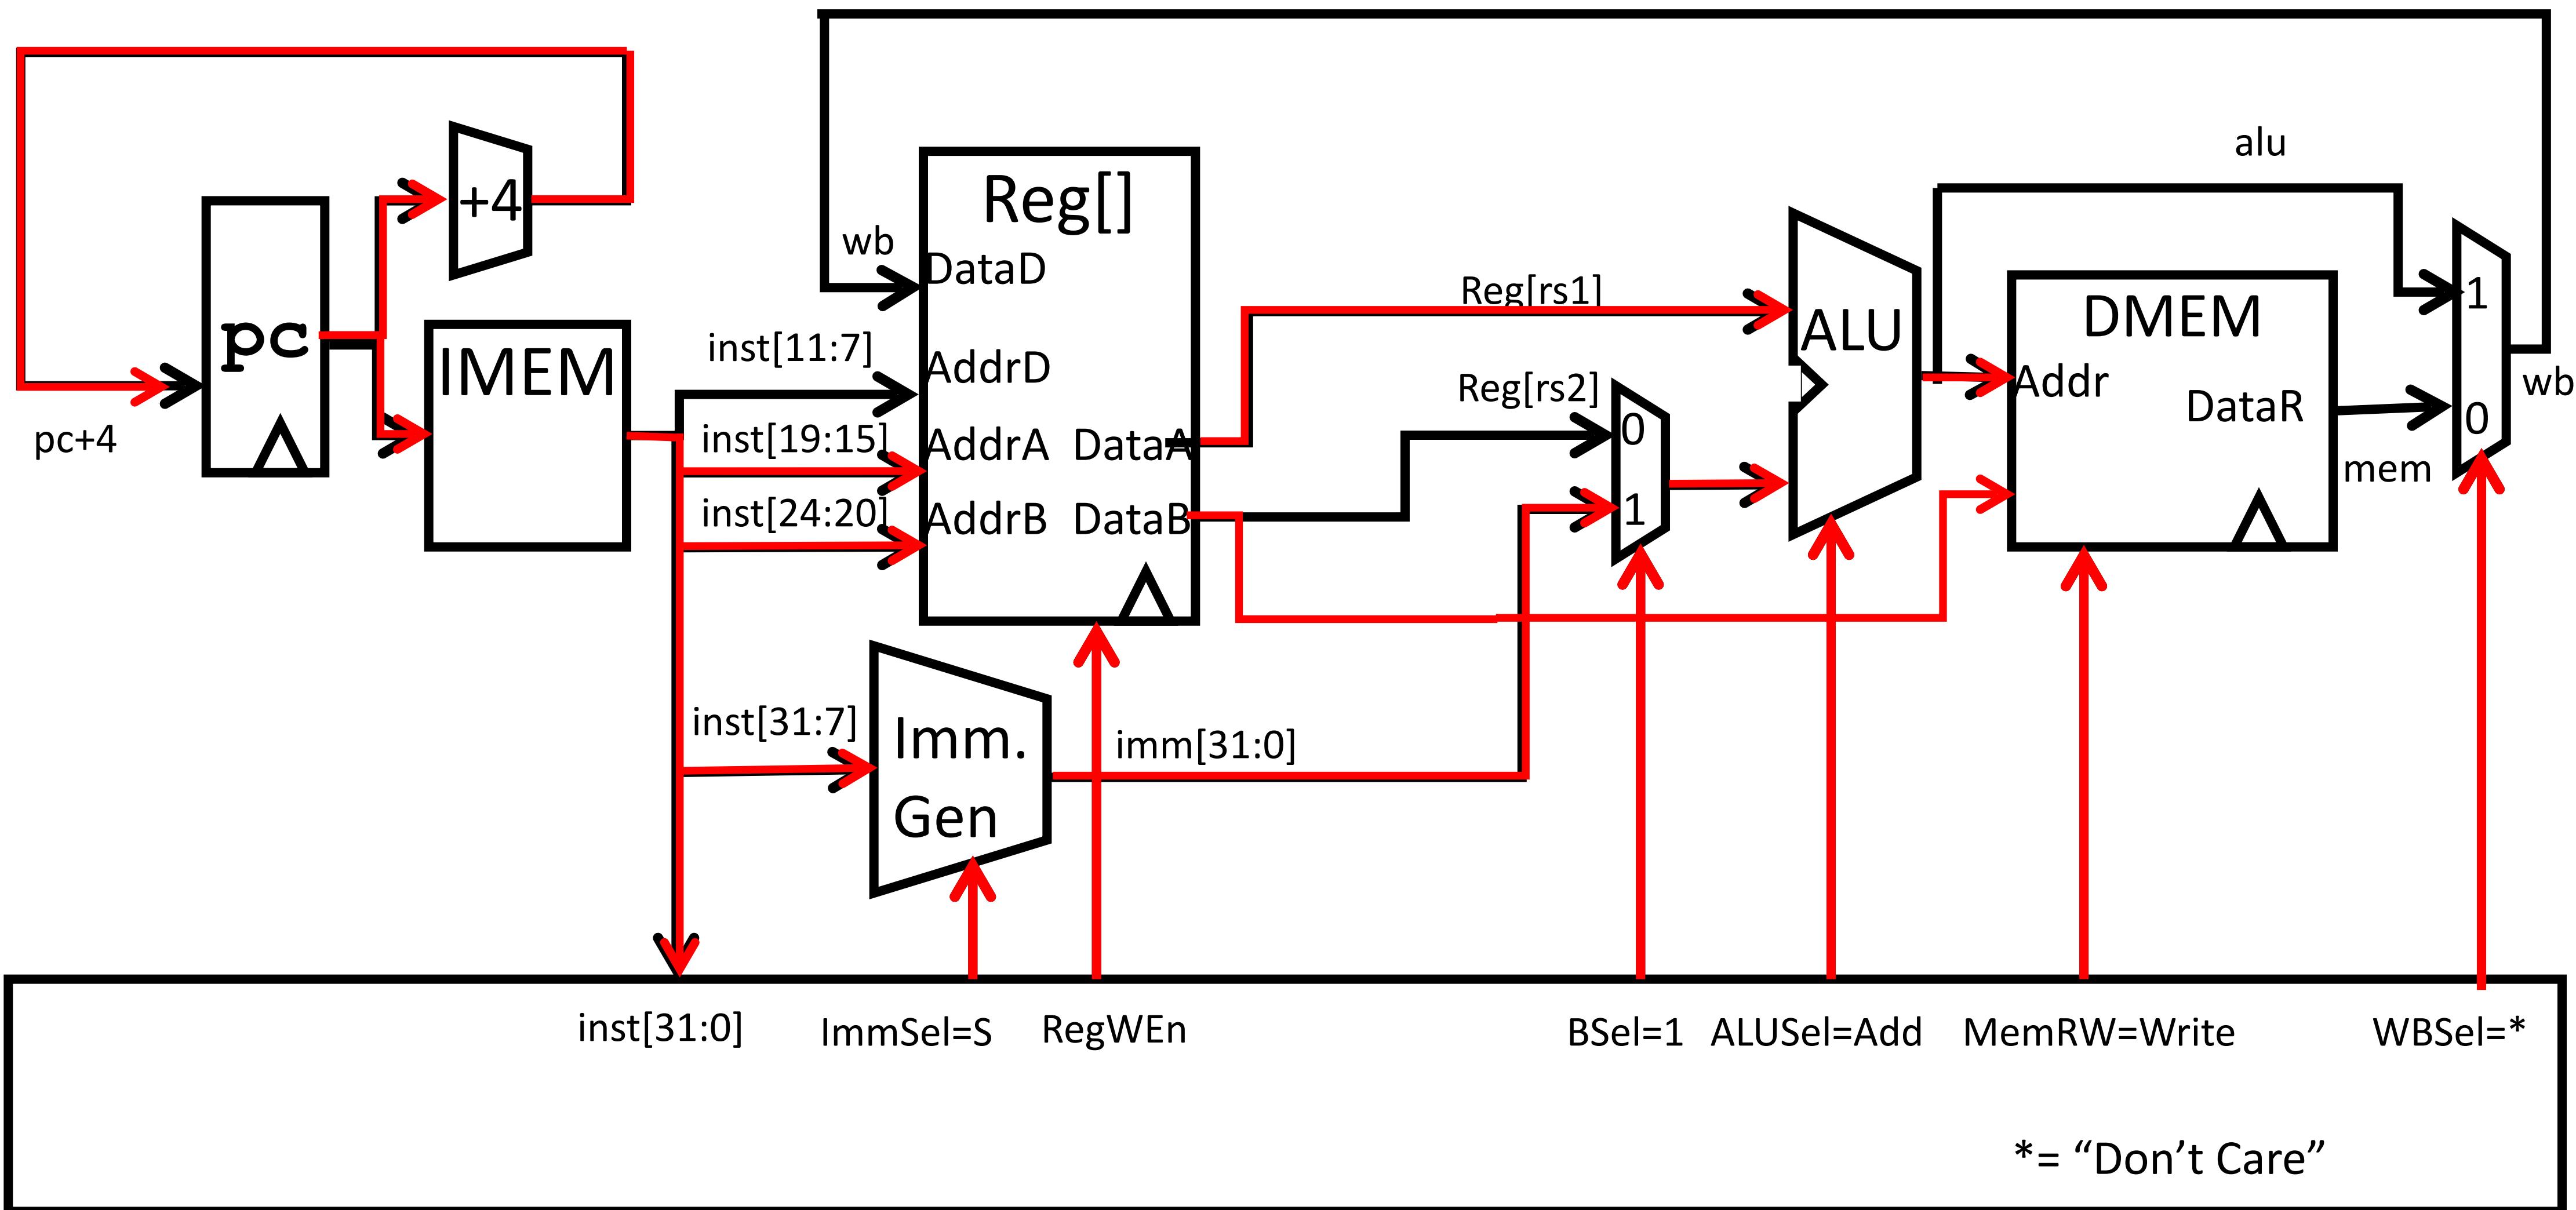
Review: Adding 1w to datapath



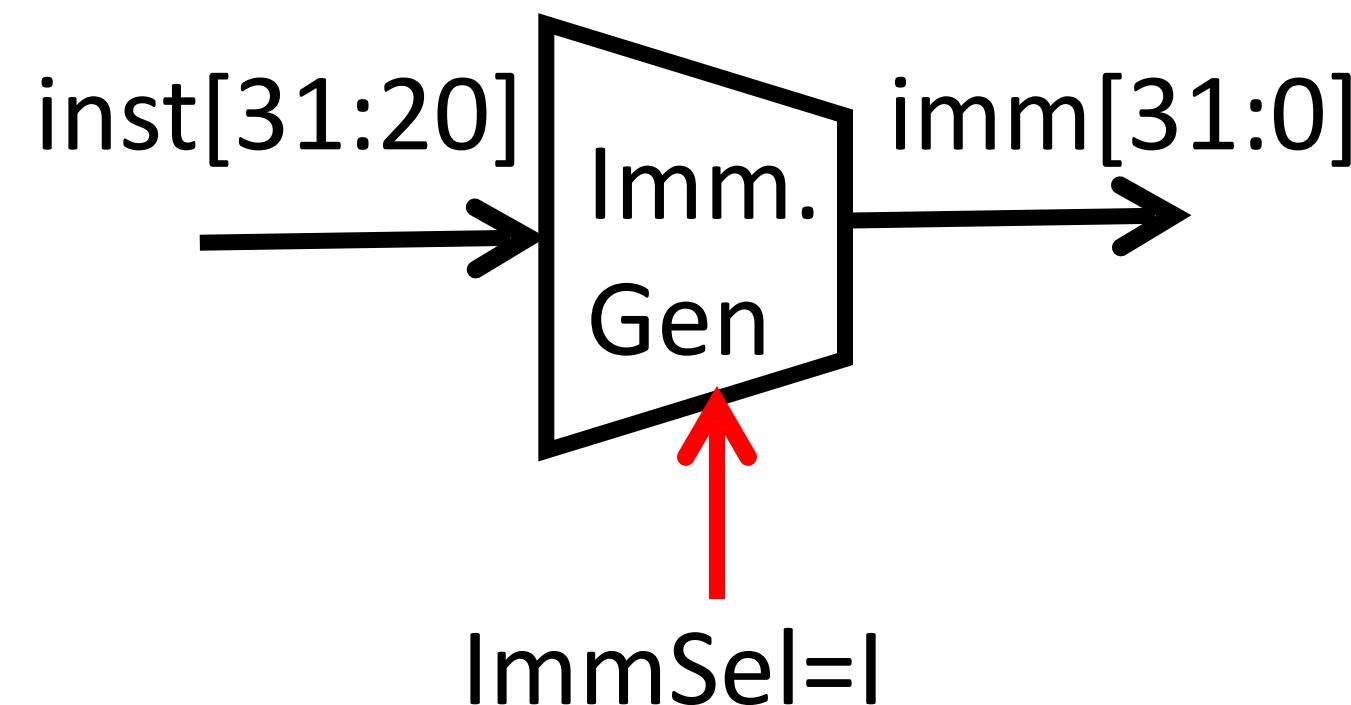
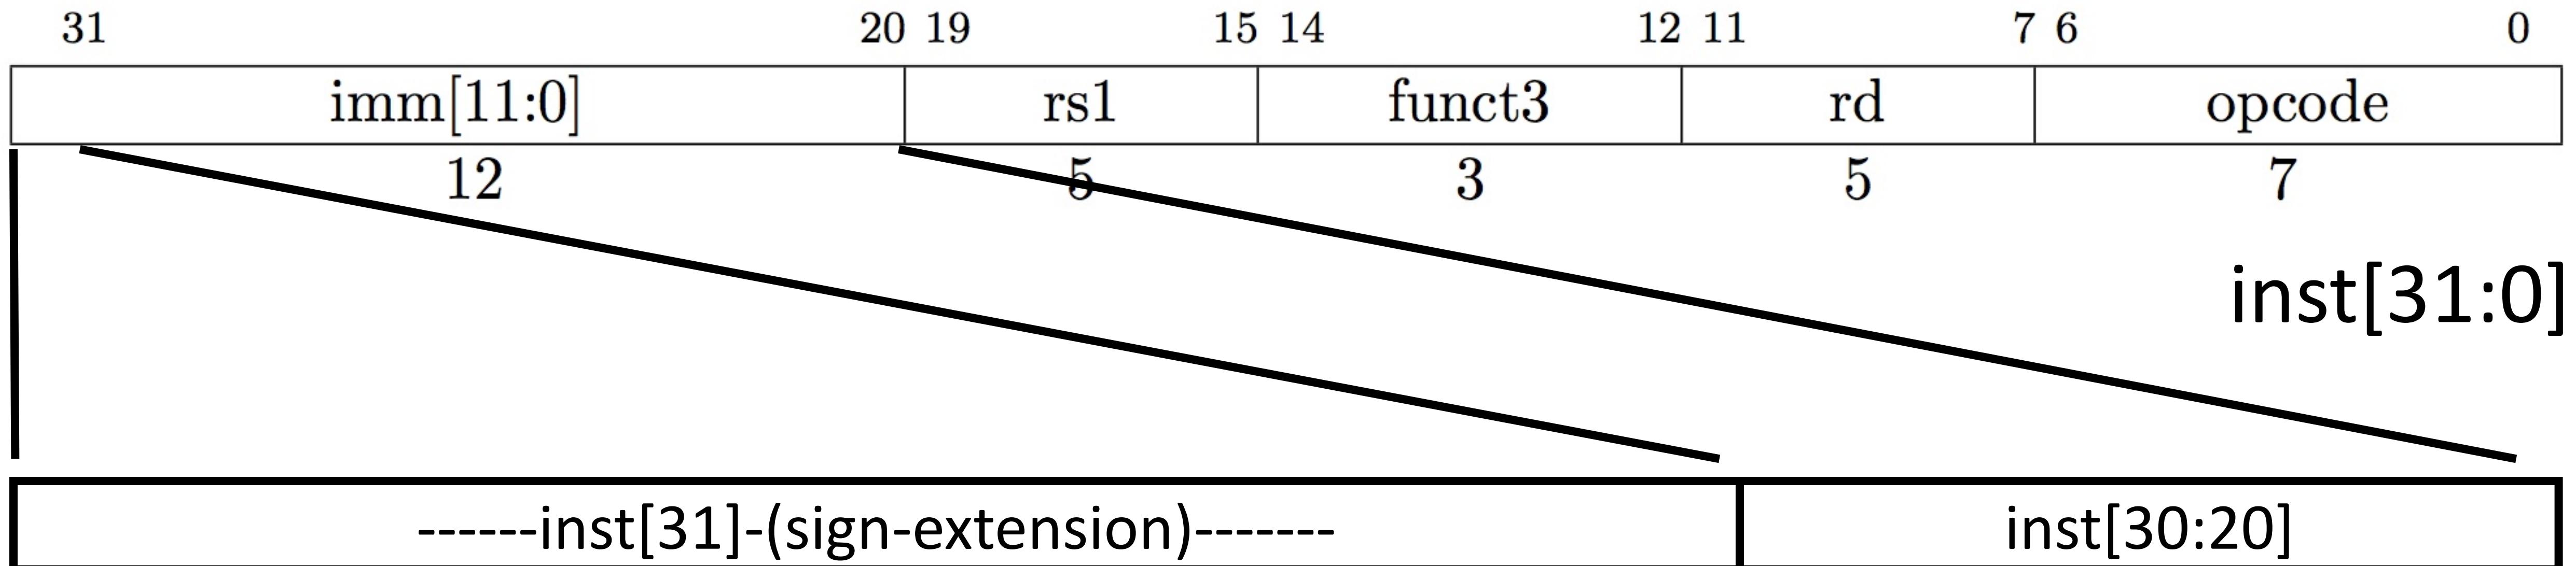
Adding **sw** to datapath



Adding *sw* to datapath



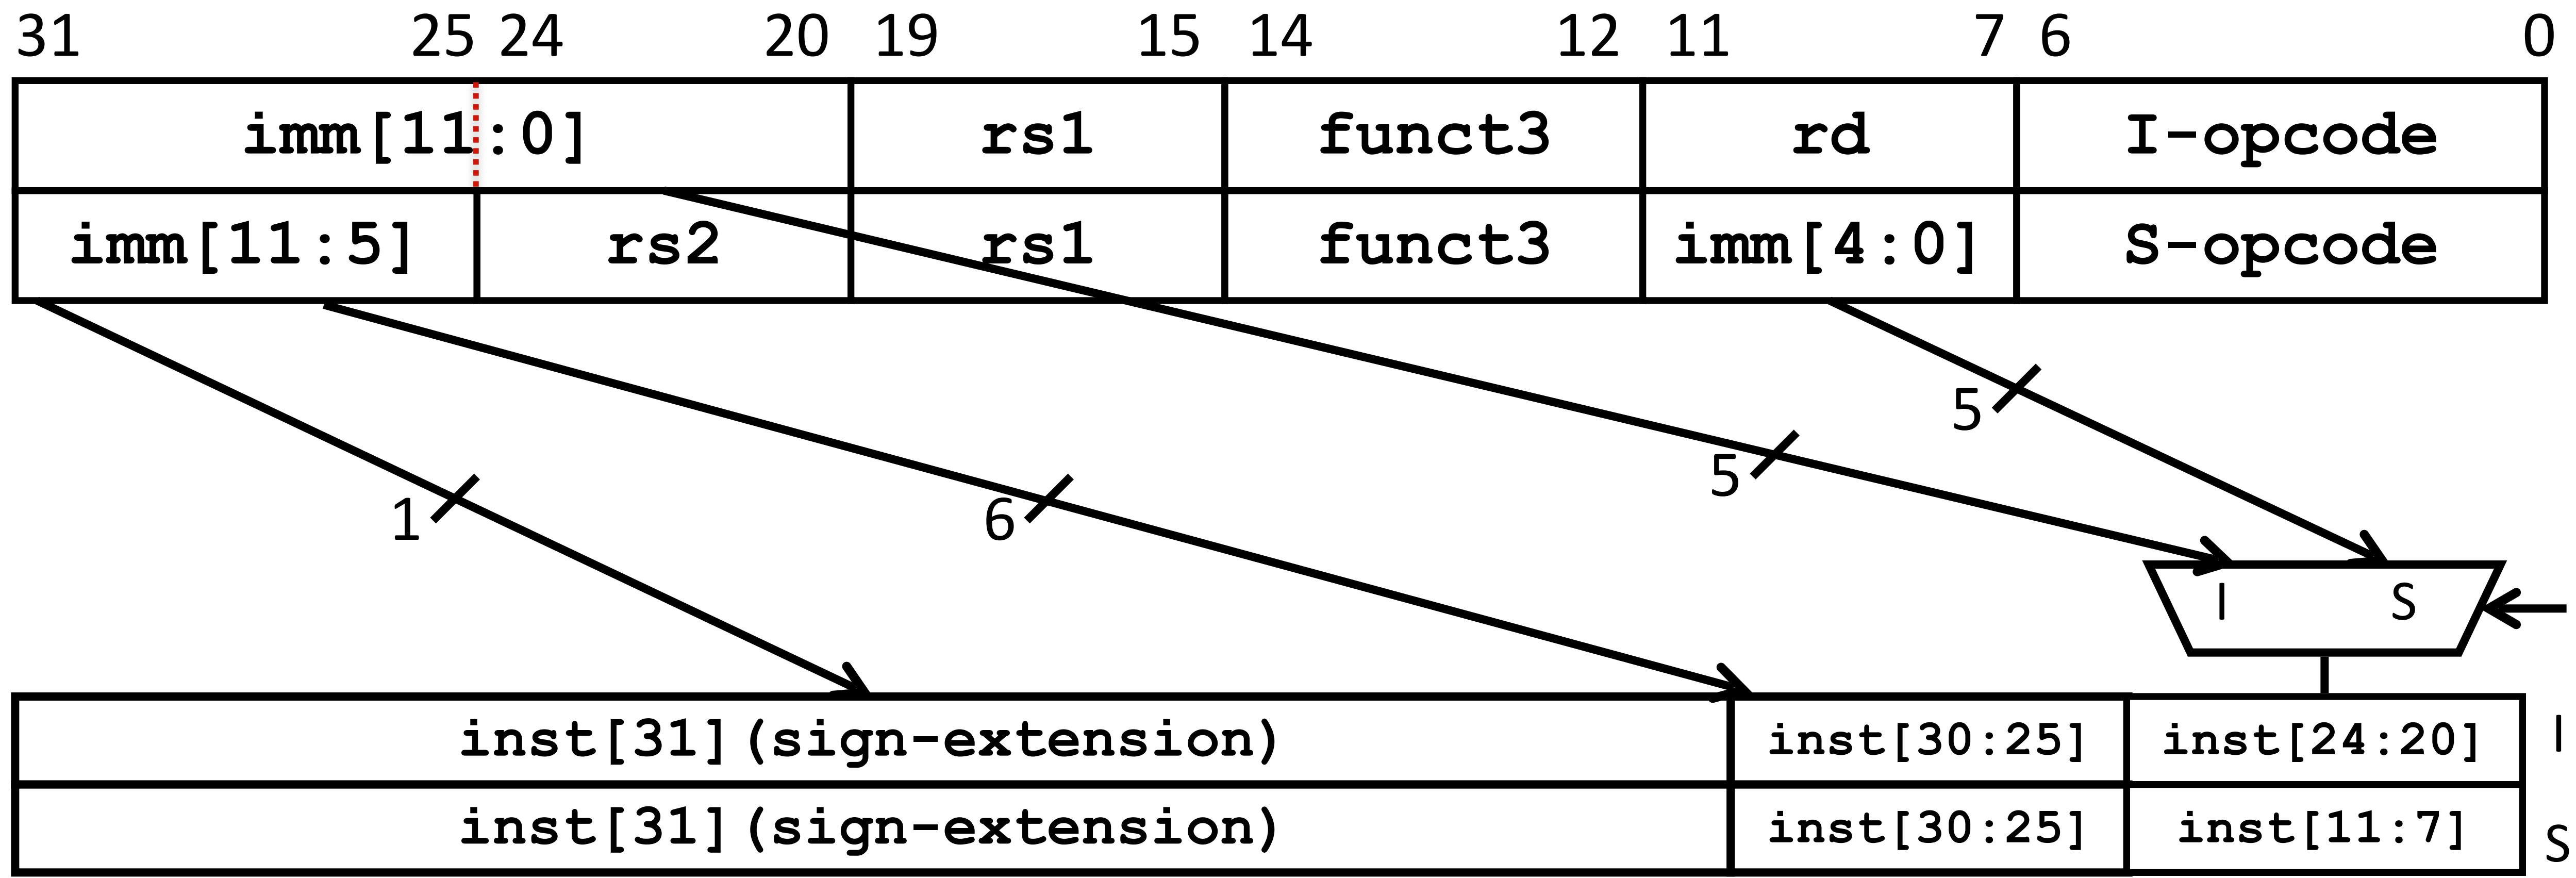
Review: I-Format immediates



- High 12 bits of instruction (inst[31:20]) copied to low 12 bits of immediate (imm[11:0])
- Immediate is sign-extended by copying value of inst[31] to fill the upper 20 bits of the immediate value (imm[31:12])

I & S -type Immediate Generator

inst[31:0]



31

11 10

5 4

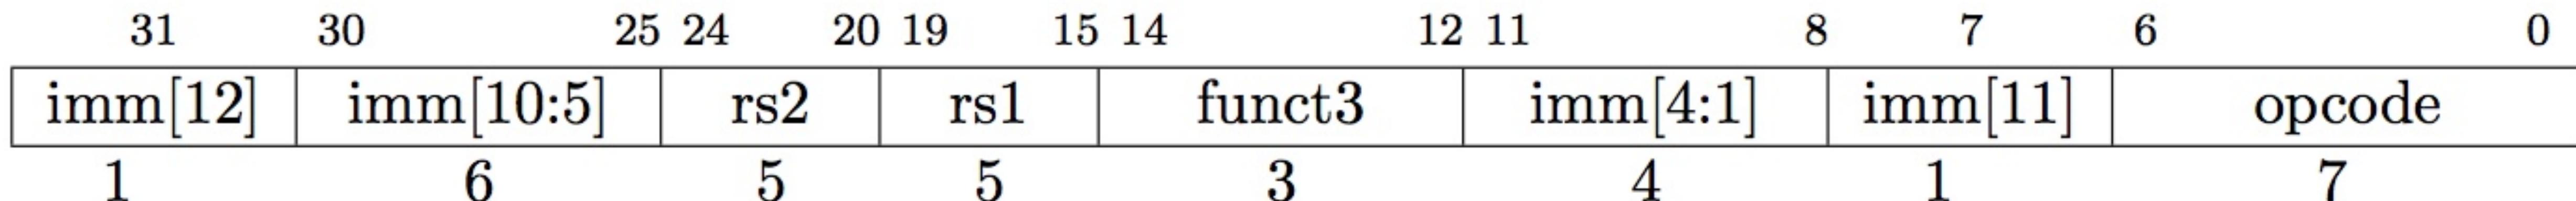
0

imm[31:0]

- Just need a 5-bit mux to select between two positions where low five bits of immediate can reside in instruction
- Other bits in immediate are wired to fixed positions in instruction

Implementing Branches

Uses the “B-type” instruction format



- RISC-V Assembly Instruction, example:

beq rs1, rs2, label

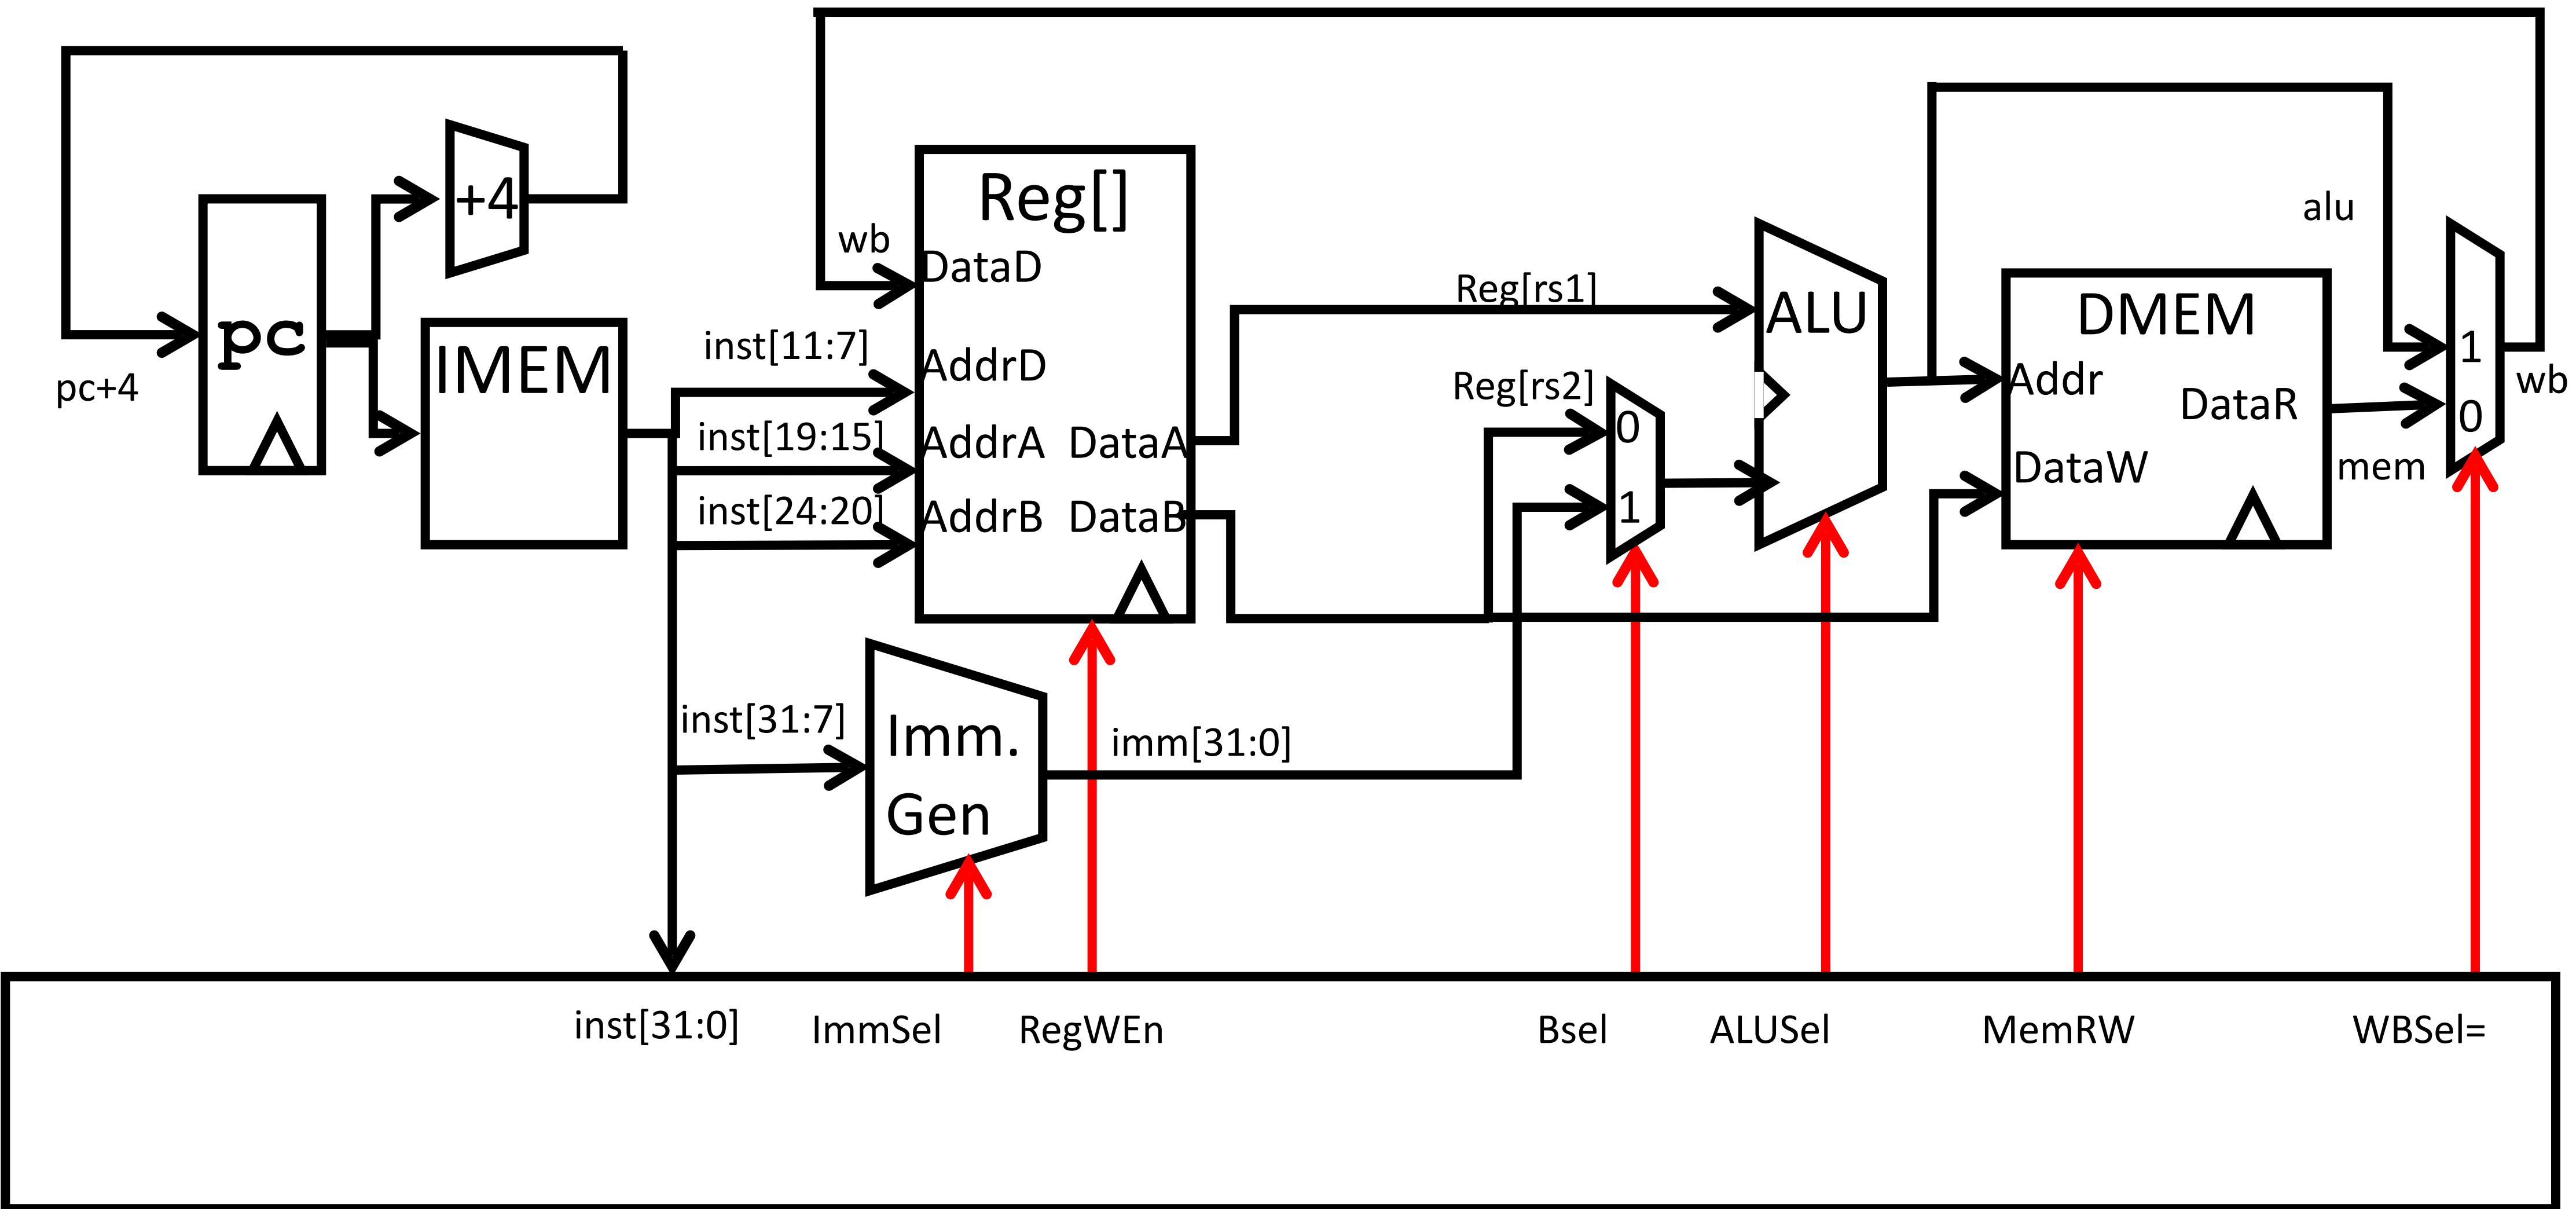
if $rs1 == rs2$ $pc \leftarrow pc + offset$ // offset computed by compiler/assembler and stored in the immediate field(s)

example:

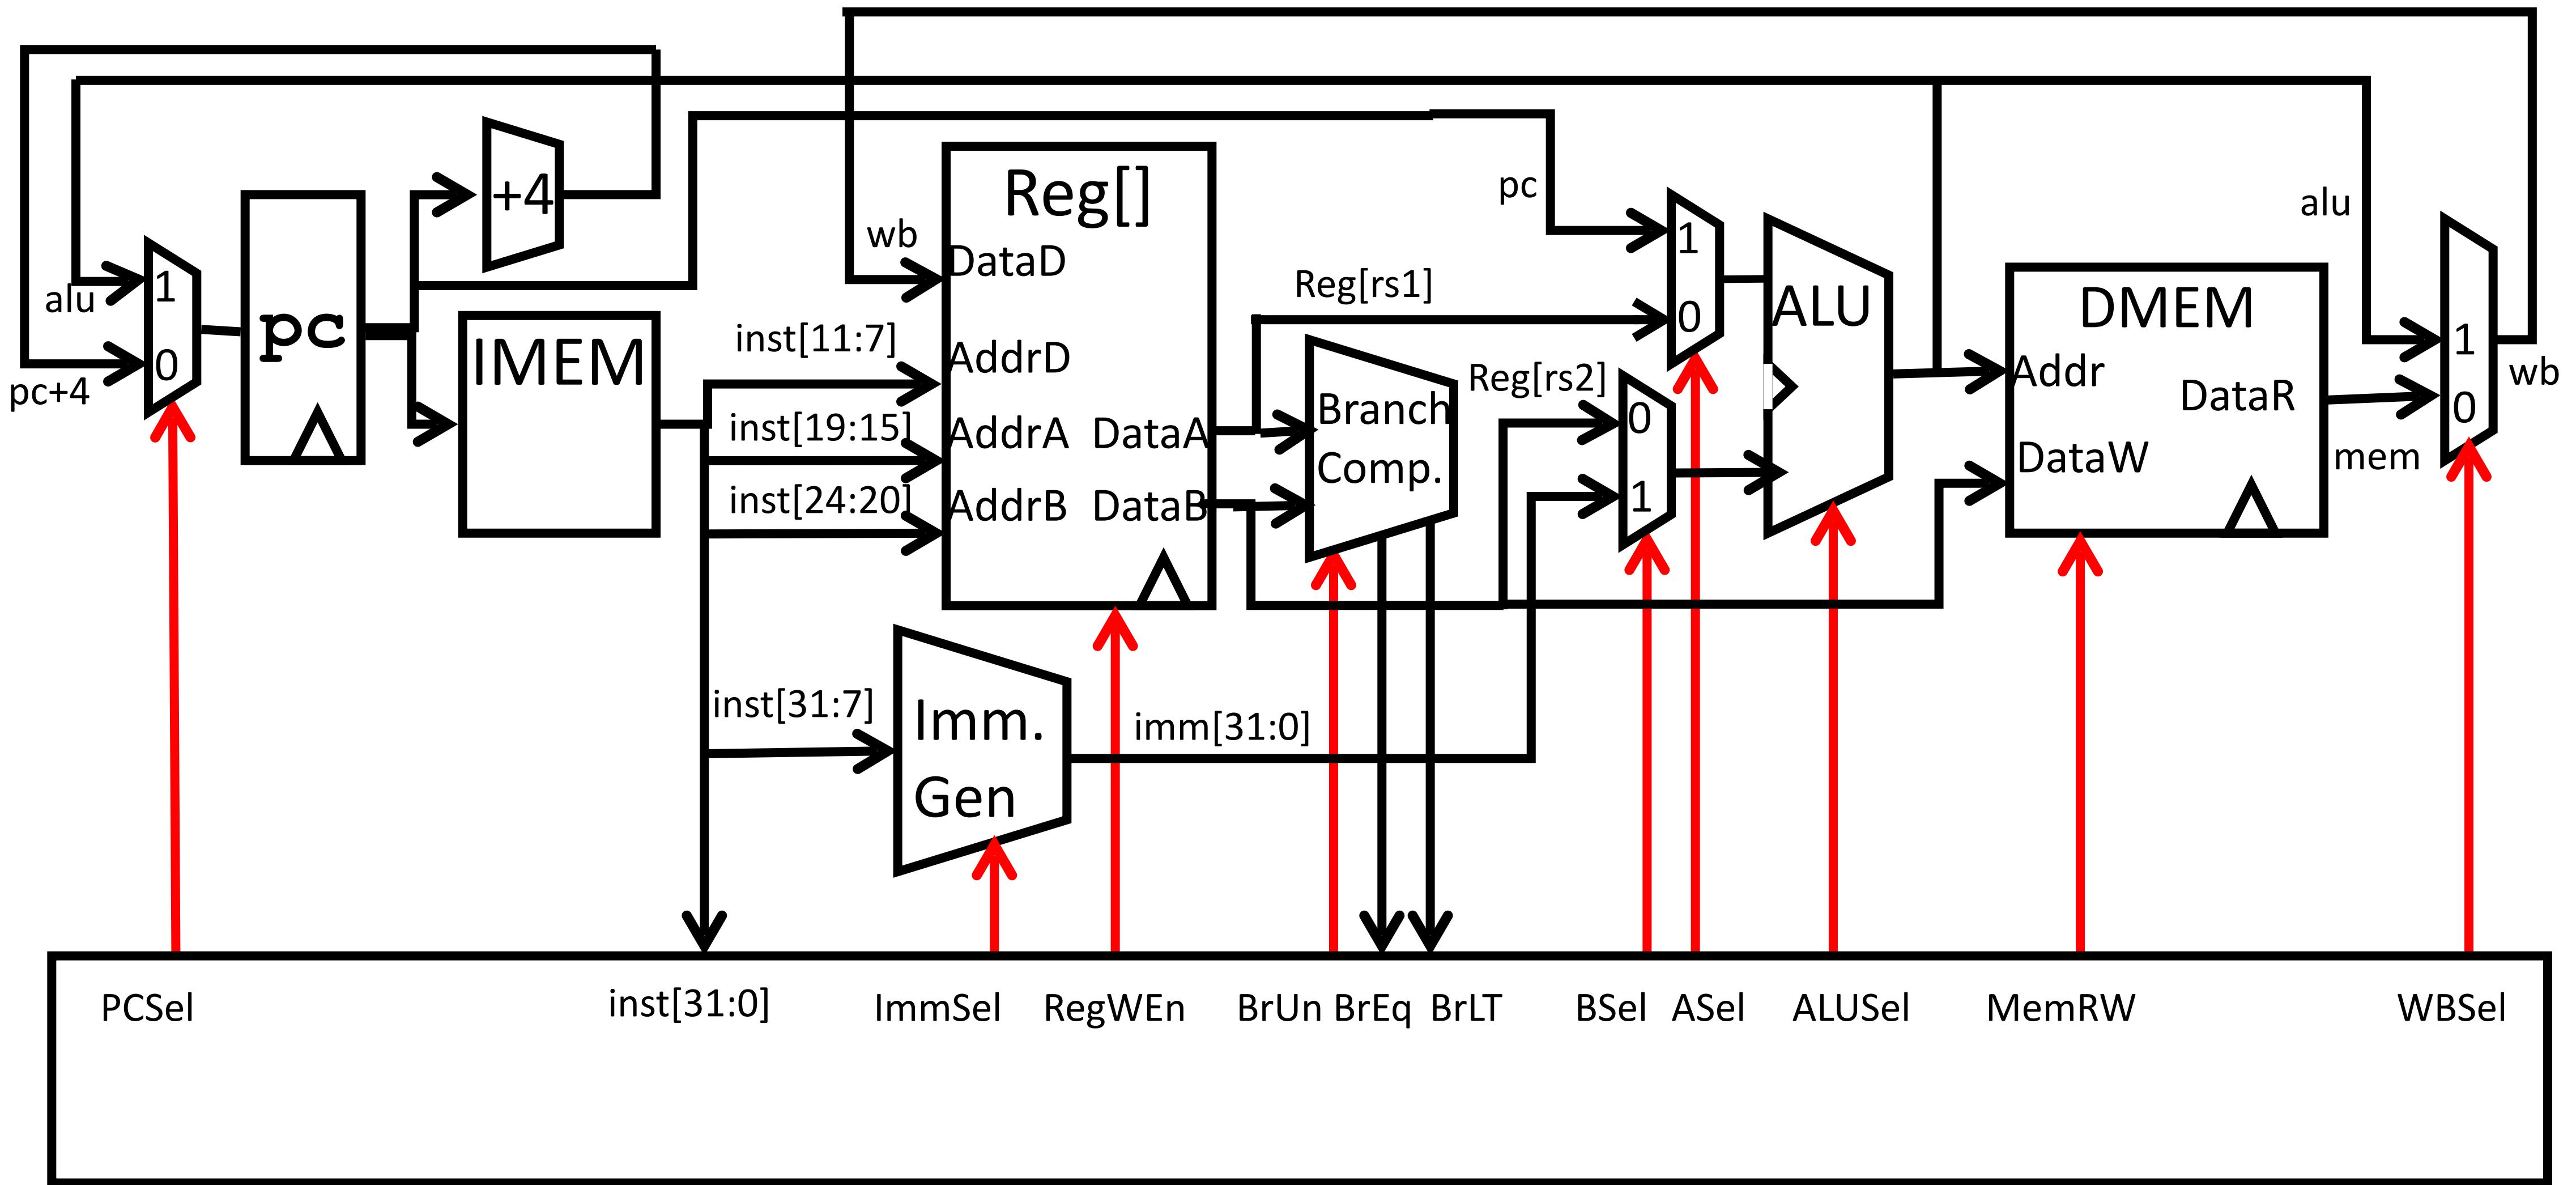
beq x1, x2, L1

- B-format is mostly same as S-Format, with two register sources (rs1/rs2) and a 12-bit immediate
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode *even* 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)

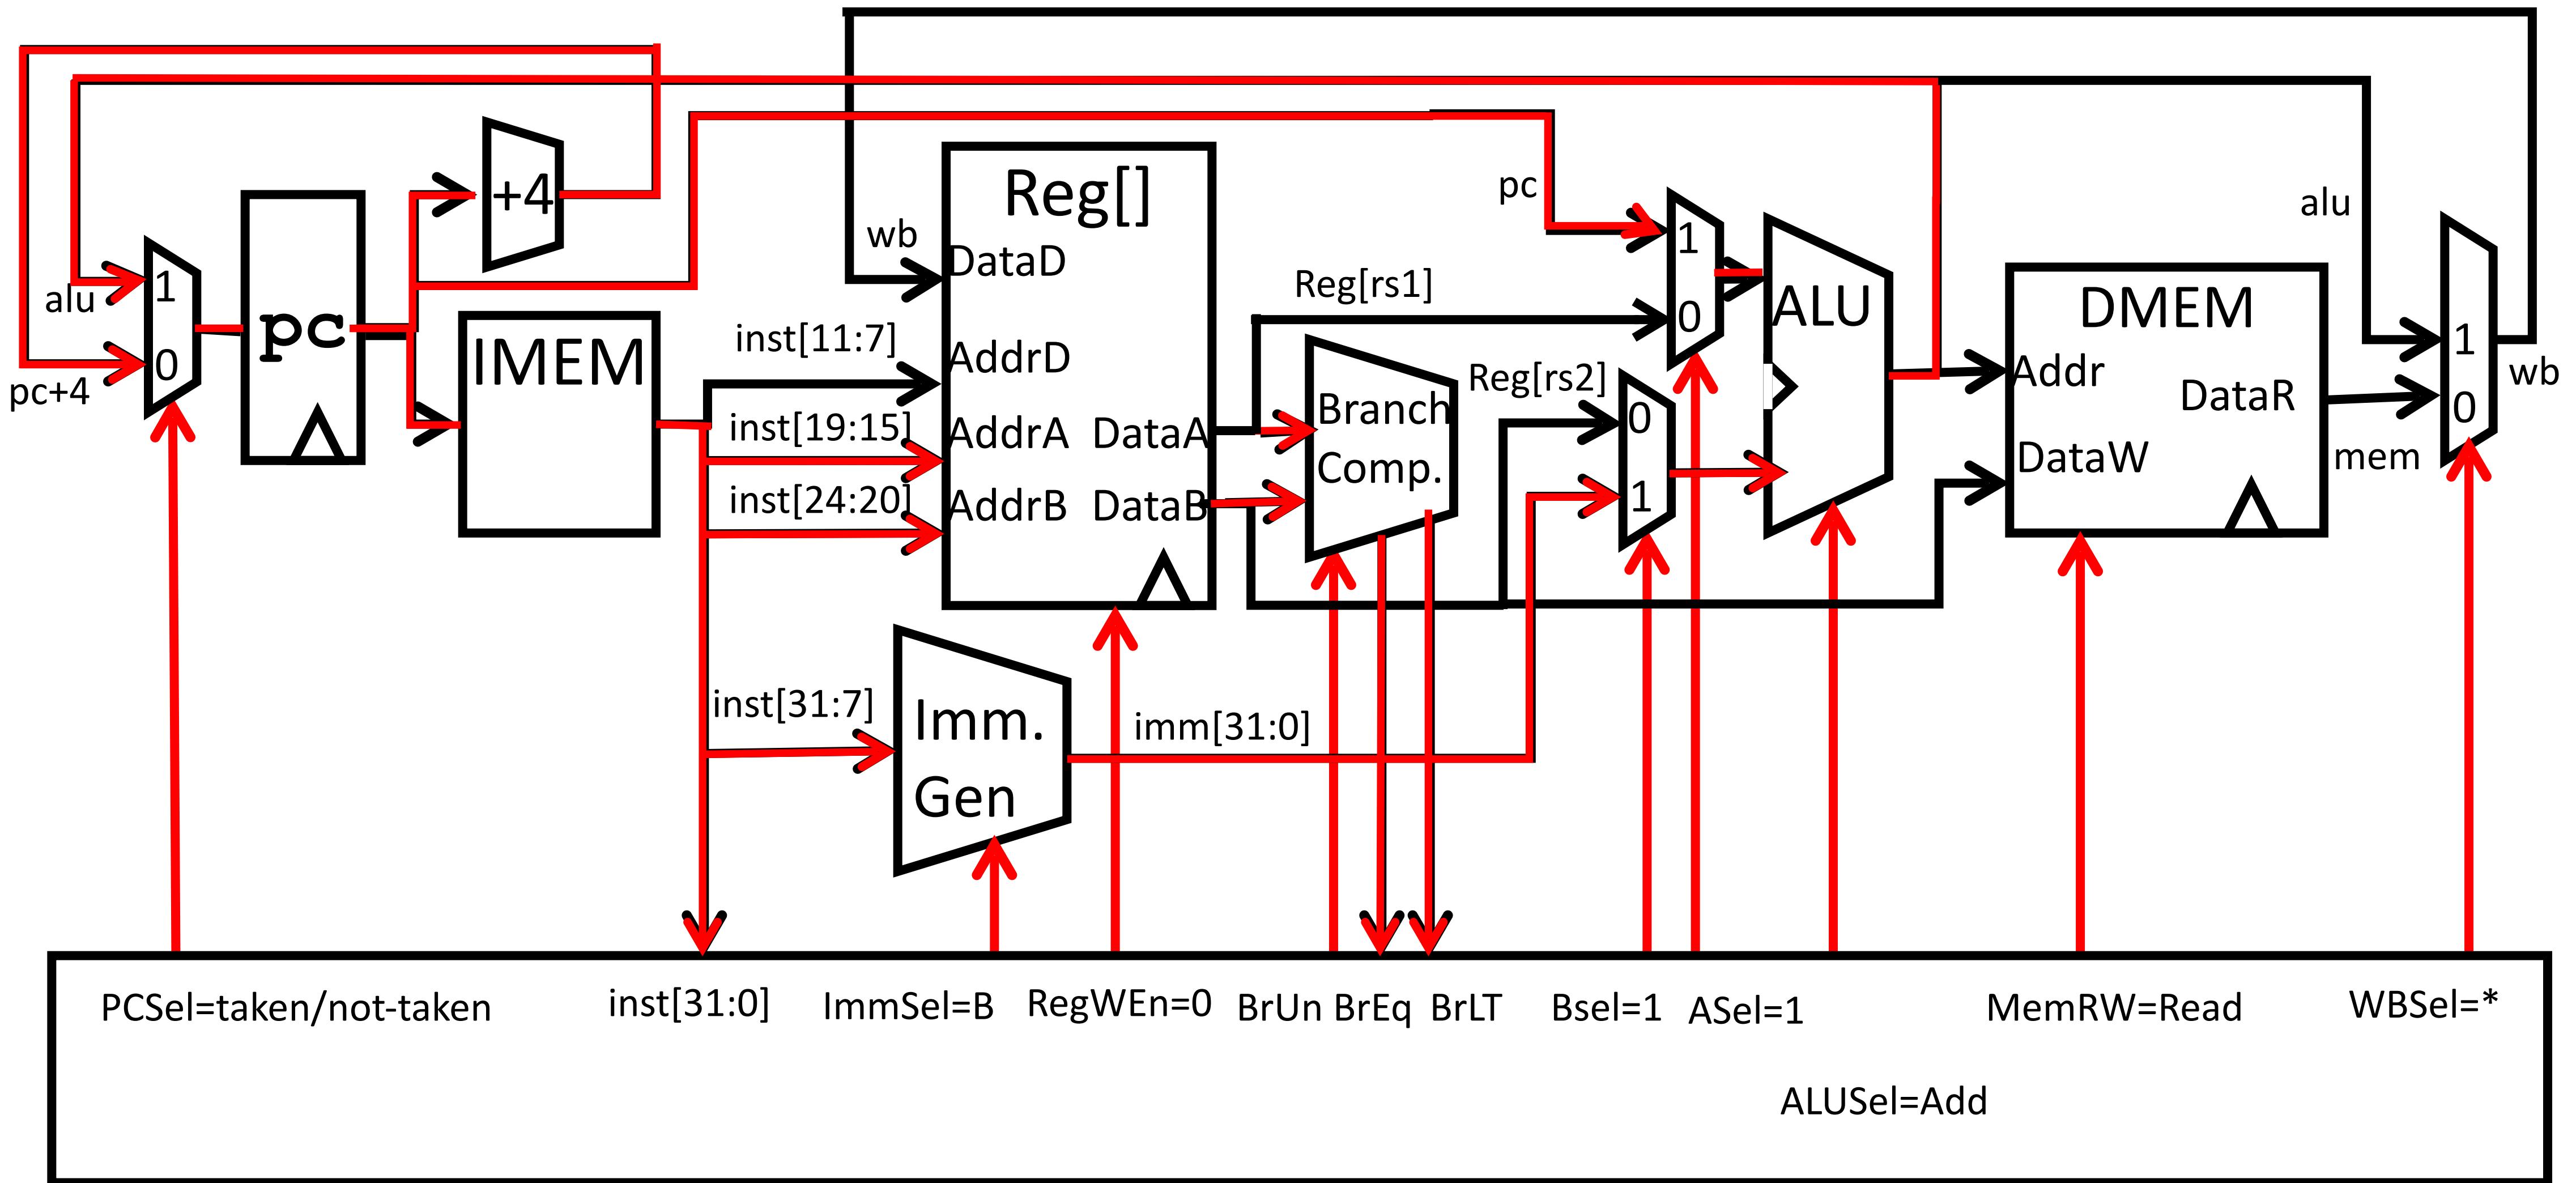
Review: Adding sw to datapath



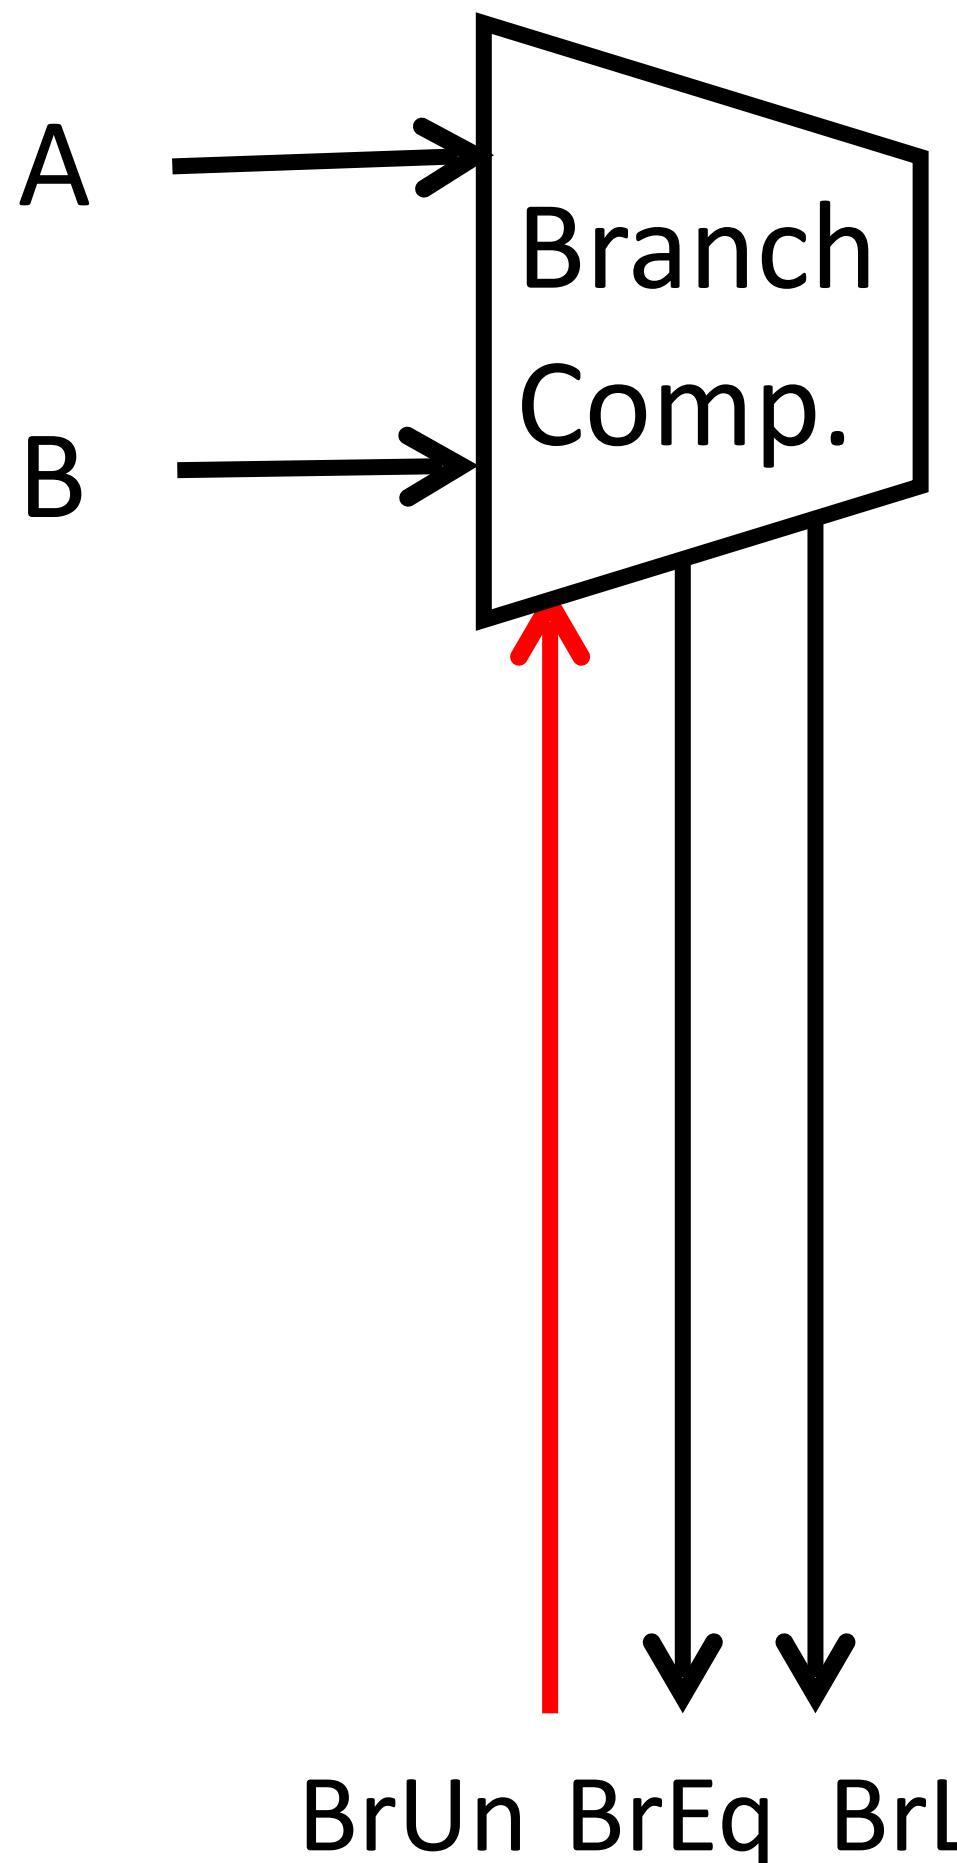
Adding branches to datapath



Adding branches to datapath



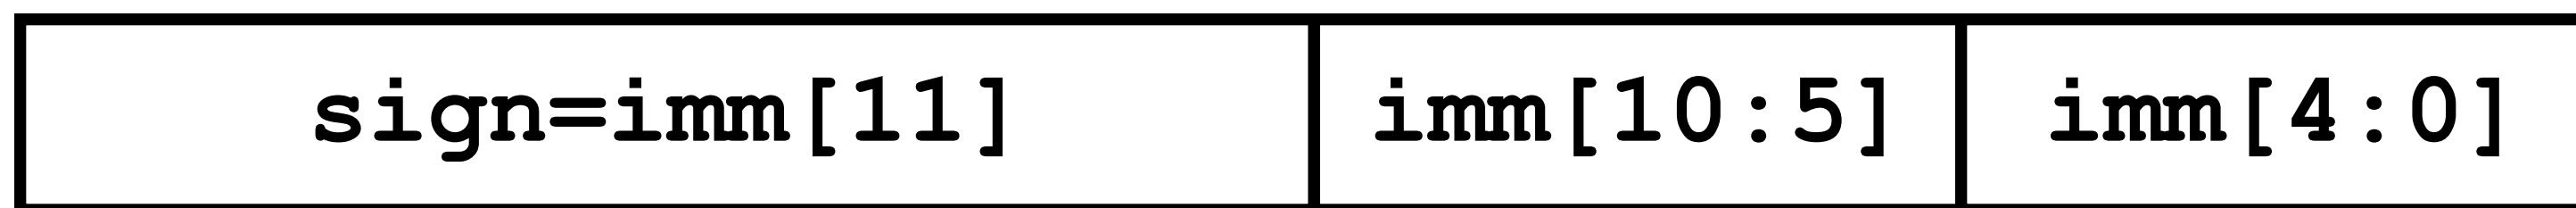
Branch Comparator



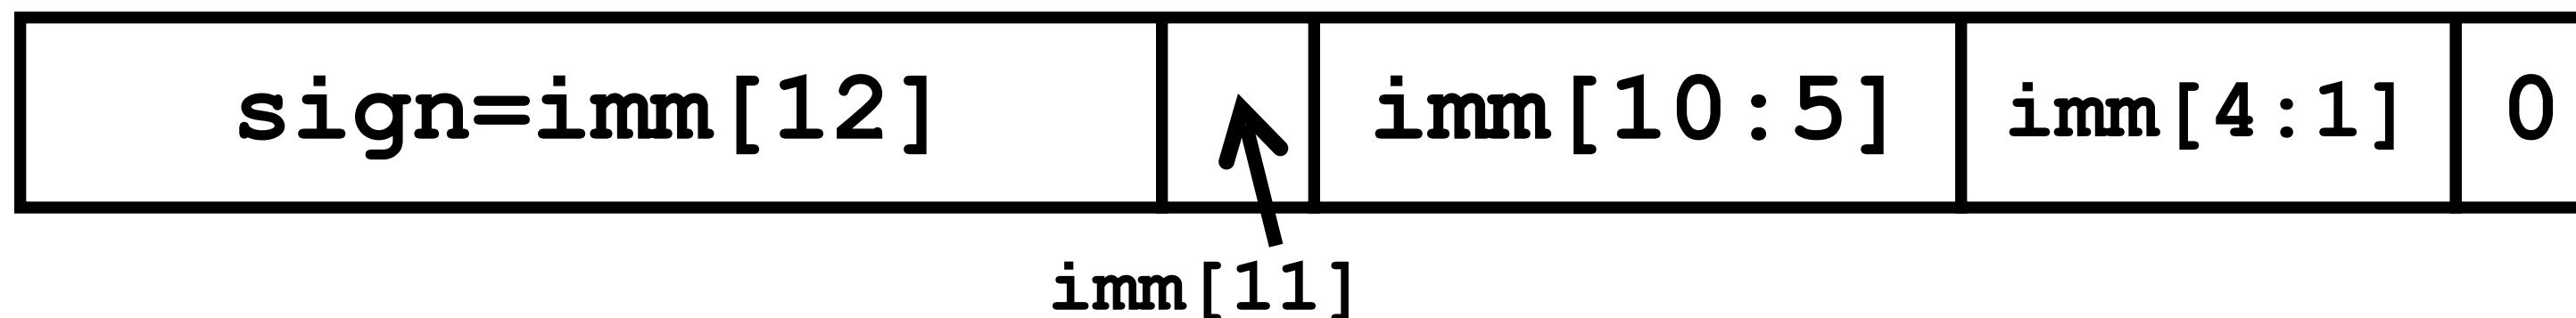
- $\text{BrEq} = 1$, if $A=B$
- $\text{BrLT} = 1$, if $A < B$
- $\text{BrUn} = 1$ selects unsigned comparison for BrLT , $0=\text{signed}$
- BGE branch: $A \geq B$, if $!(A < B)$

RISC-V Branch Immediates

- 12-bit immediate encodes PC-relative offset of -4096 to +4094 bytes in multiples of 2 bytes
- RISC-V approach: keep 11 immediate bits in fixed position in output value, and rotate LSB of S-format to be bit 12 of B-format



S-Immediate



B-Immediate (shift left by 1)

Only one bit changes position between S and B, so only need a single-bit 2-way mux

RISC-V Immediate Encoding

Instruction Encodings, $\text{inst}[31:0]$

31	30	25 24	21	20	19	15 14	12 11	8	7	6	0	
												opcode
												R-type
												I-type
												S-type
												B-type

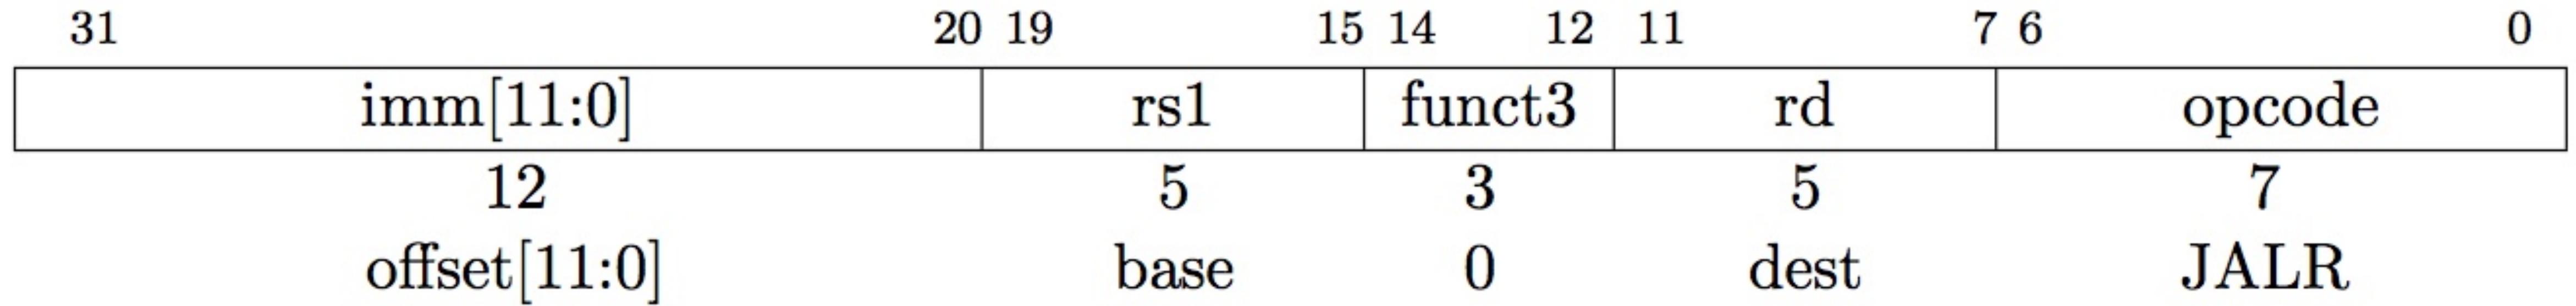
32-bit immediates produced, $\text{imm}[31:0]$

31	30	20 19	12	11	10	5	4	1	0	
		— $\text{inst}[31]$ —				$\text{inst}[30:25]$	$\text{inst}[24:21]$	$\text{inst}[20]$		I-immediate
		— $\text{inst}[31]$ —				$\text{inst}[30:25]$	$\text{inst}[11:8]$	$\text{inst}[7]$		S-immediate
		— $\text{inst}[31]$ —		$\text{inst}[7]$	$\text{inst}[30:25]$	$\text{inst}[11:8]$	0			B-immediate

Upper bits sign-extended from $\text{inst}[31]$ always

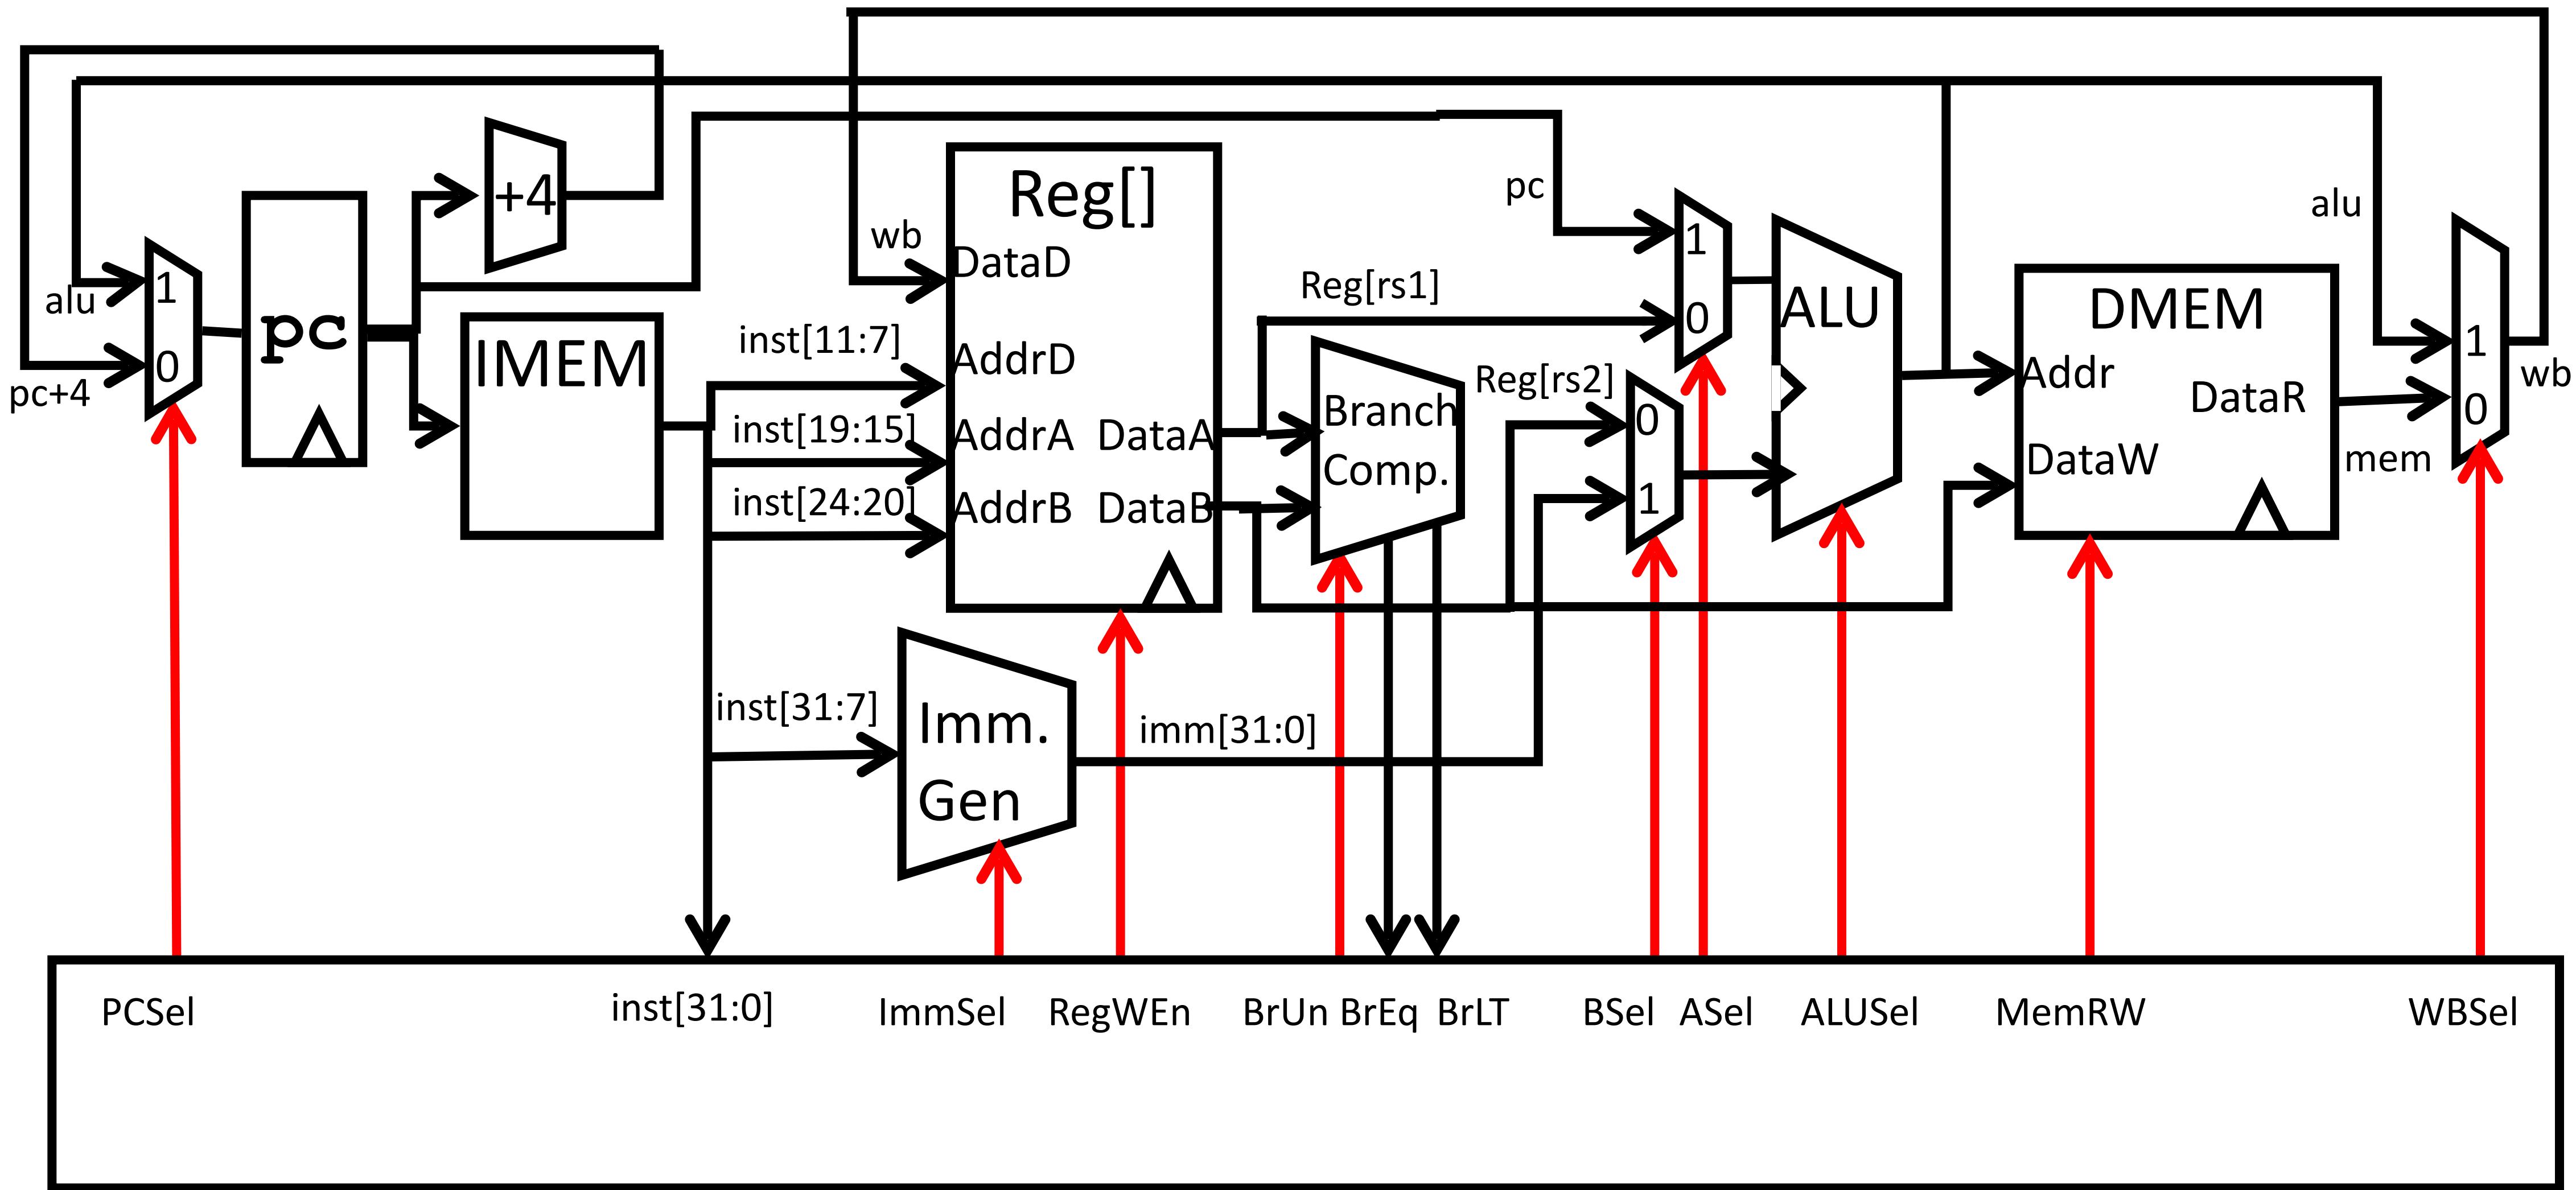
Only bit 7 of instruction changes role in immediate between S and B

Implementing JALR Instruction (I-Format)

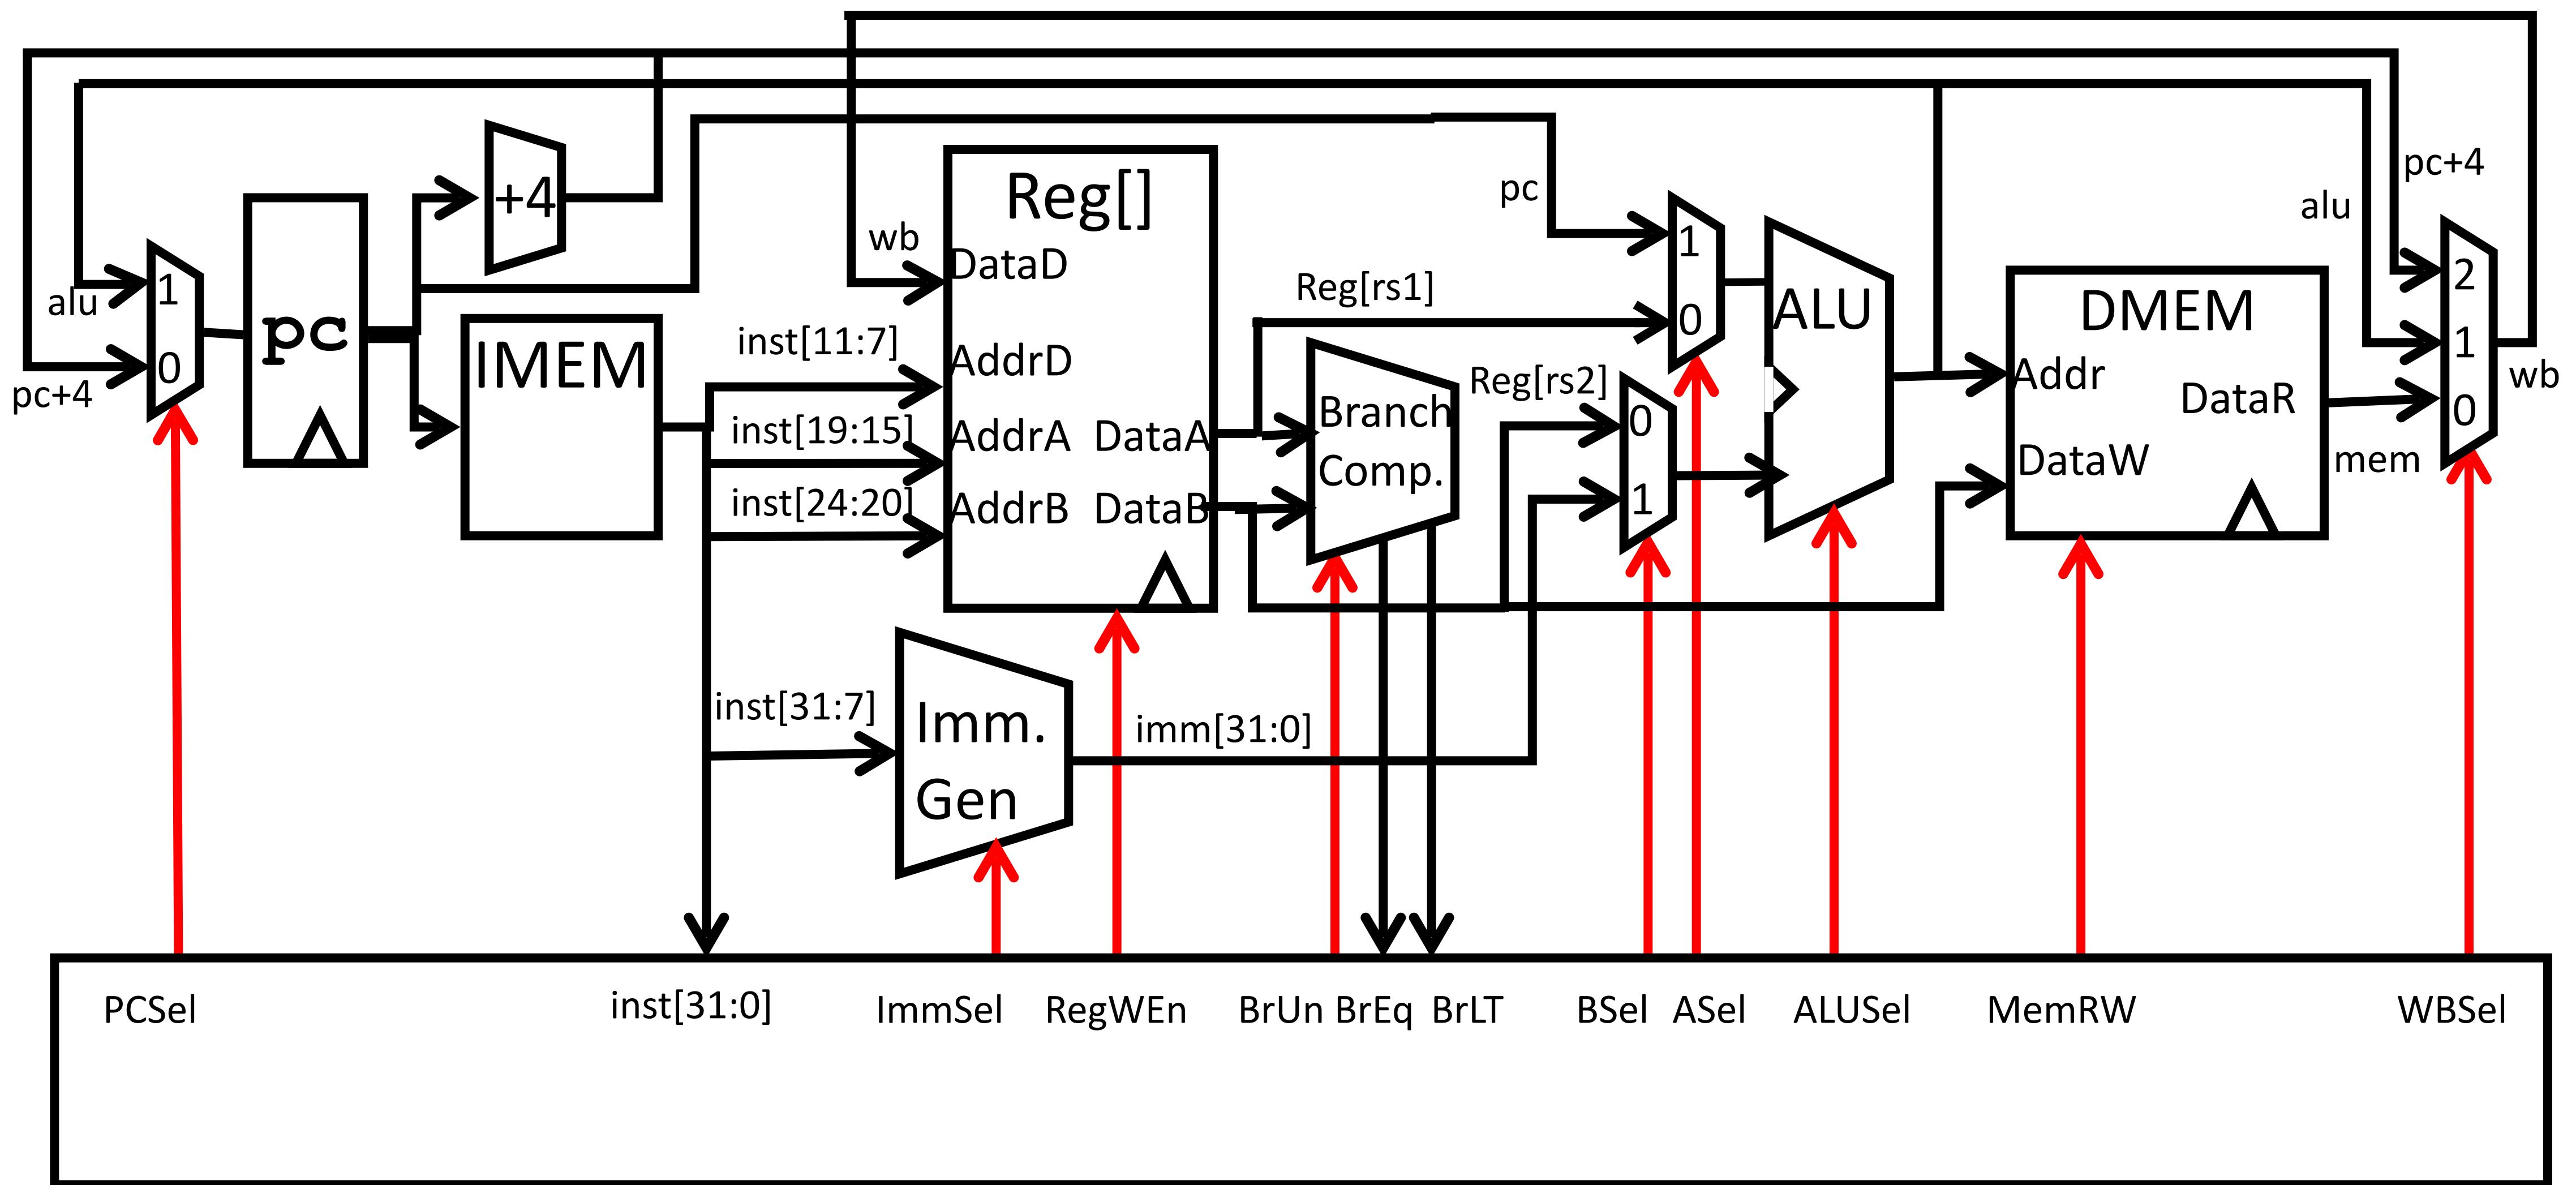


- JALR rd, rs, immediate
 - Writes PC+4 to Reg[rd] (return address)
 - Sets PC = Reg[rs1] + offset
 - Uses same immediates as arithmetic and loads
 - *no* multiplication by 2 bytes

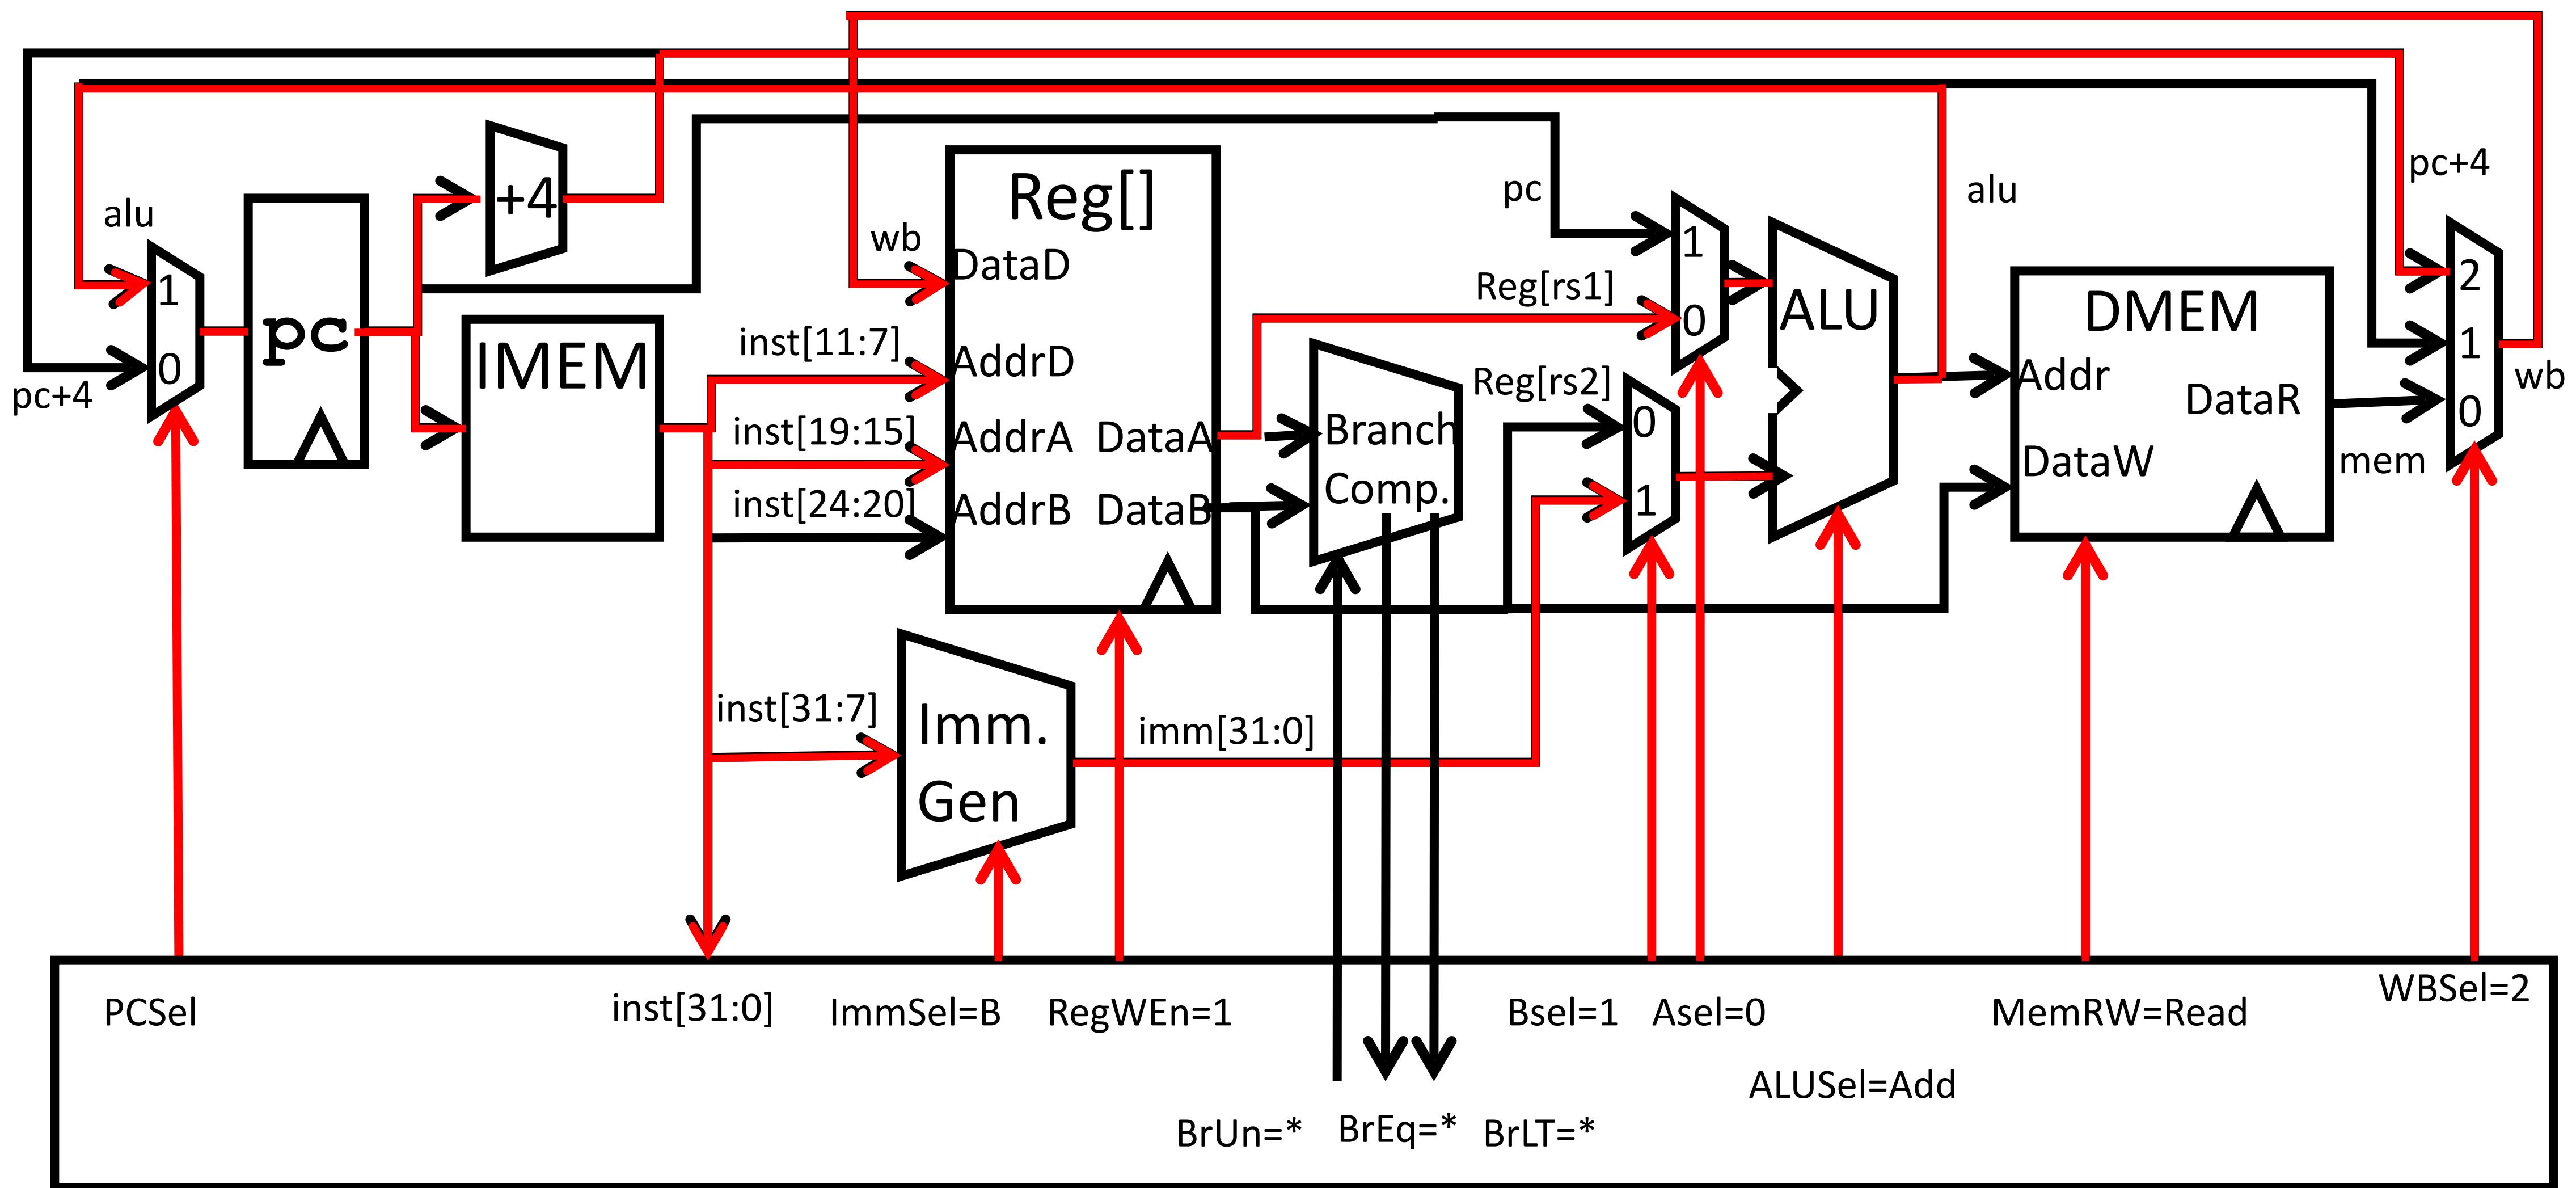
Review: Adding branches to datapath



Adding jalr to datapath

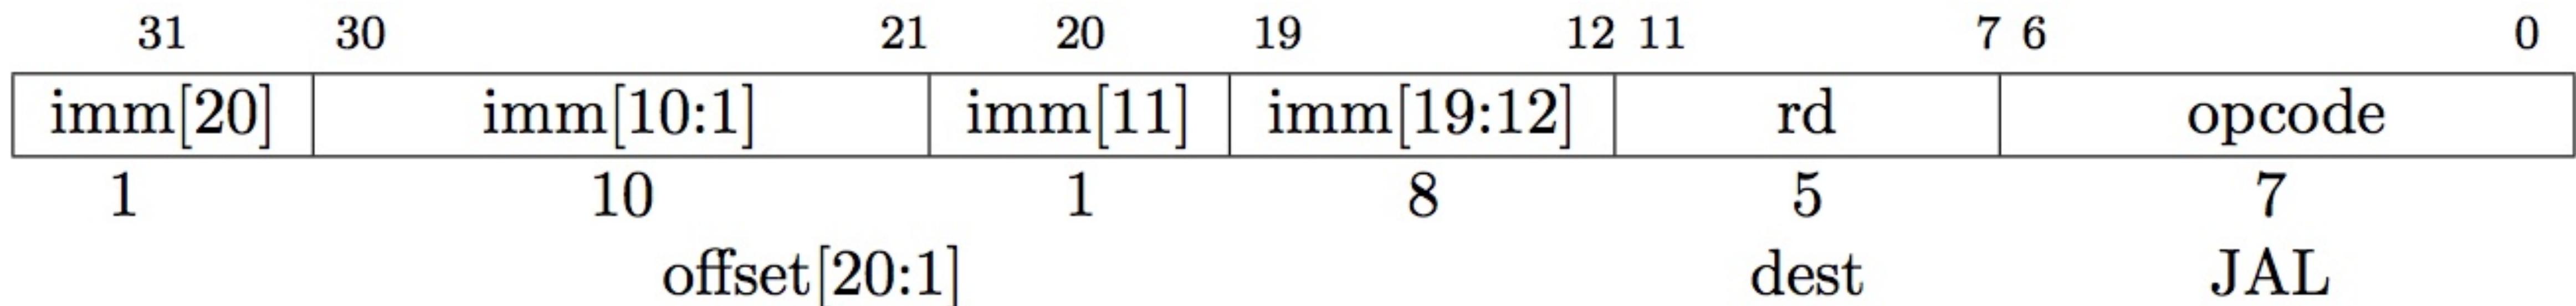


Adding jalr to datapath



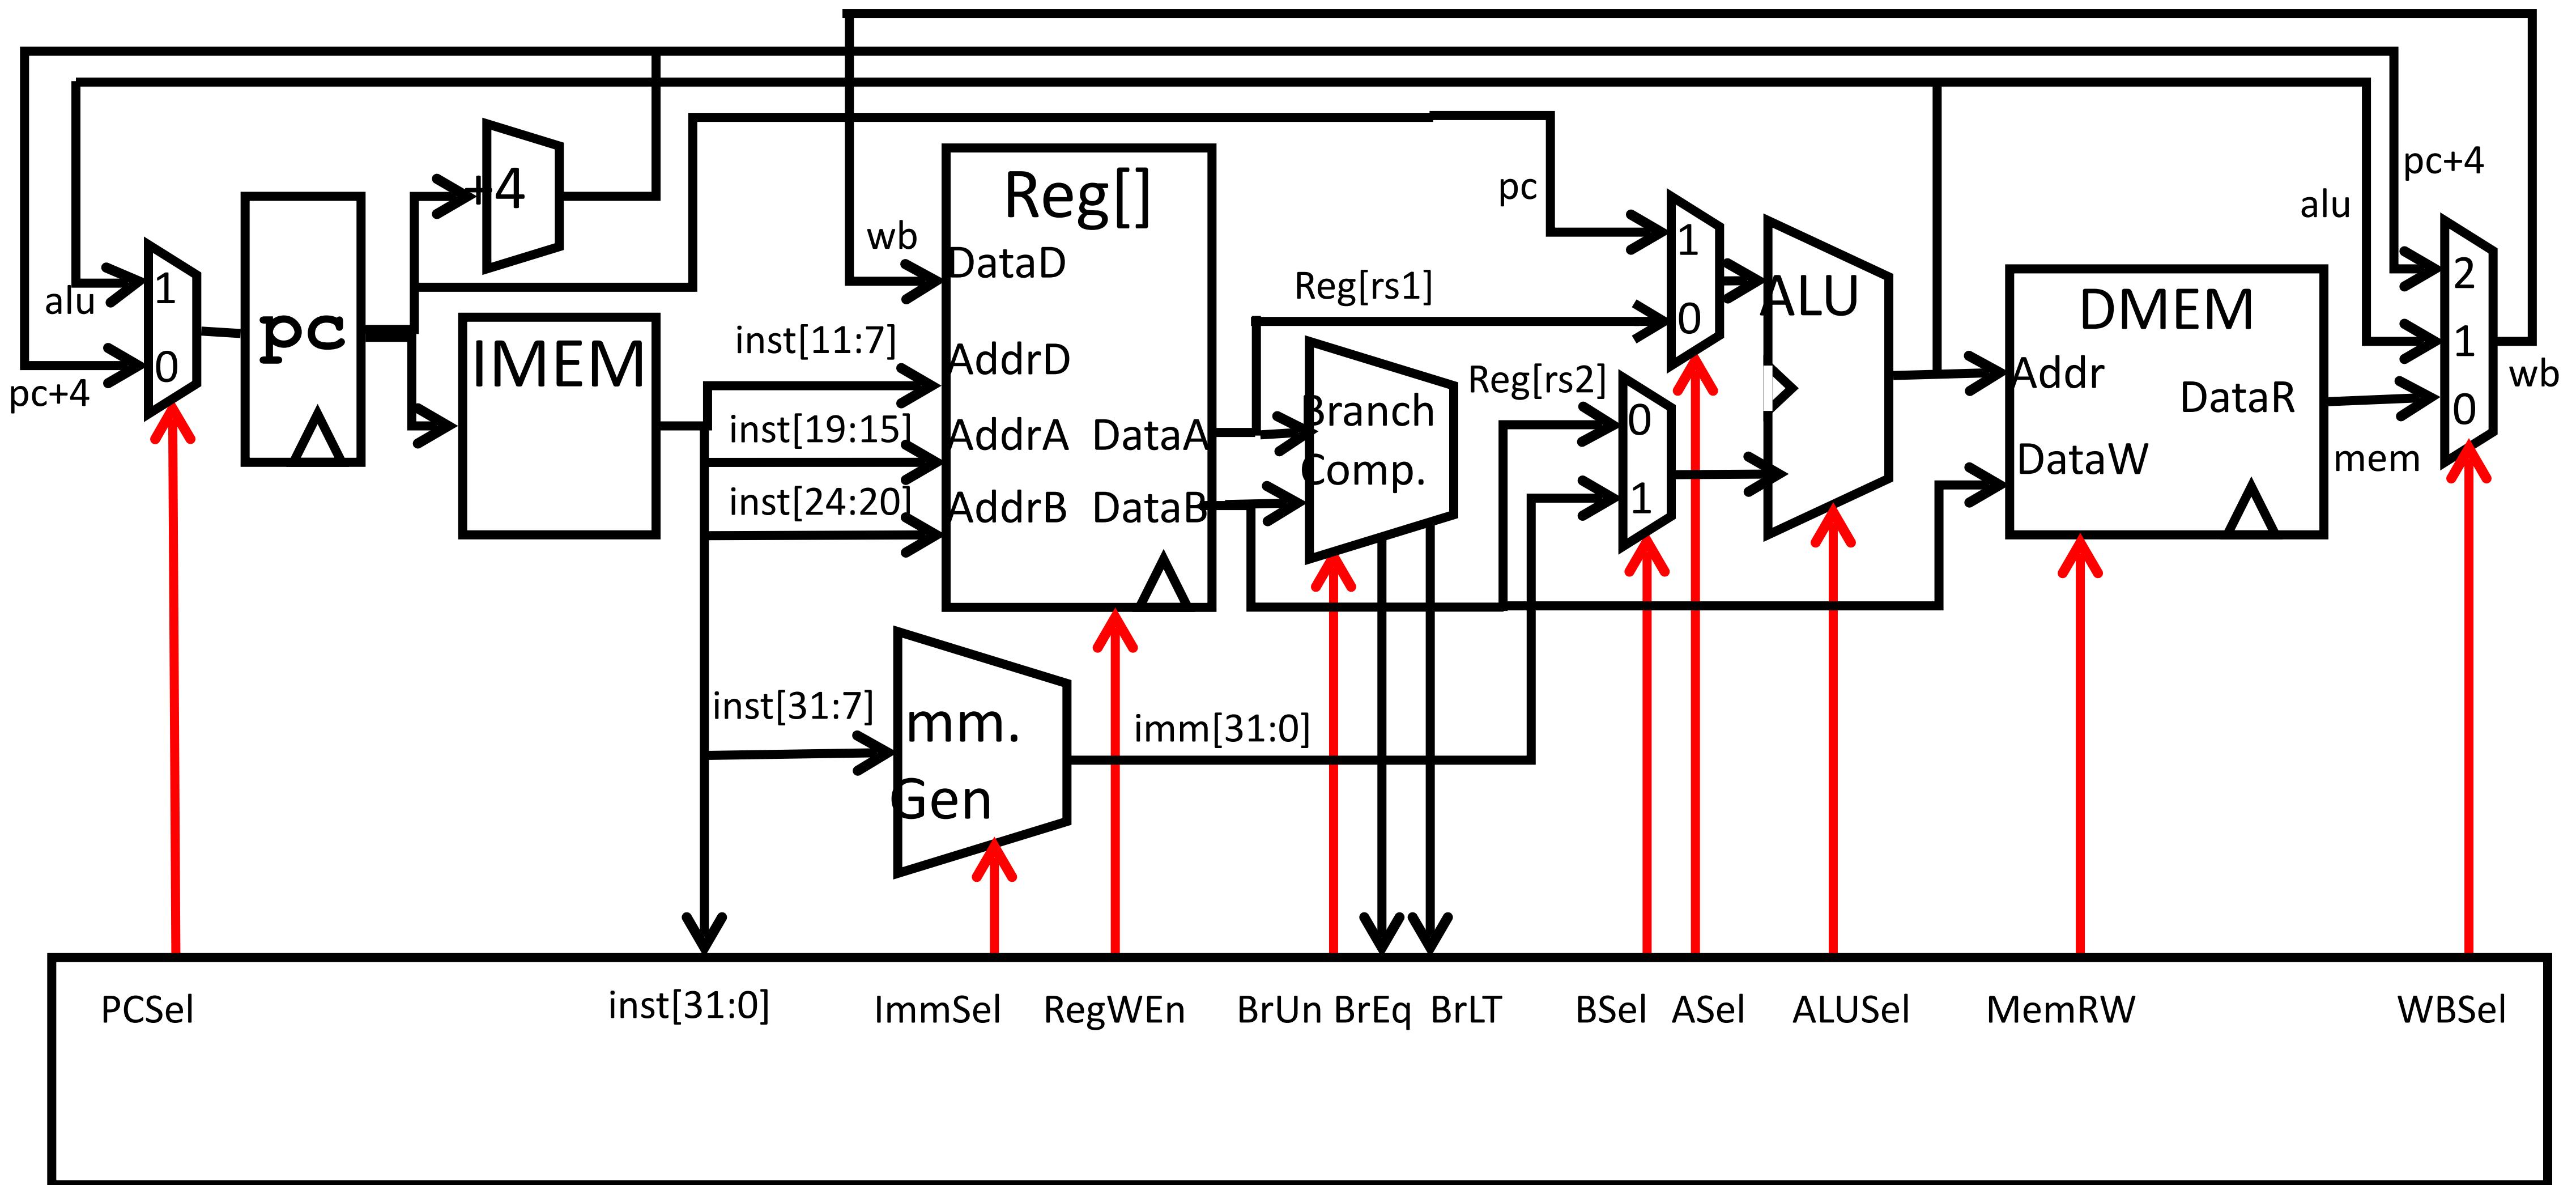
Implementing jal Instruction

Uses the “J-type” instruction format

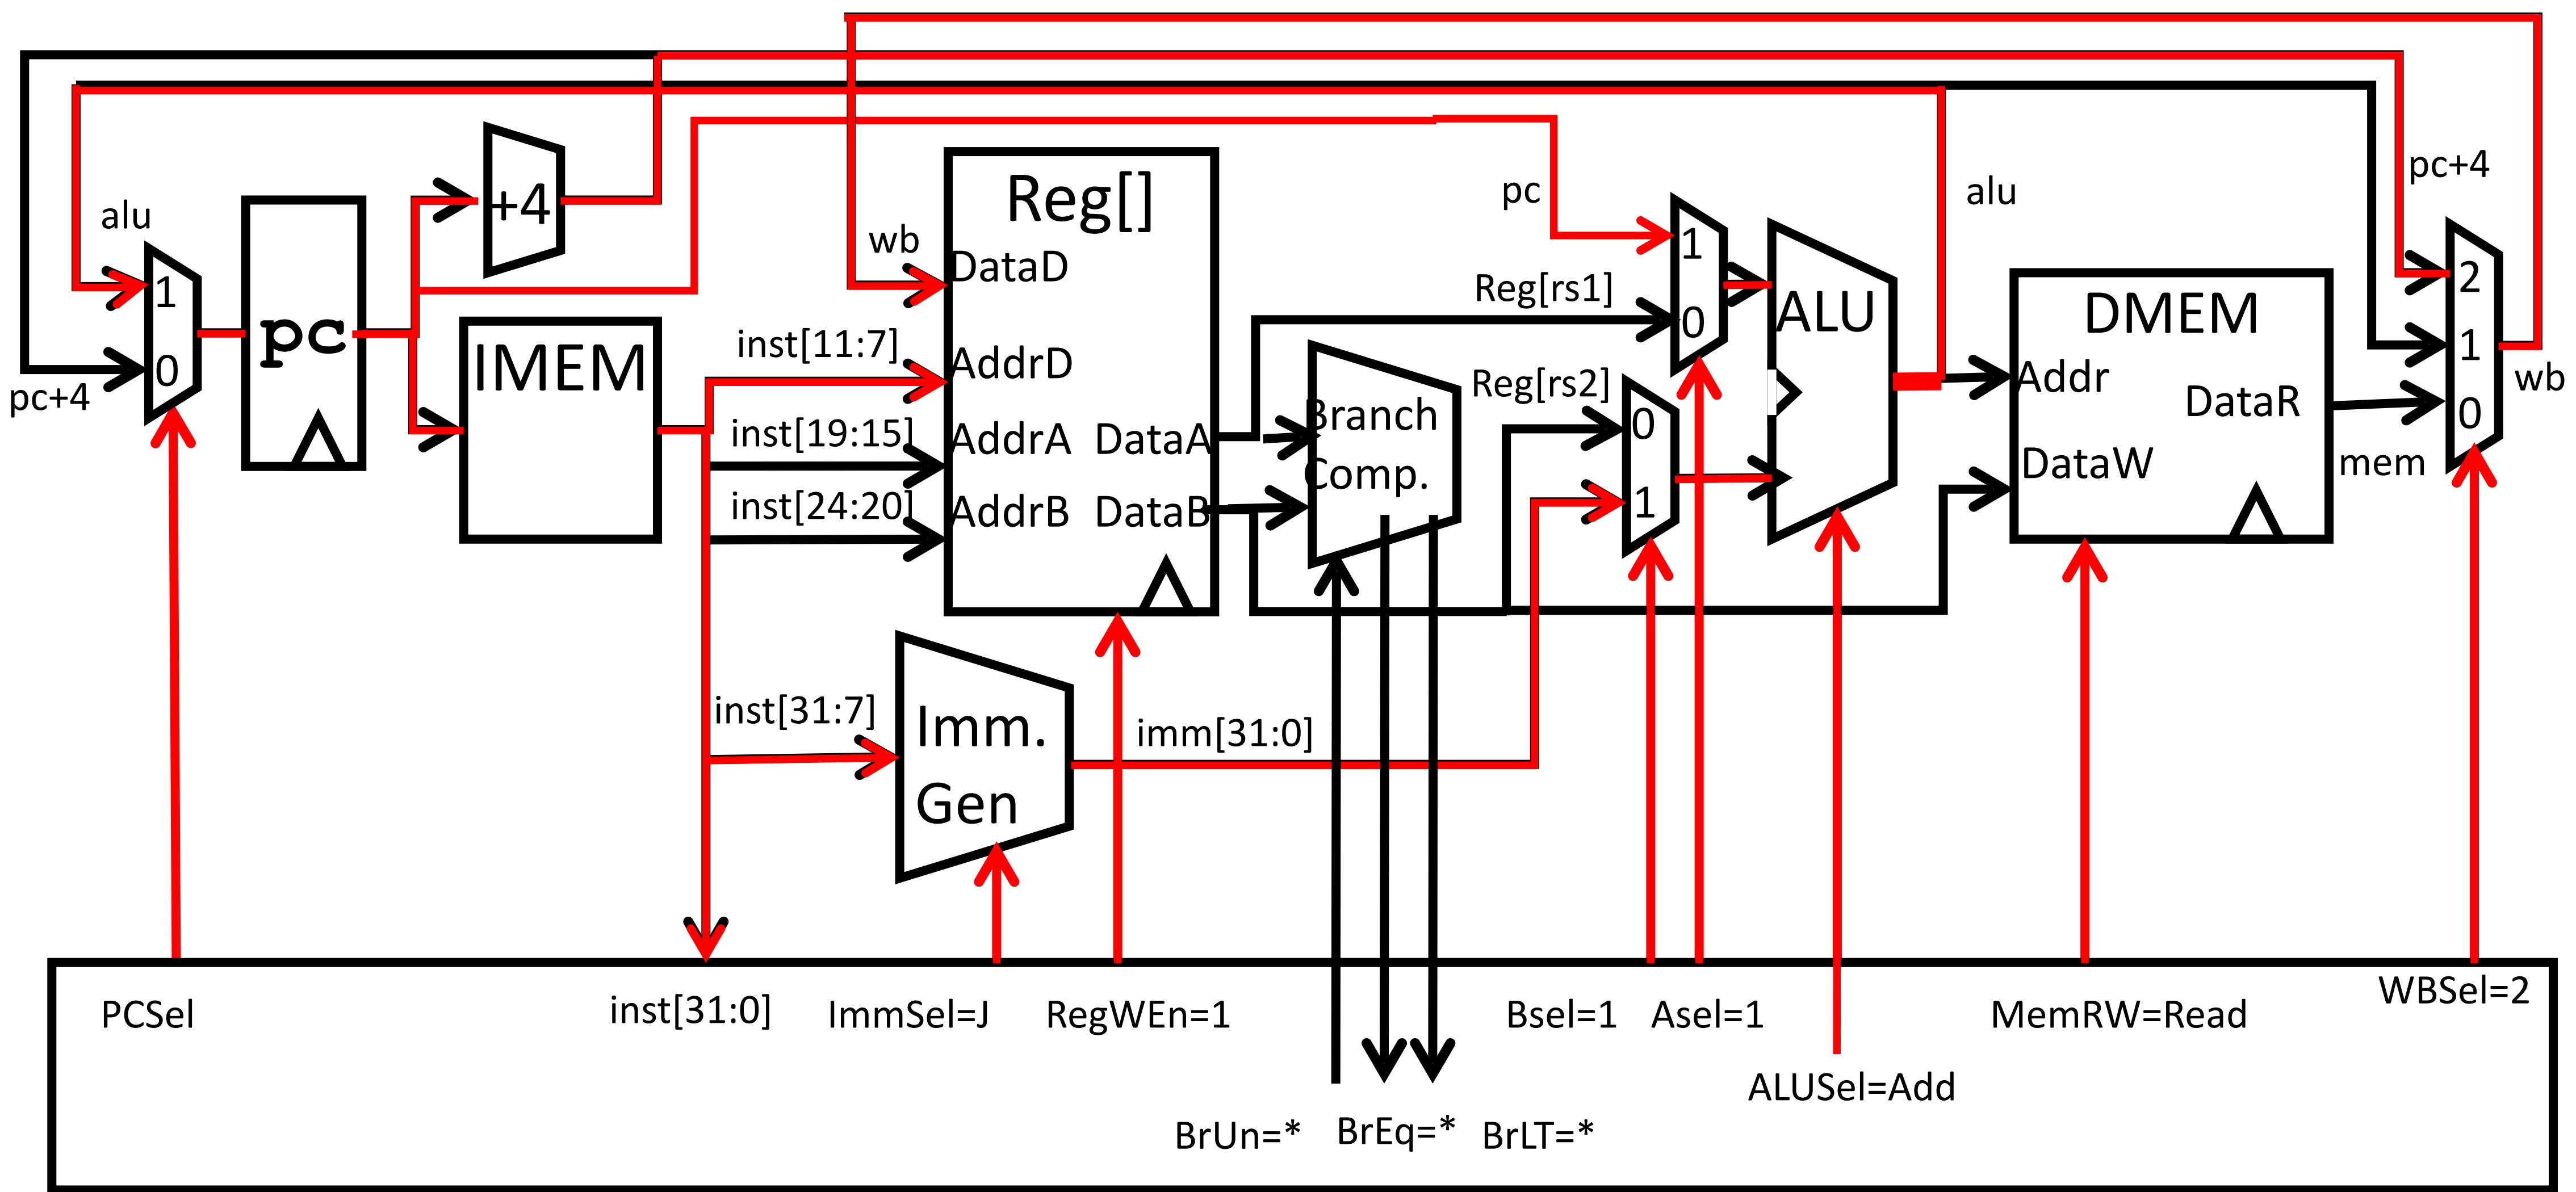


- JAL saves PC+4 in Reg[rd] (the return address)
- Set PC = PC + offset (PC-relative jump)
- Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart
 - $\pm 2^{18}$ 32-bit instructions
- Immediate encoding optimized similarly to branch instruction to reduce hardware cost

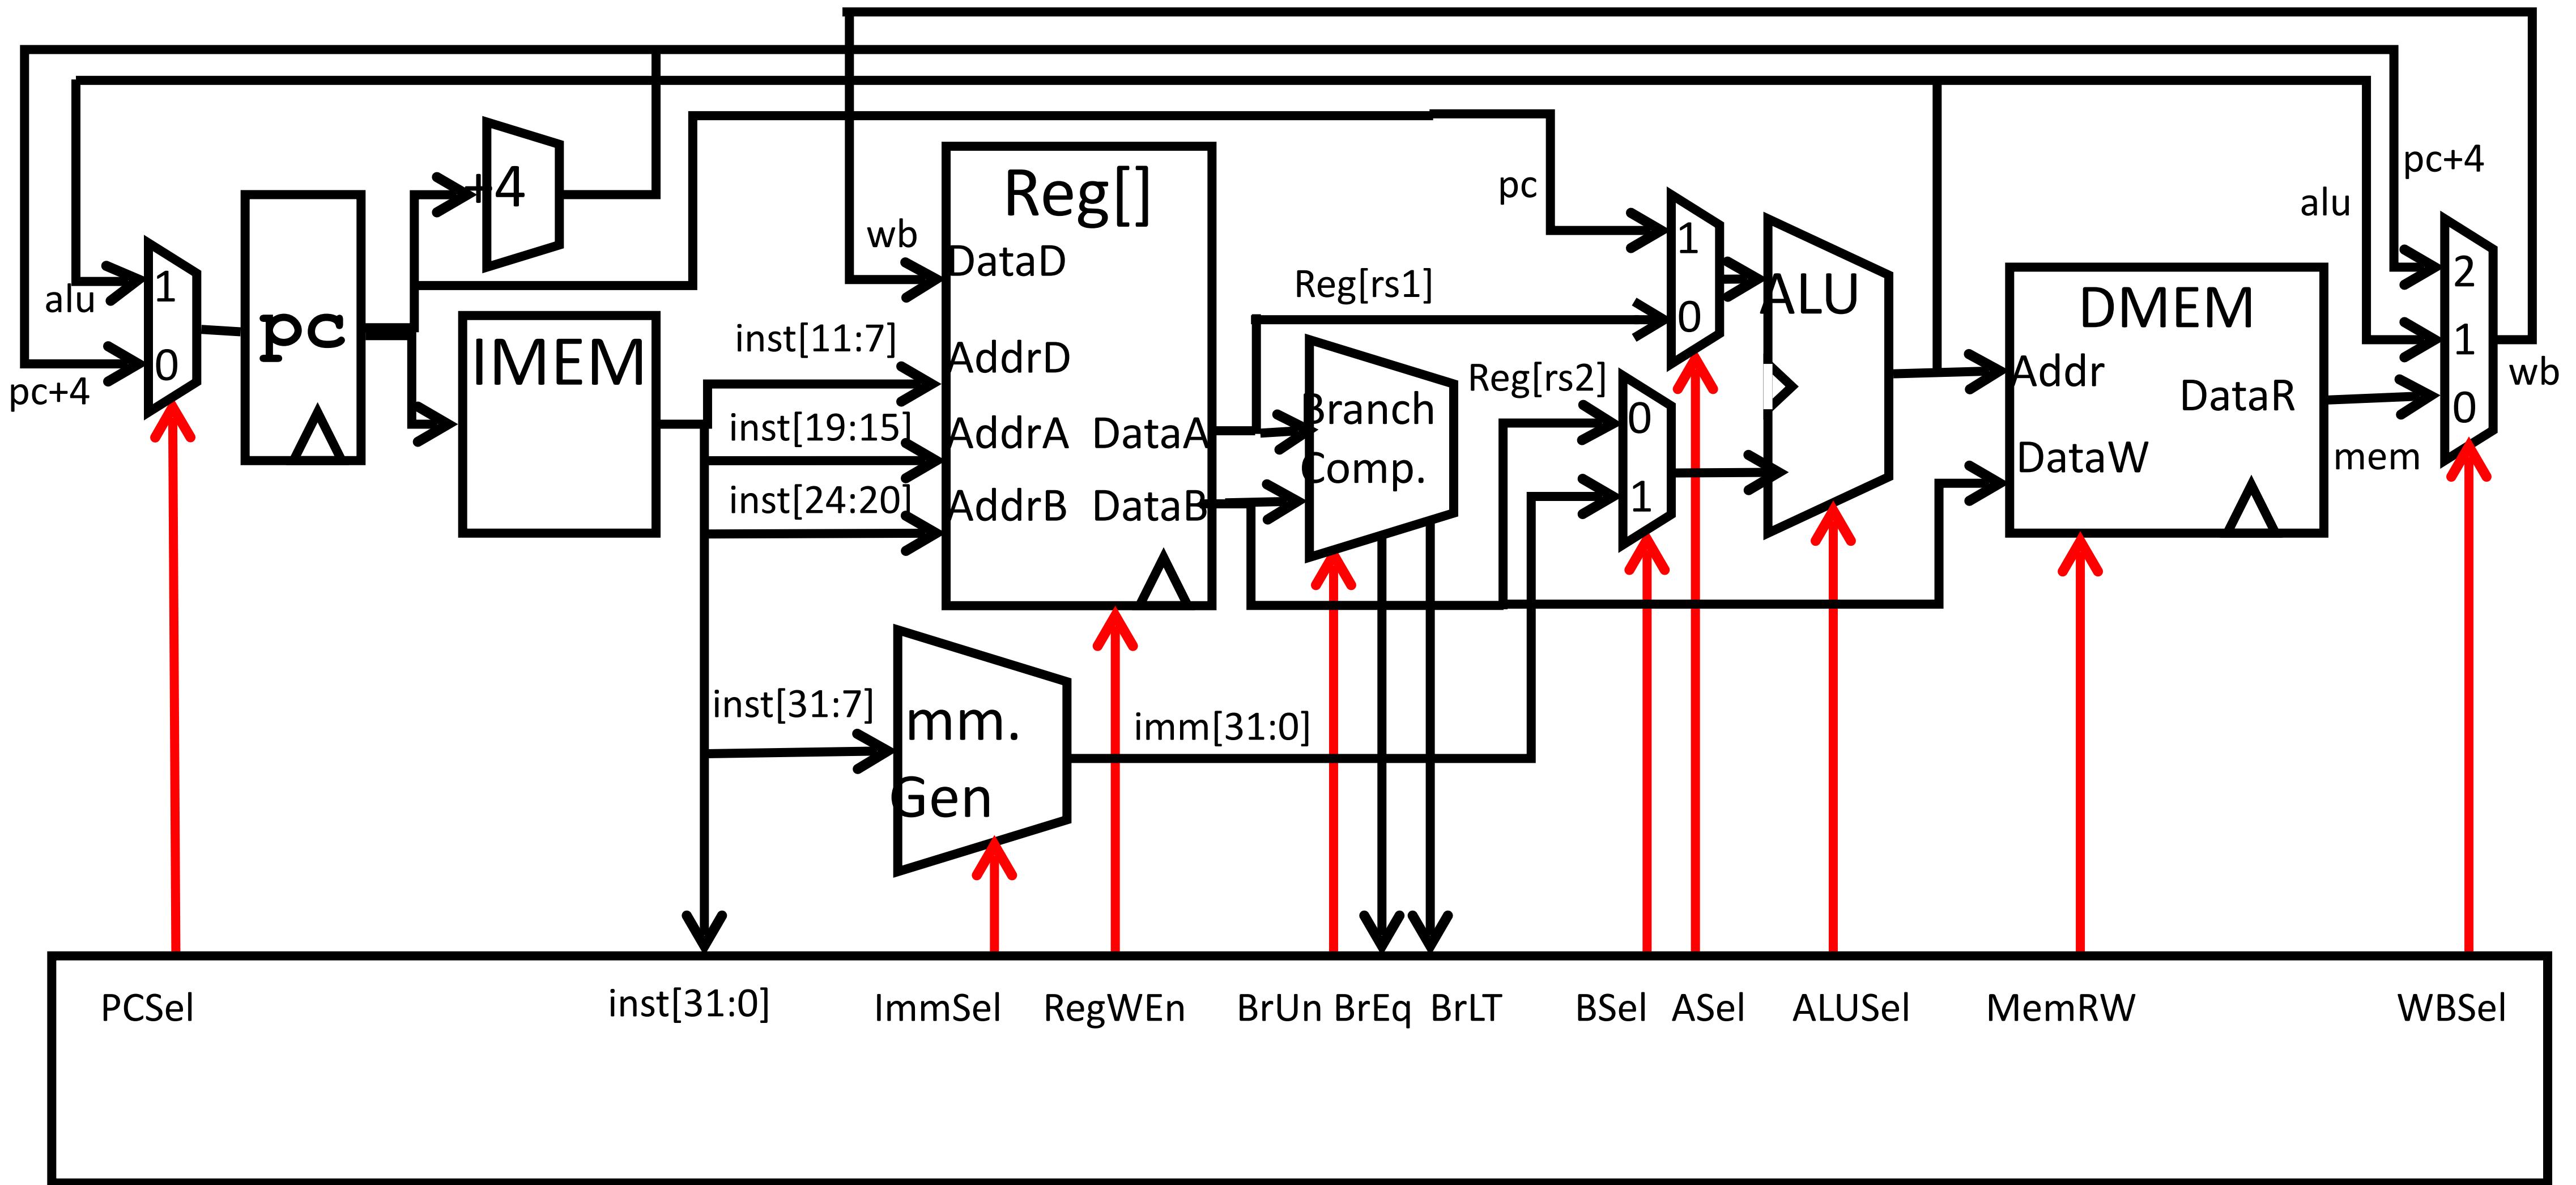
Adding jal to datapath



Adding jail to datapath



Single-Cycle RISC-V RV32I Datapath



Controller Implementation:

- ❑ Control logic works really well as a case statement...

```
always @* begin
    op = instr[26:31];
    imm = instr[15:0]; ...

    reg_dst = 1'bx;    // Don't care
    reg_write = 1'b0; // By default don't write
    ...
    case (op)
        6'b000000: begin reg_write = 1; ... end
        ...
    endcase
end
```

COVID-19 Announcements (tentative)

- ❑ Lectures:
 - ❑ All lectures are recorded and available online.
 - ❑ If you feel ill, *please do not come to class*.
 - ❑ If you feel uncomfortable coming to class, or would simply like to do your part to help stop the spread, please feel free not to come.
- ❑ Labs:
 - ❑ Lab meetings will continue, however, ...
 - ❑ If you feel ill, *please do not come to lab*, but do let us know.
 - ❑ All lab software is available remotely, so feel free to not work in 125 Cory.
 - ❑ FPGA boards are available for check-out.
 - ❑ Check-offs will be done in person as before, or now can be done with Zoom/Skype (contact Tan).

COVID-19 Announcements (tentative)

- ❑ Exam:
 - ❑ To be held during previously scheduled time slot, 6-9PM Thursday.
 - ❑ However, it will be take-at-home style.
 - ❑ Exam will be posted at 6PM and your solution must be turned in using gradescope by 9PM.
 - ❑ You will be allowed to use lecture notes and other resources, but not allowed to discuss exam questions or answers with other students during the exam time.
- ❑ Office hours:
 - ❑ for meeting the GSIs or professor, in addition to in person meetings, teleconference (Skype/Zoom) will be available
- ❑ Discussion Session:
 - ❑ Will be held as usual, and Zoom will be used for remote participation.



Processor Pipelining

Review: Processor Performance

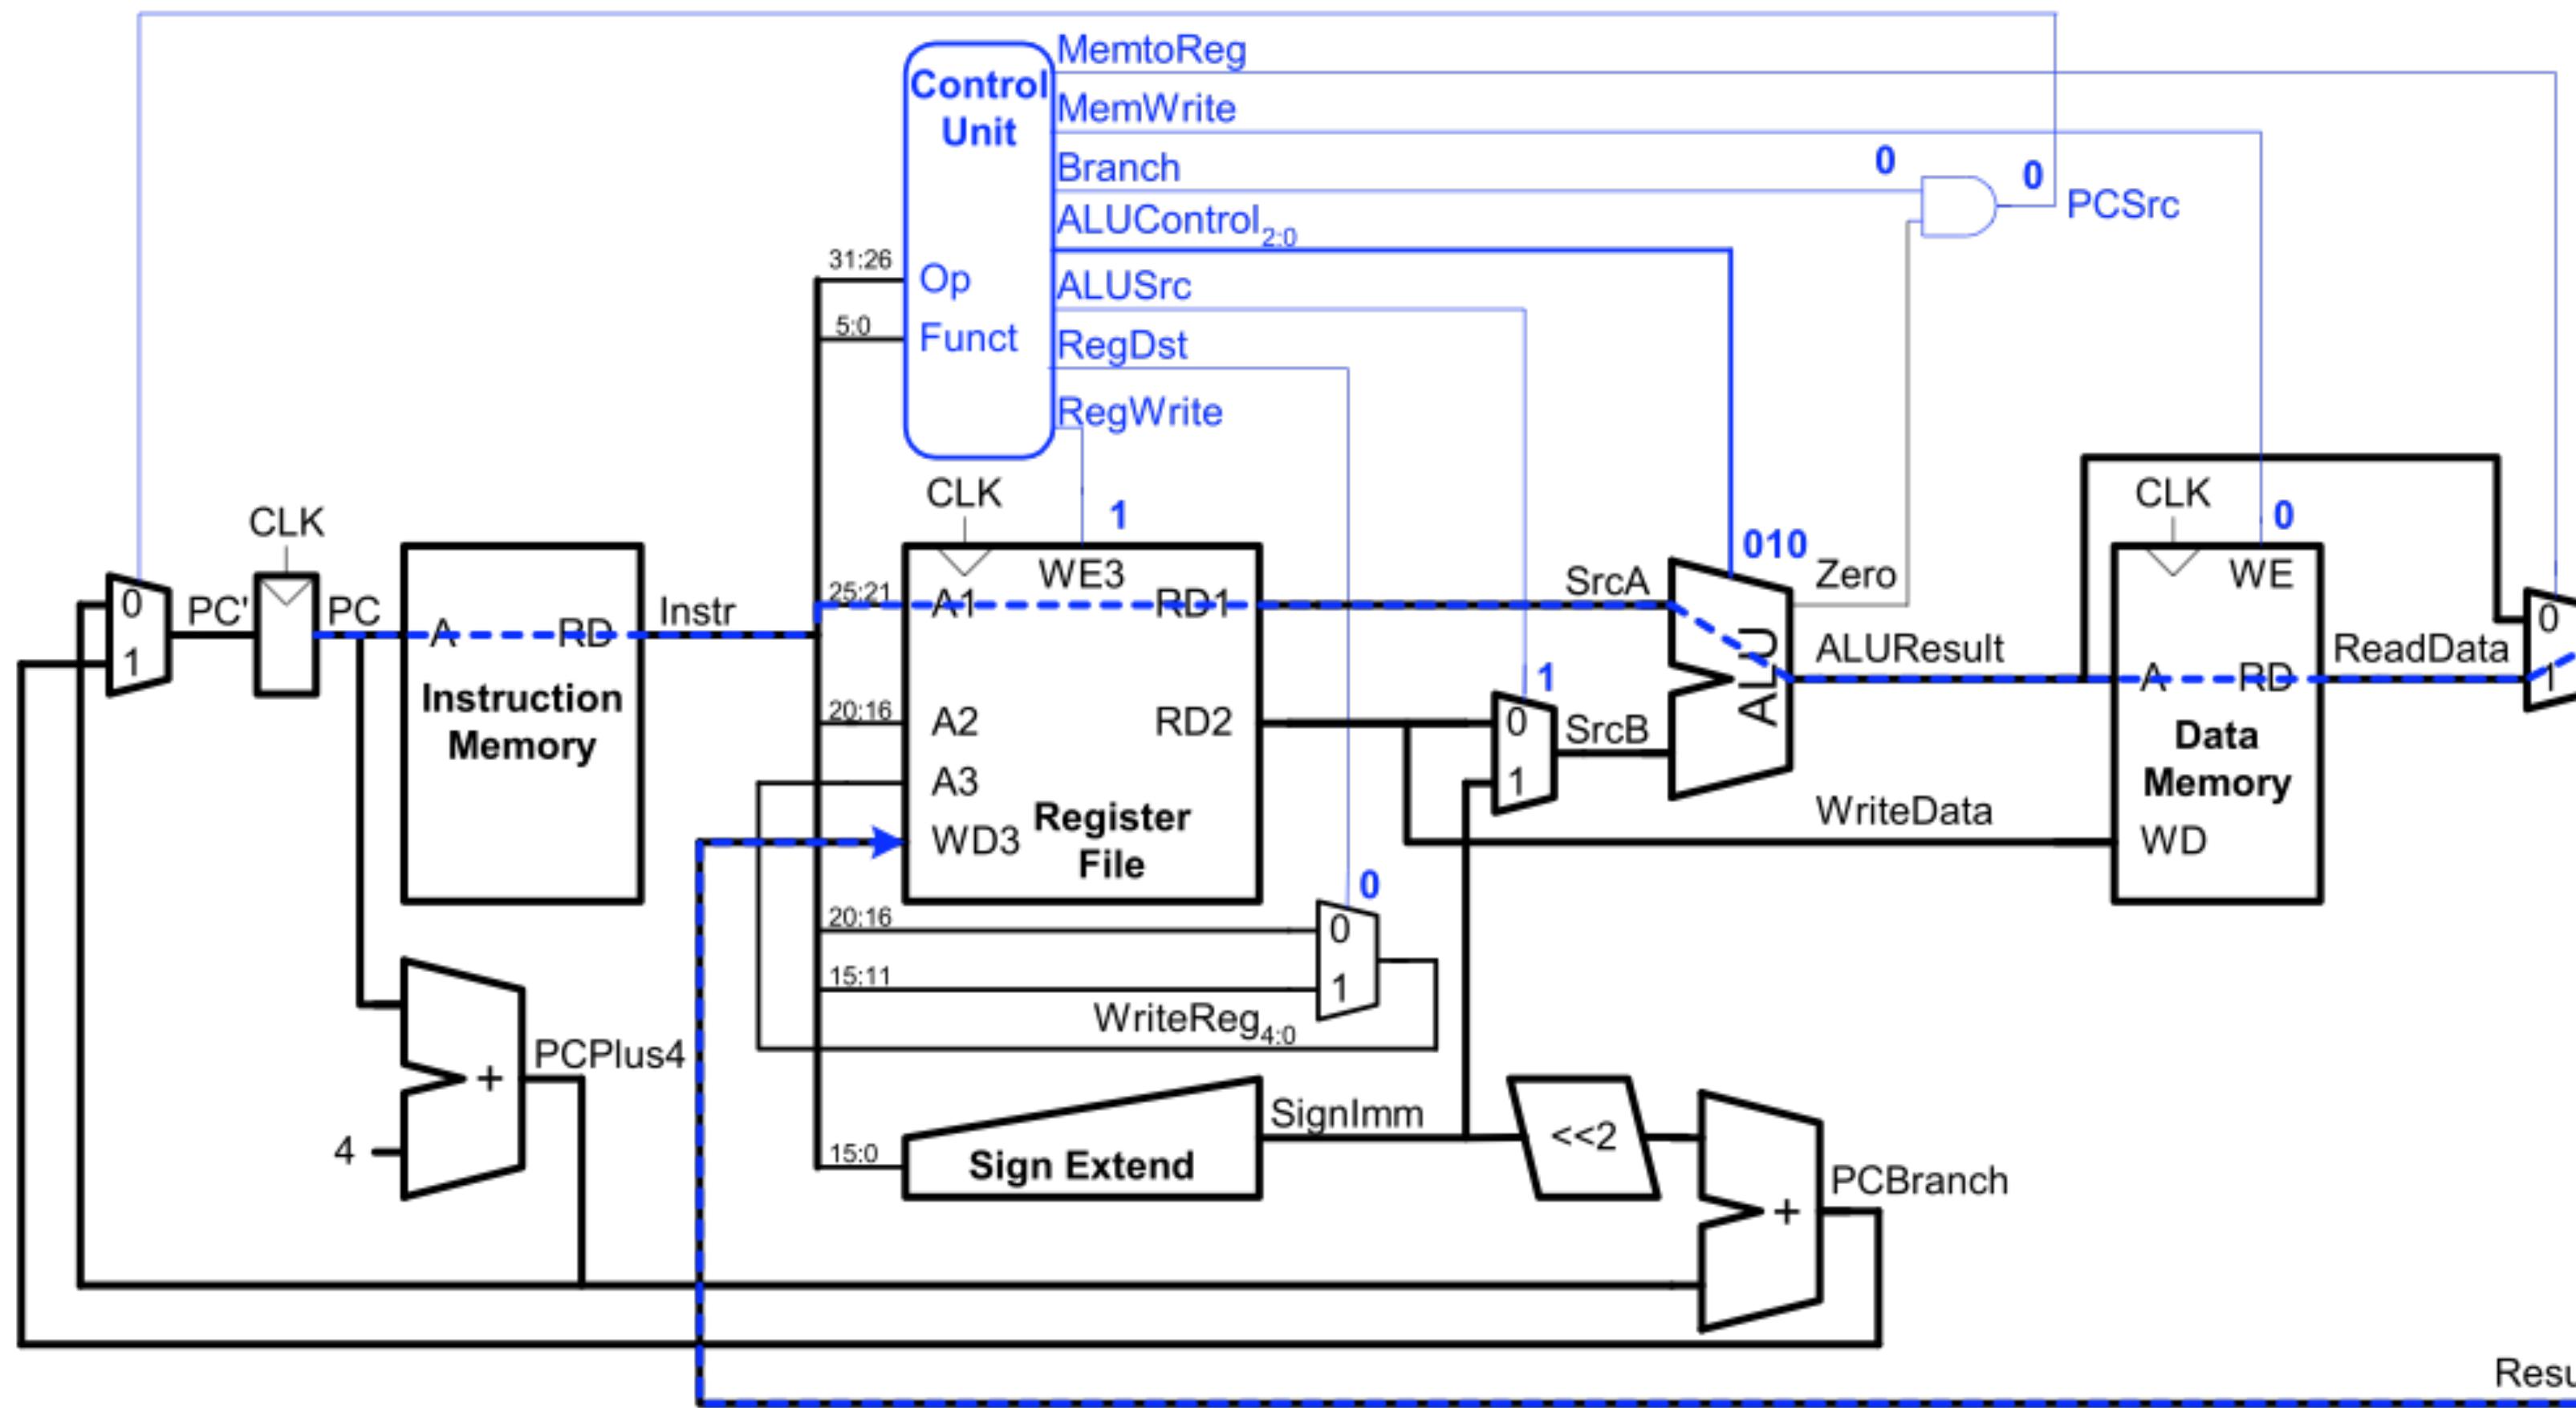
Program Execution Time

$$= (\# \text{ instructions})(\text{cycles/instruction})(\text{seconds/cycle})$$

$$= \# \text{ instructions} \times \text{CPI} \times T_C$$

Single-Cycle Performance

- T_C is limited by the critical path ($1w$)



Single-Cycle Performance

- *Single-cycle critical path:*

$$T_c = t_{q_PC} + t_{mem} + \max(t_{RFread}, t_{sext} + t_{mux}) + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$$

- *In most implementations, limiting paths are:*

- *memory, ALU, register file.*

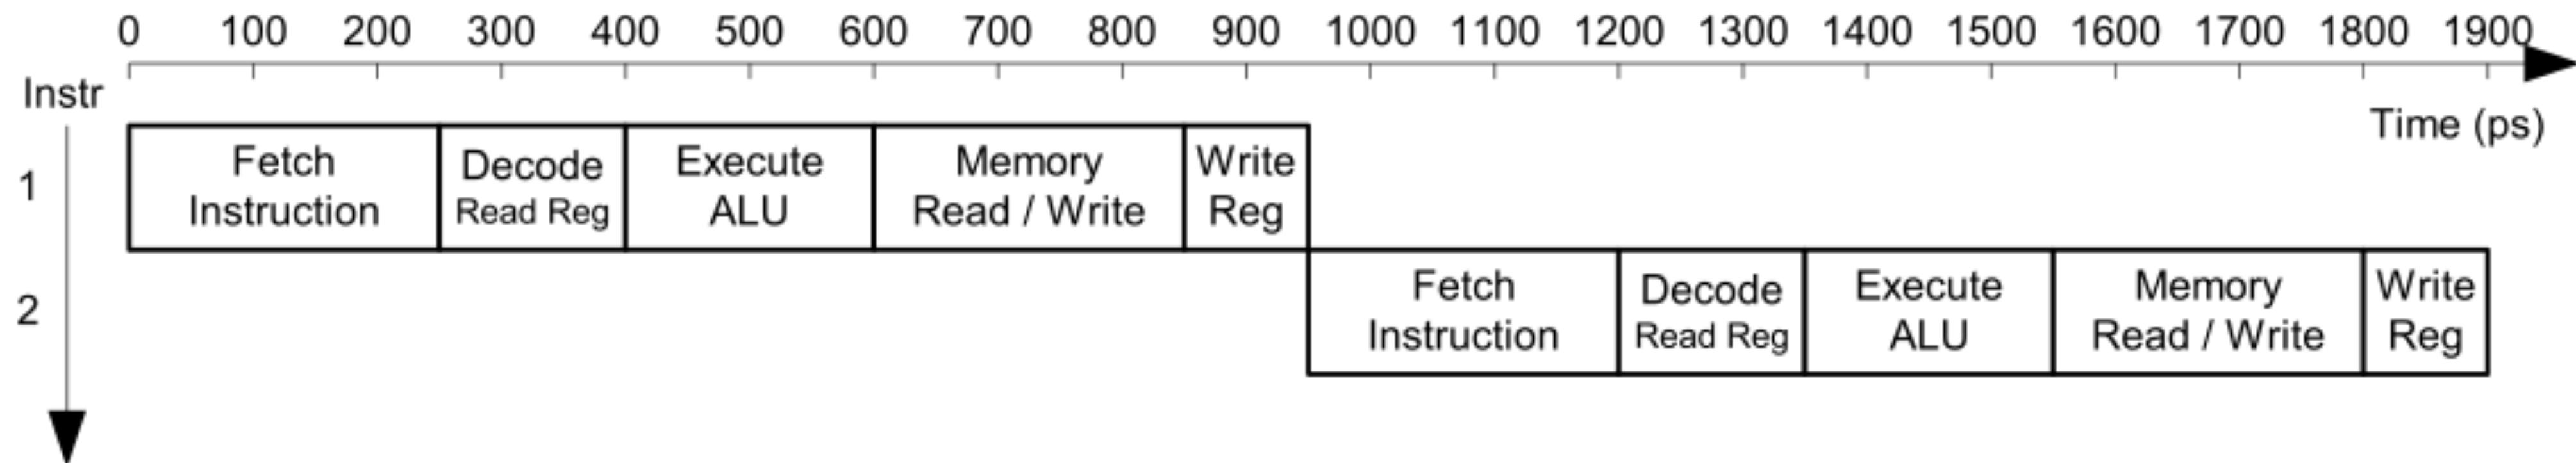
- $T_c = t_{q_PC} + 2t_{mem} + t_{RFread} + t_{mux} + t_{ALU} + t_{RFsetup}$

Pipelined Processor

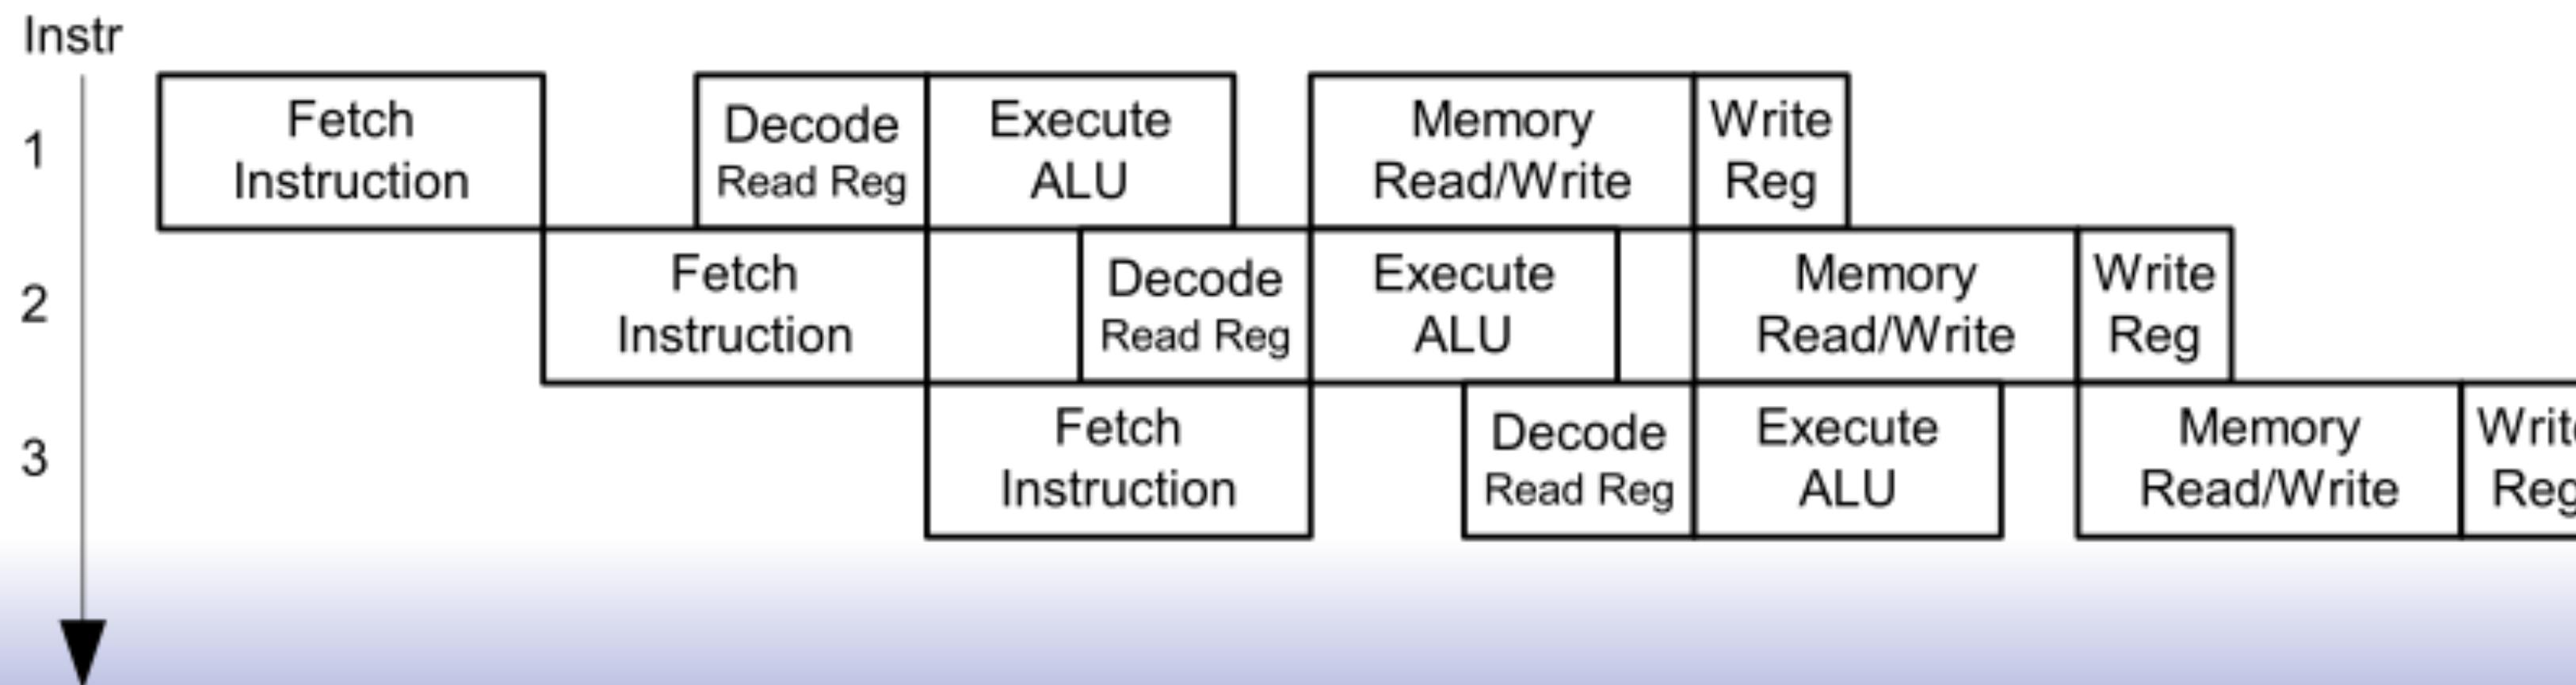
- Use temporal parallelism
- Divide single-cycle processor into 5 stages:
 - Fetch
 - Decode
 - Execute
 - Memory
 - Writeback
- Add pipeline registers between stages

Single-Cycle vs. Pipelined Performance

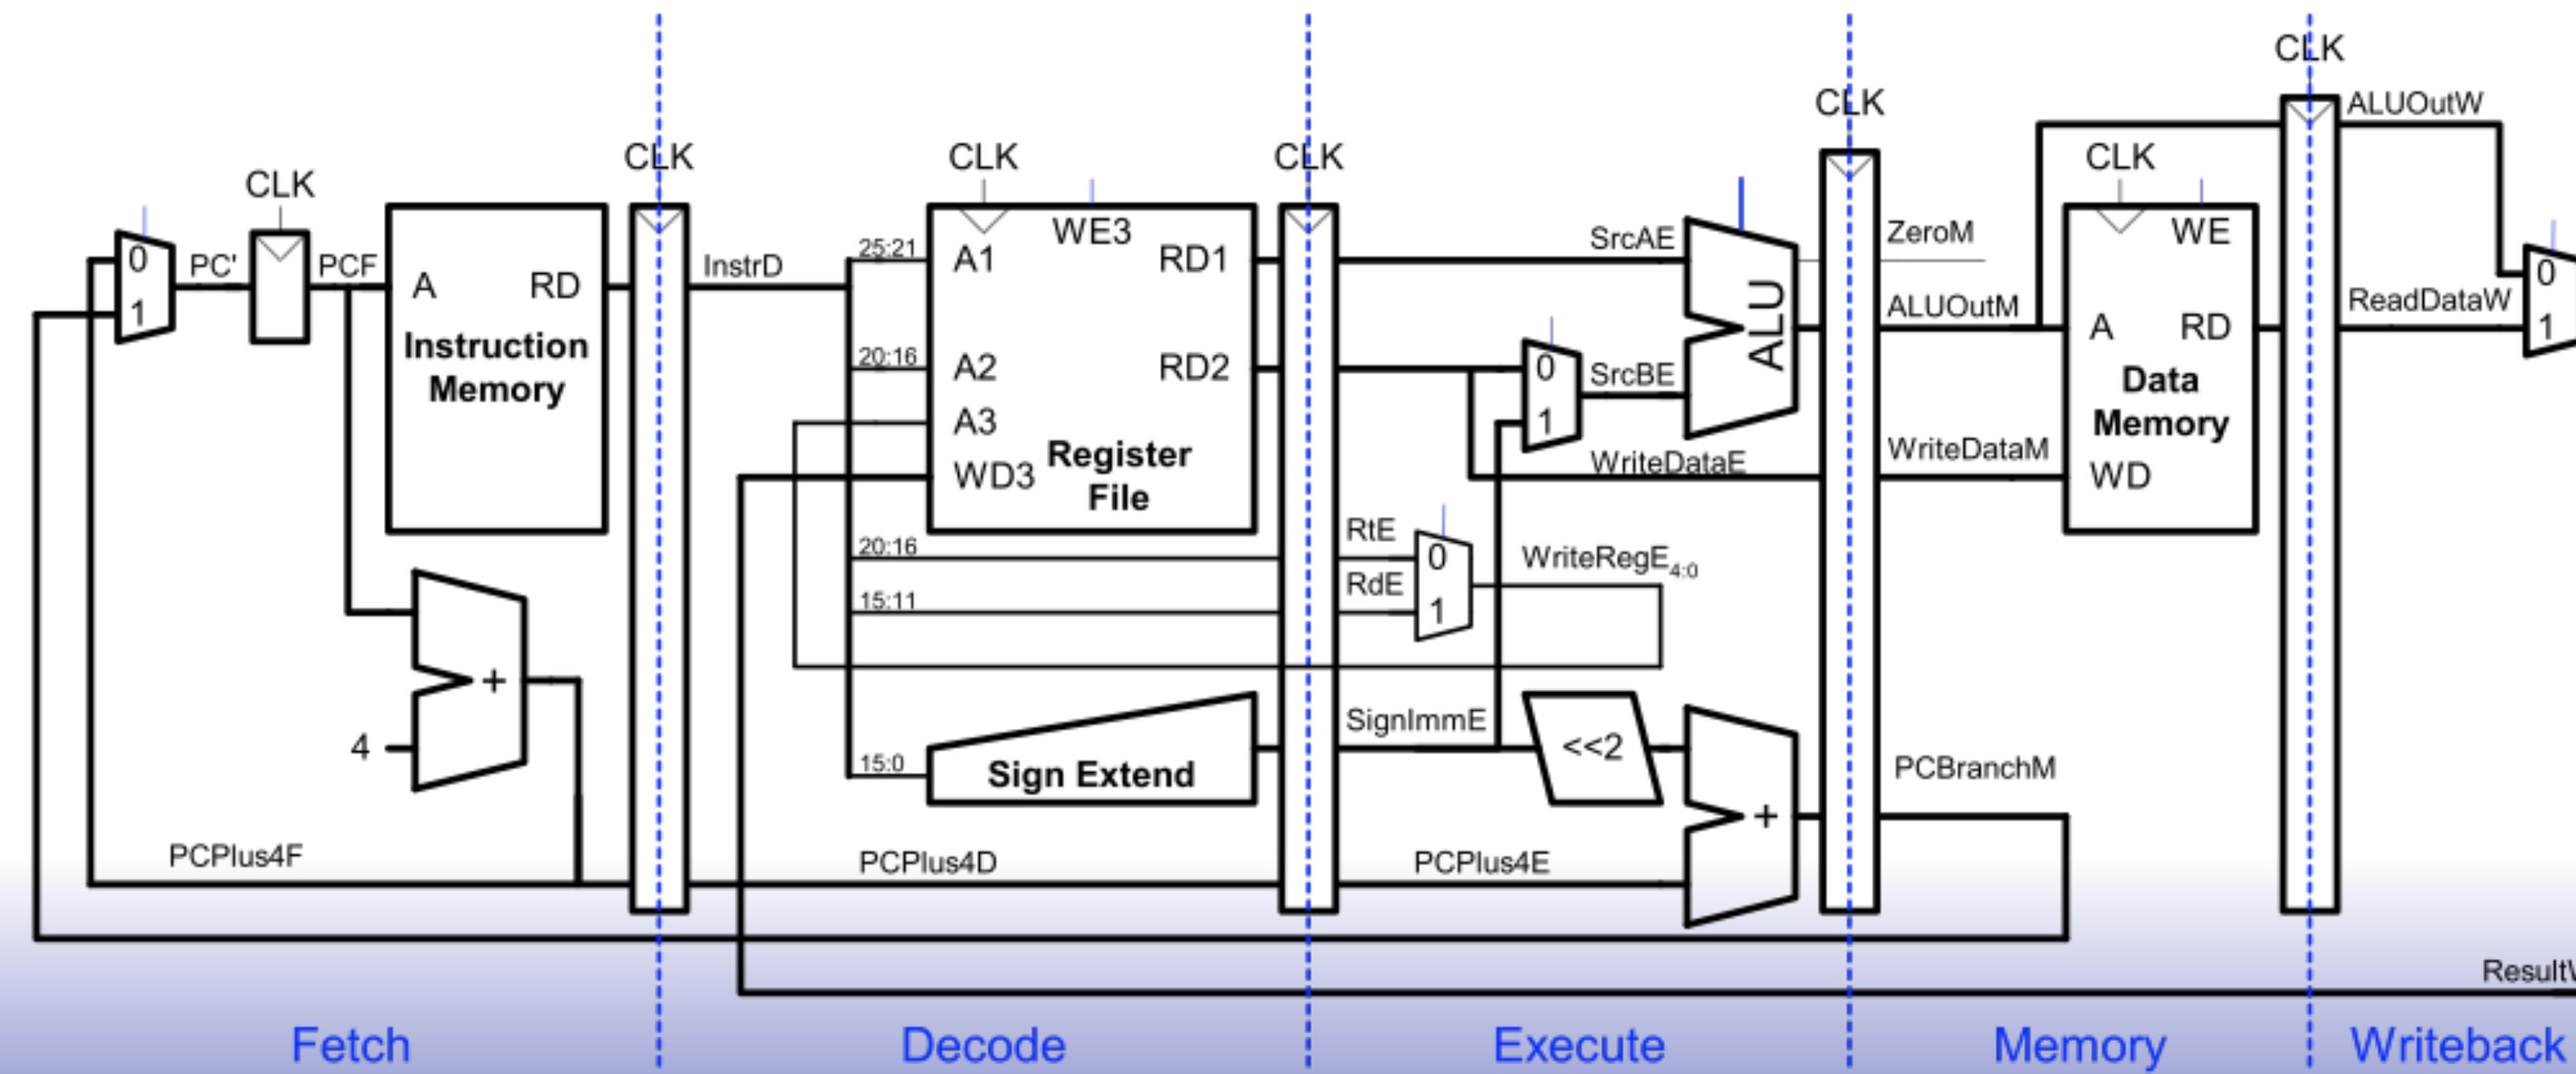
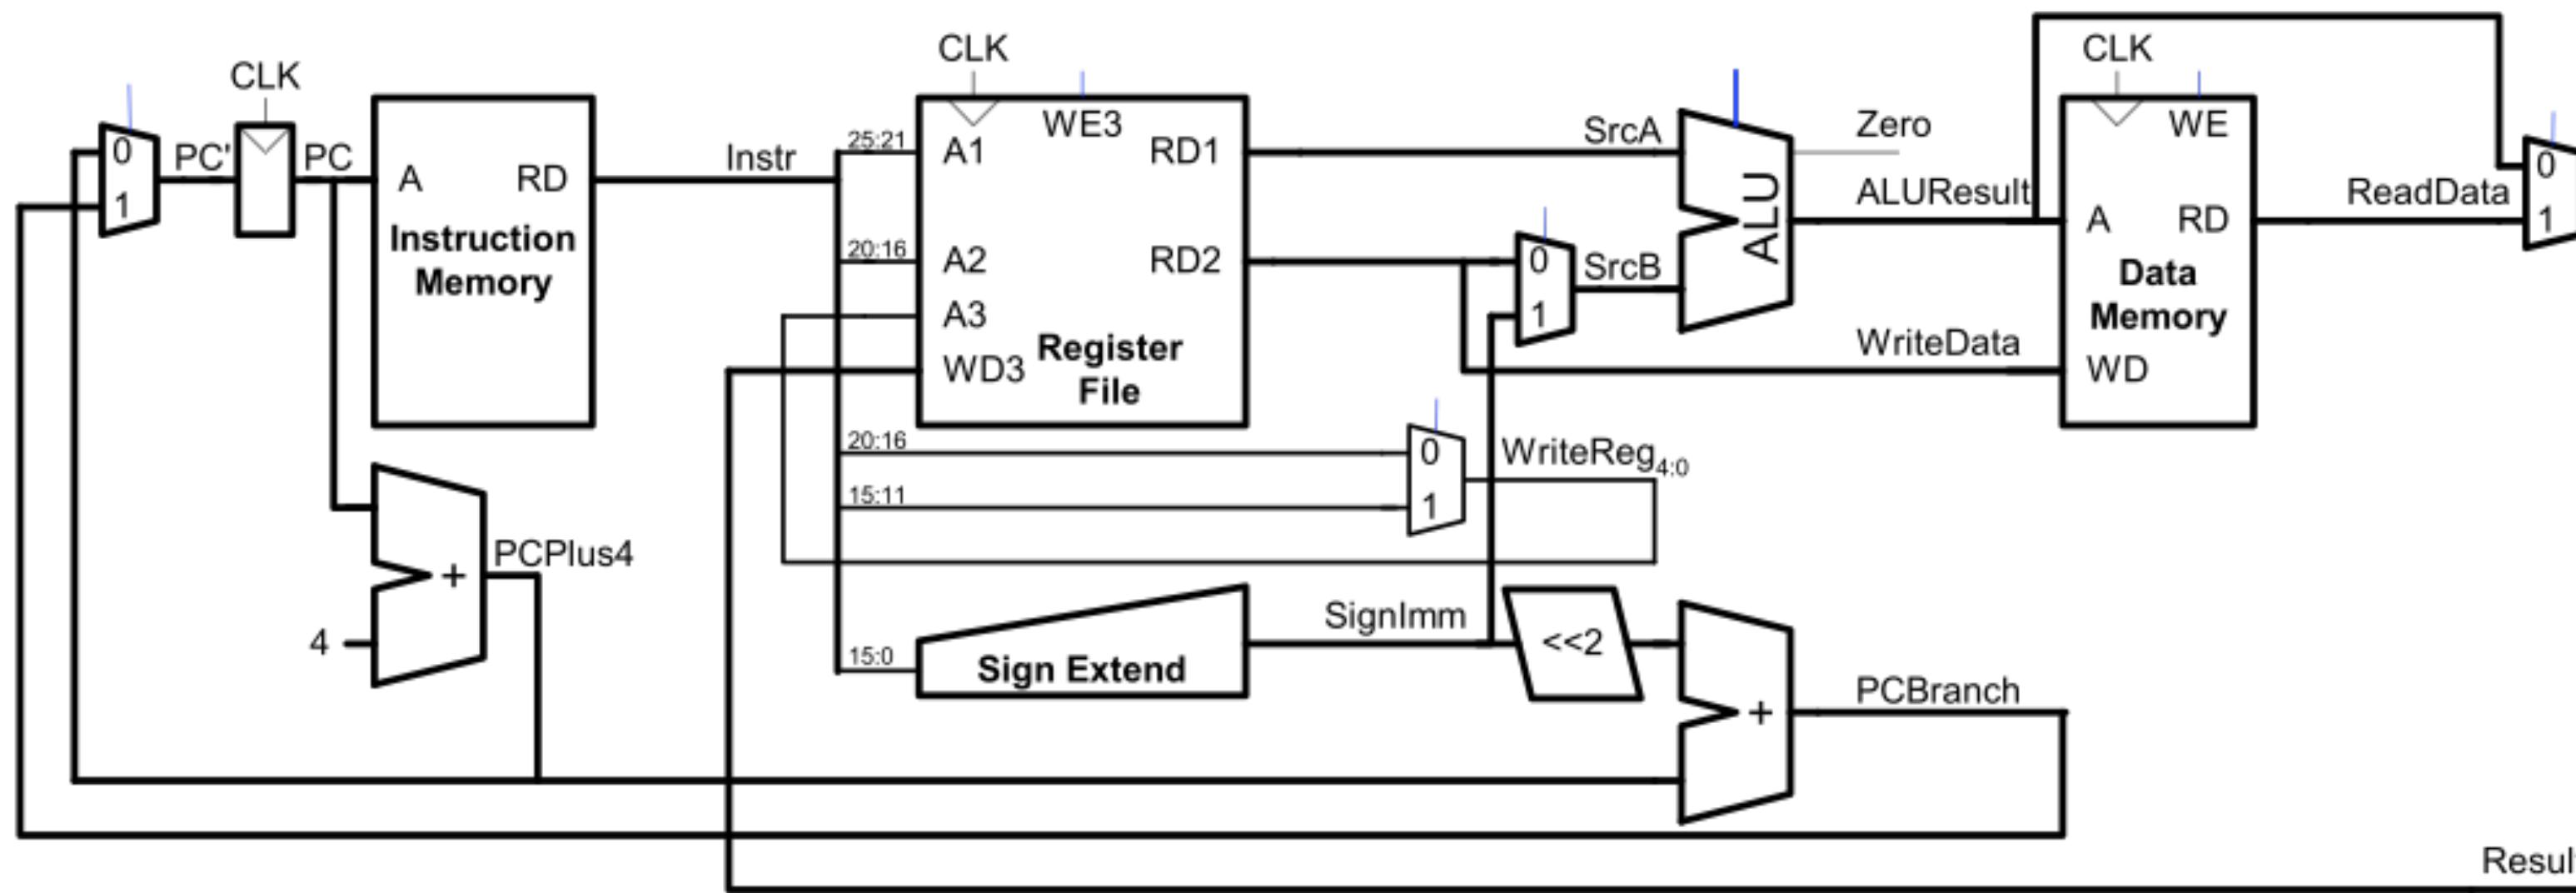
Single-Cycle



Pipelined

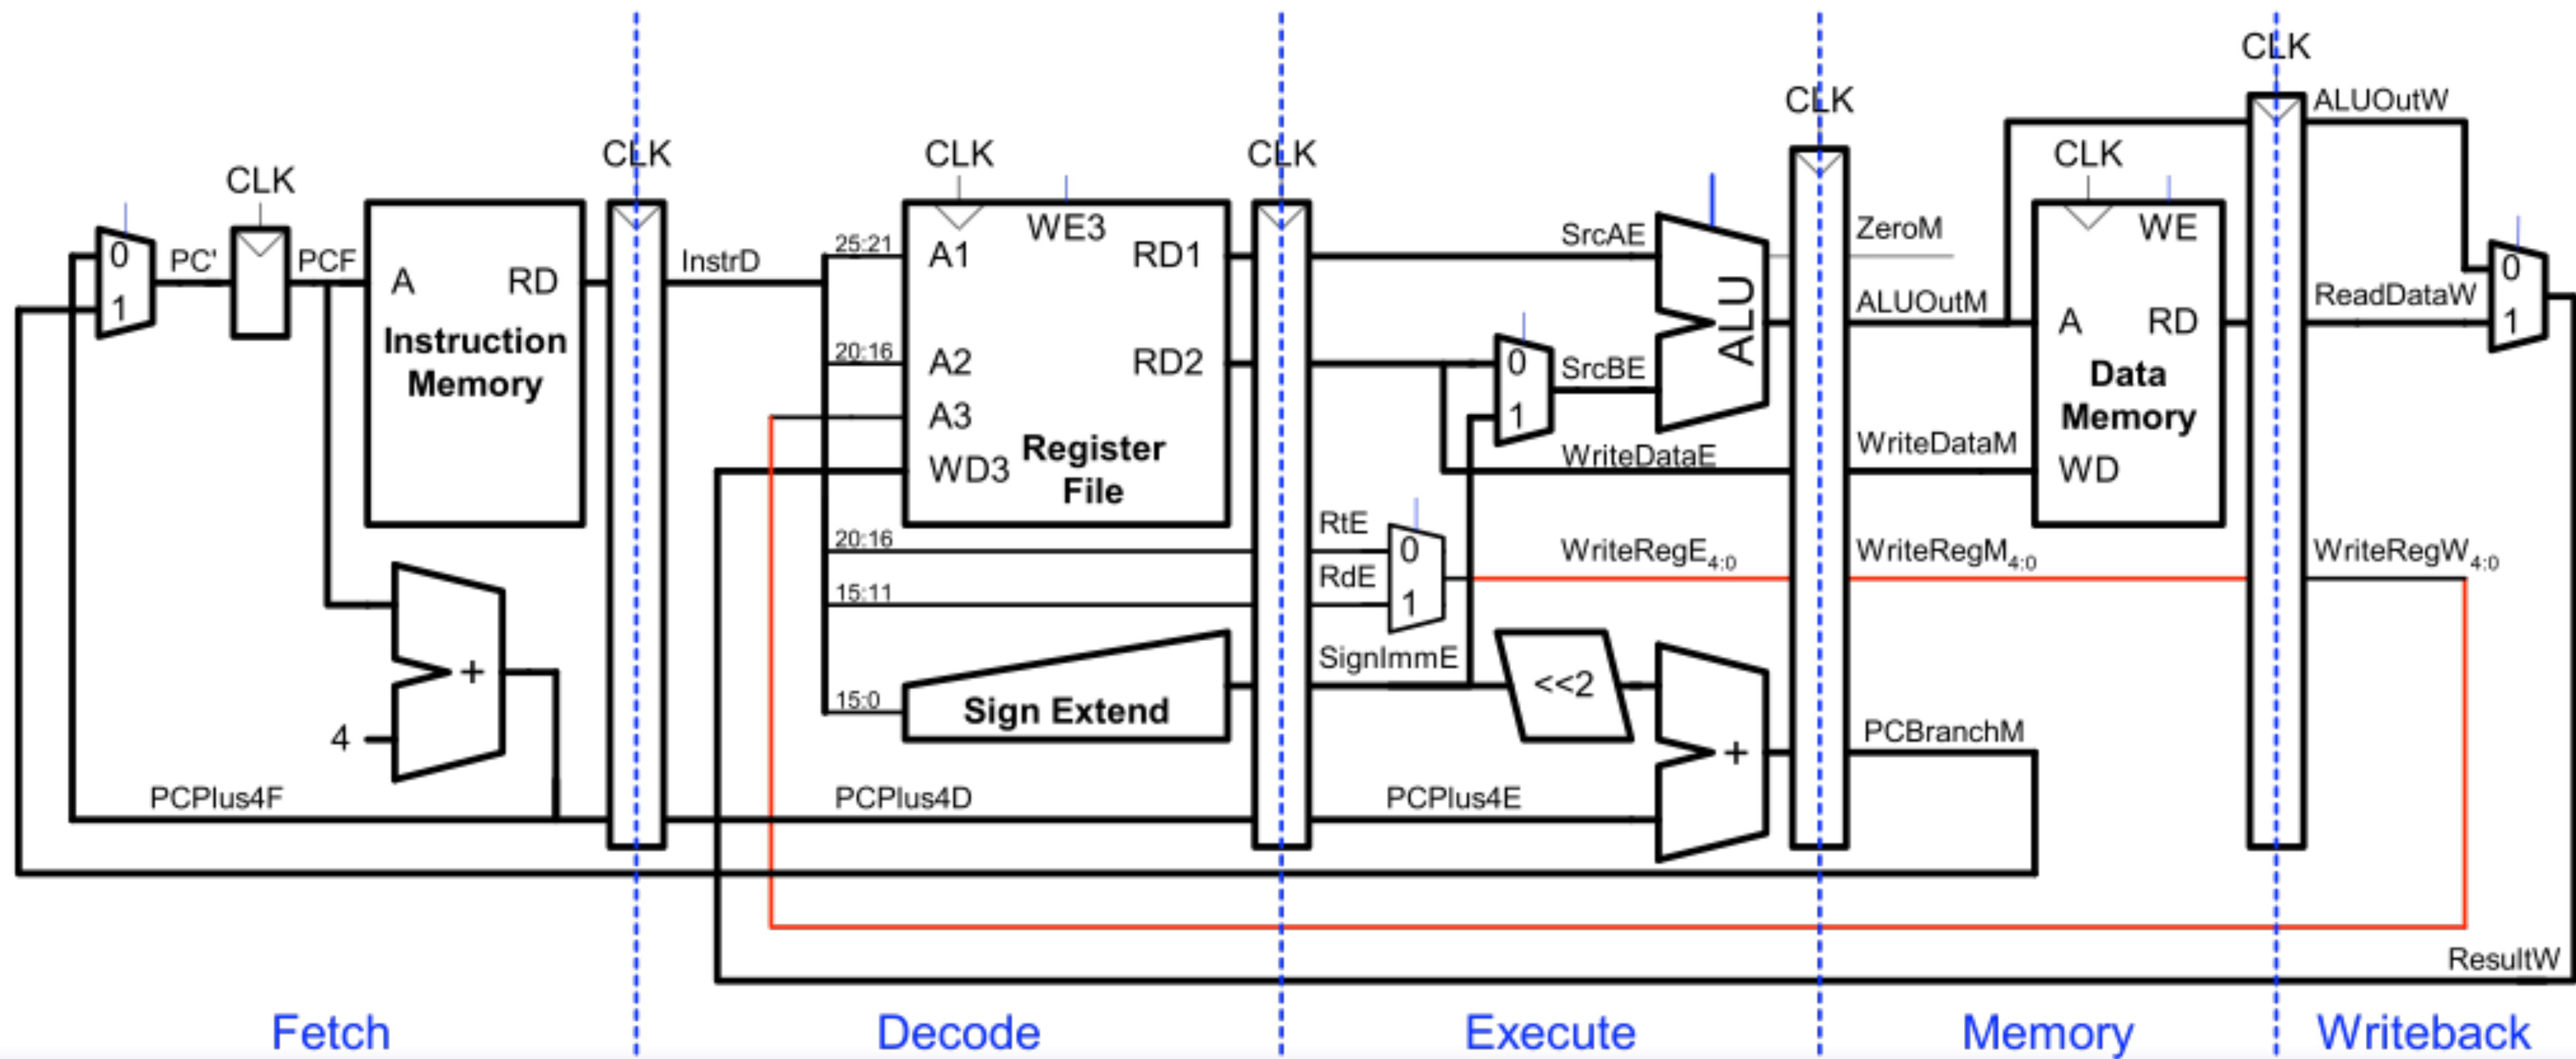


Single-Cycle and Pipelined Datapath

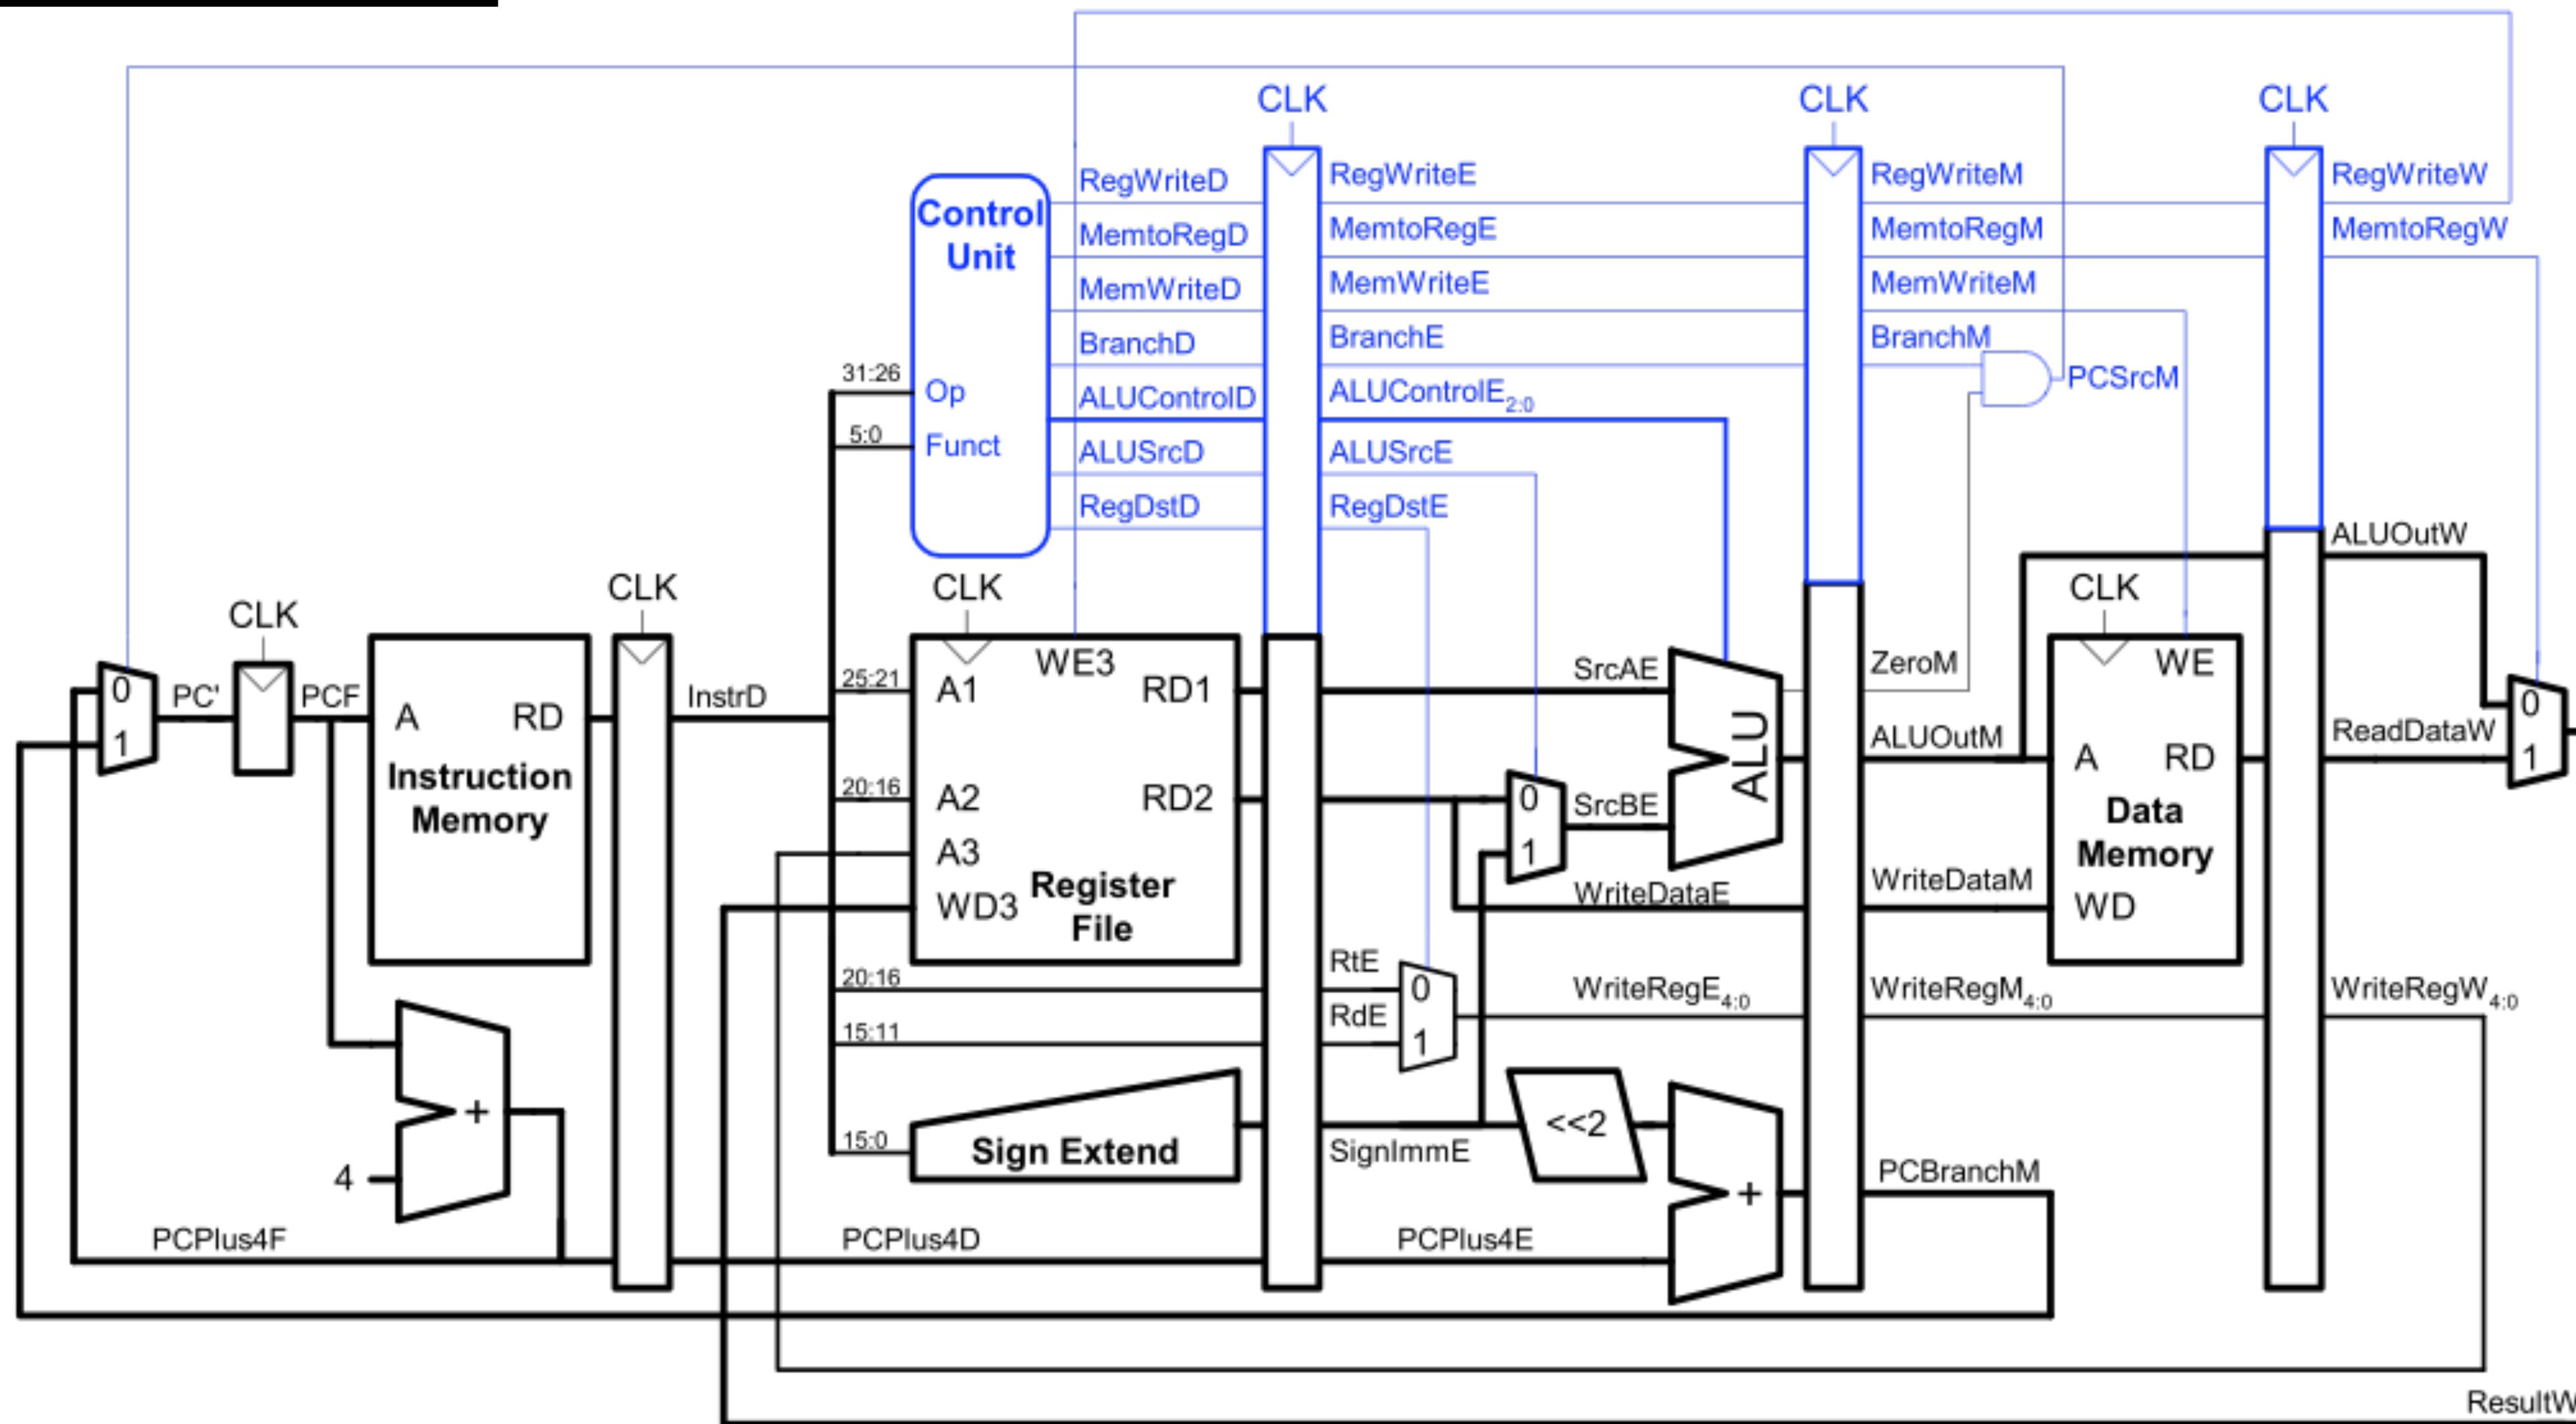


Corrected Pipelined Datapath

- *WriteReg must arrive at the same time as Result*



Pipelined Control



Same control unit as single-cycle processor
Control delayed to proper pipeline stage

Pipeline Hazards

- Occurs when an instruction depends on results from previous instruction that hasn't completed.
- Types of hazards:
 - **Data hazard:** register value not written back to register file yet
 - **Control hazard:** next instruction not decided yet (caused by branches)

We need to design ways to avoid hazards, else we pay the price in CPI (cycles per instruction) and processor performance suffers.

Processor Pipelining

Deeper pipeline example.

IF1	IF2	ID	X1	X2	M1	M2	WB
			IF1	IF2	ID	X1	X2

Deeper pipelines => less logic per stage => high clock rate.

But

Deeper pipelines => more hazards => more cost and/or higher CPI.*

Cycles per instruction might go up because of unresolvable hazards.

Remember, Performance = # instructions X Frequency_{clk} / CPI

**Many designs included pipelines as long as 7, 10 and even 20 stages (like in the [Intel Pentium 4](#)). The later "Prescott" and "Cedar Mill" Pentium 4 cores (and their [Pentium D](#) derivatives) had a 31-stage pipeline.*

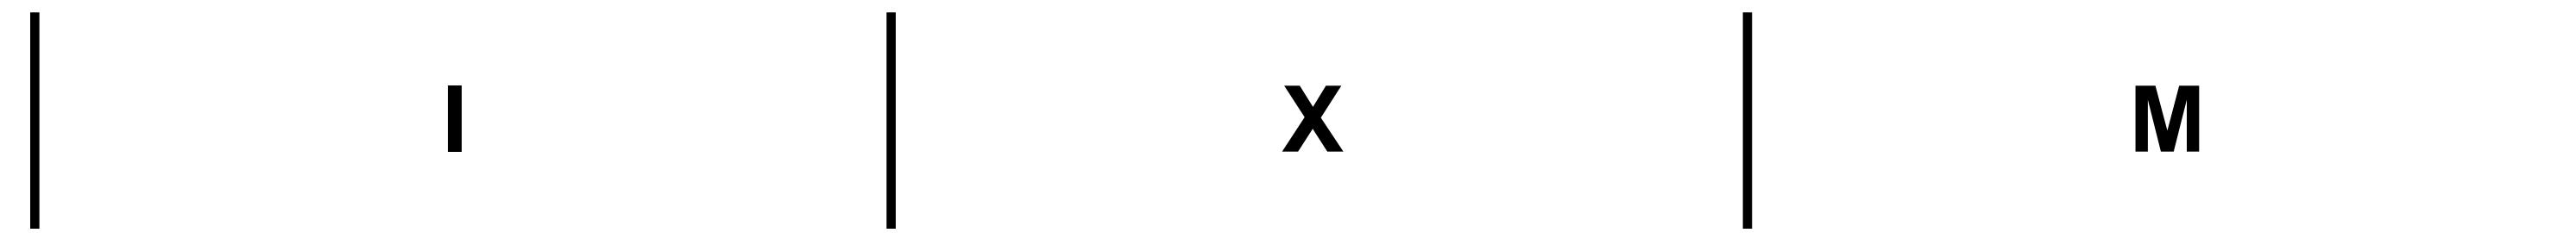
How about shorter pipelines ... Less cost, less performance (but higher cost efficiency)



3-Stage Pipeline

3-Stage Pipeline (used for FPGA/ASIC project)

The blocks in the datapath with the greatest delay are: IMEM, ALU, and DMEM. Allocate one pipeline stage to each:



Use PC register as address to IMEM and retrieve next instruction. Instruction gets stored in a pipeline register, also called “instruction register”, in this case.

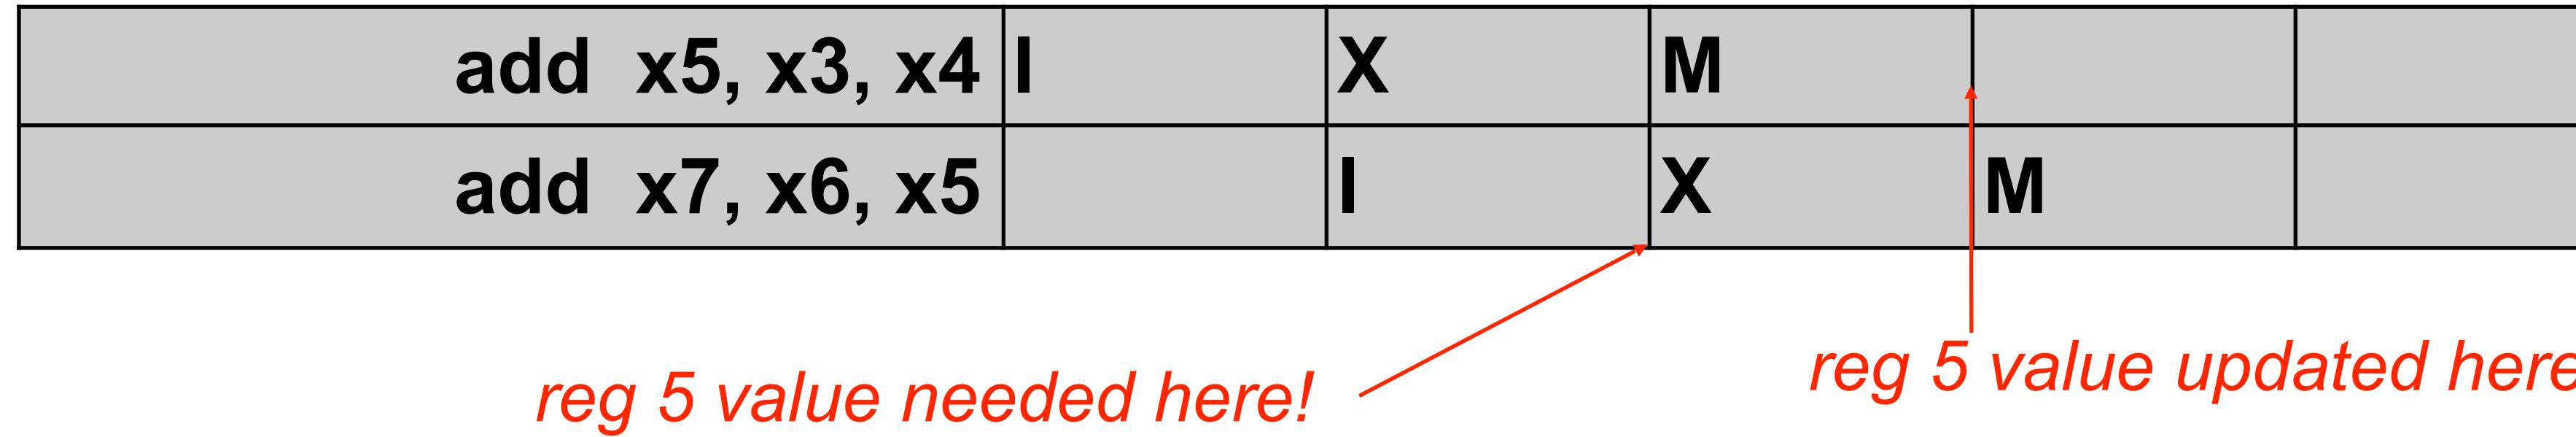
Use ALU to compute result, memory address, or branch target address.

Access data memory or I/O device for load or store. Allow for setup time for register file write.

Most details you will need to work out for yourself. Some details to follow ... In particular, let's look at hazards.

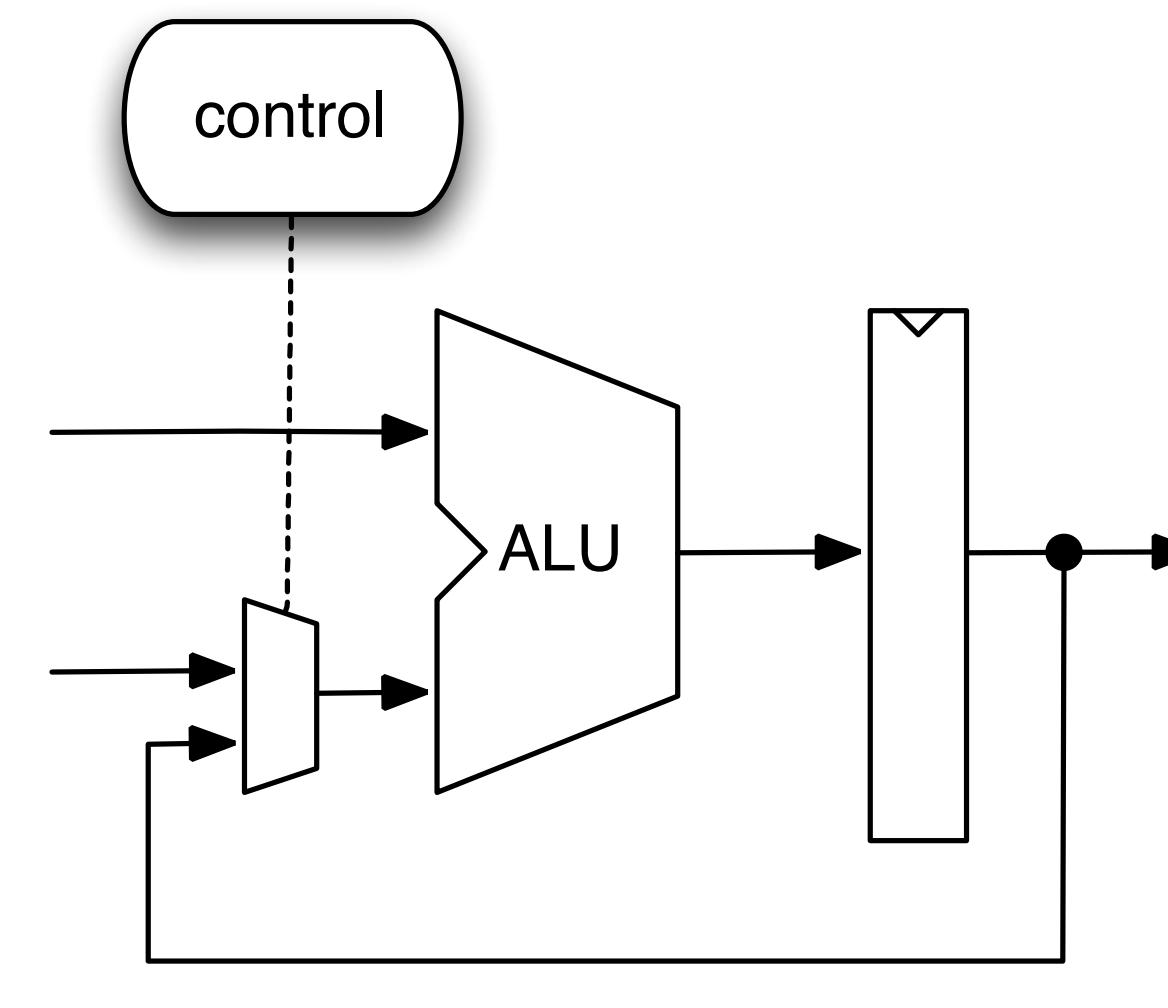
3-stage Pipeline

Data Hazard



The fix:

Selectively forward ALU result back to input of ALU.



- Need to add mux at input to ALU, add control logic to sense when to activate. Check reference for details.

3-stage Pipeline

Load Hazard

Iw x5, offset(x4)	I	X	M		
add x7, x6, x5		I	X	M	

value needed here!

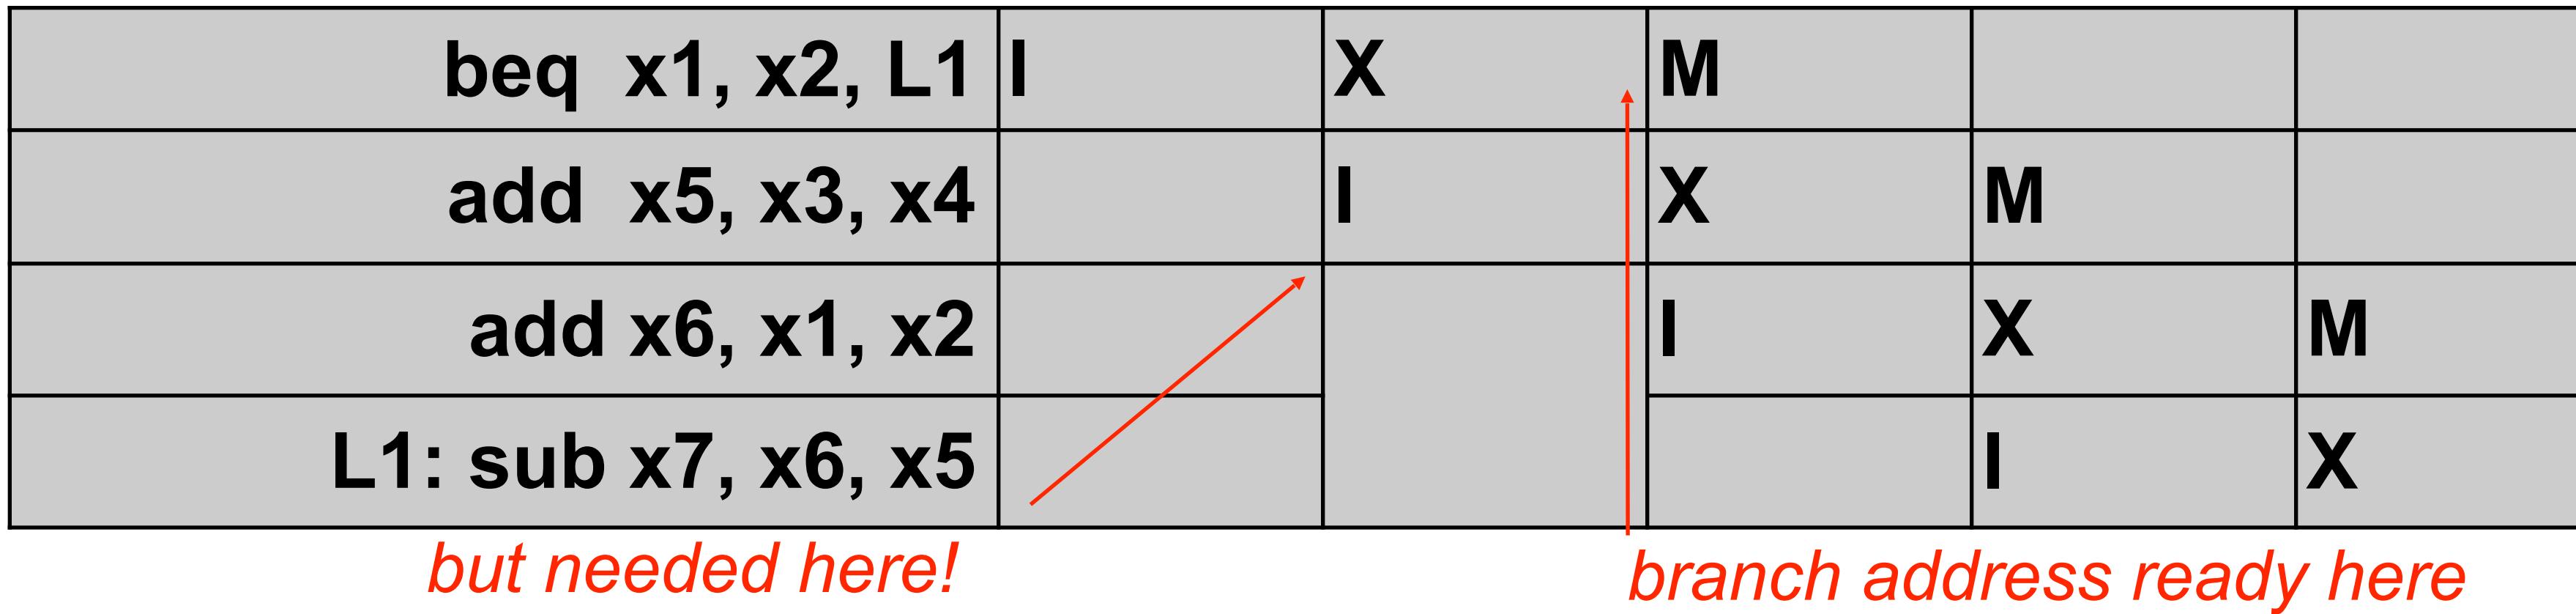
Memory value known here. It is written into the regfile on this edge.

The fix: *Delay the dependent instruction by one cycle to allow the load to complete, send the result of load directly to the ALU (and to the regfile). No delay if not dependent!*

Iw x5, offset(x4)	I	X	M		
add x7, x6, x5		I	nop	nop	
add x7, x6, x5			I	X	M

3-stage Pipeline

Control Hazard



Several Possibilities:^{*}

The fix:

1. Always delay fetch of instruction after branch
 2. Assume branch “not taken”, continue with instruction at PC+4, and correct later if wrong.
 3. Predict branch taken or not based on history (state) and correct later if wrong.
-
1. Simple, but all branches now take 2 cycles (lowers performance)
 2. Simple, only some branches take 2 cycles (better performance)
 3. Complex, very few branches take 2 cycles (best performance)

* MIPS defines “branch delay slot”, RISC-V doesn’t

Predict “not taken”

Control Hazard

Branch address ready at end of X stage:

- If branch “not taken”, do nothing.
- If branch “taken”, then kill instruction in I stage (about to enter X stage) and fetch at new target address (PC)

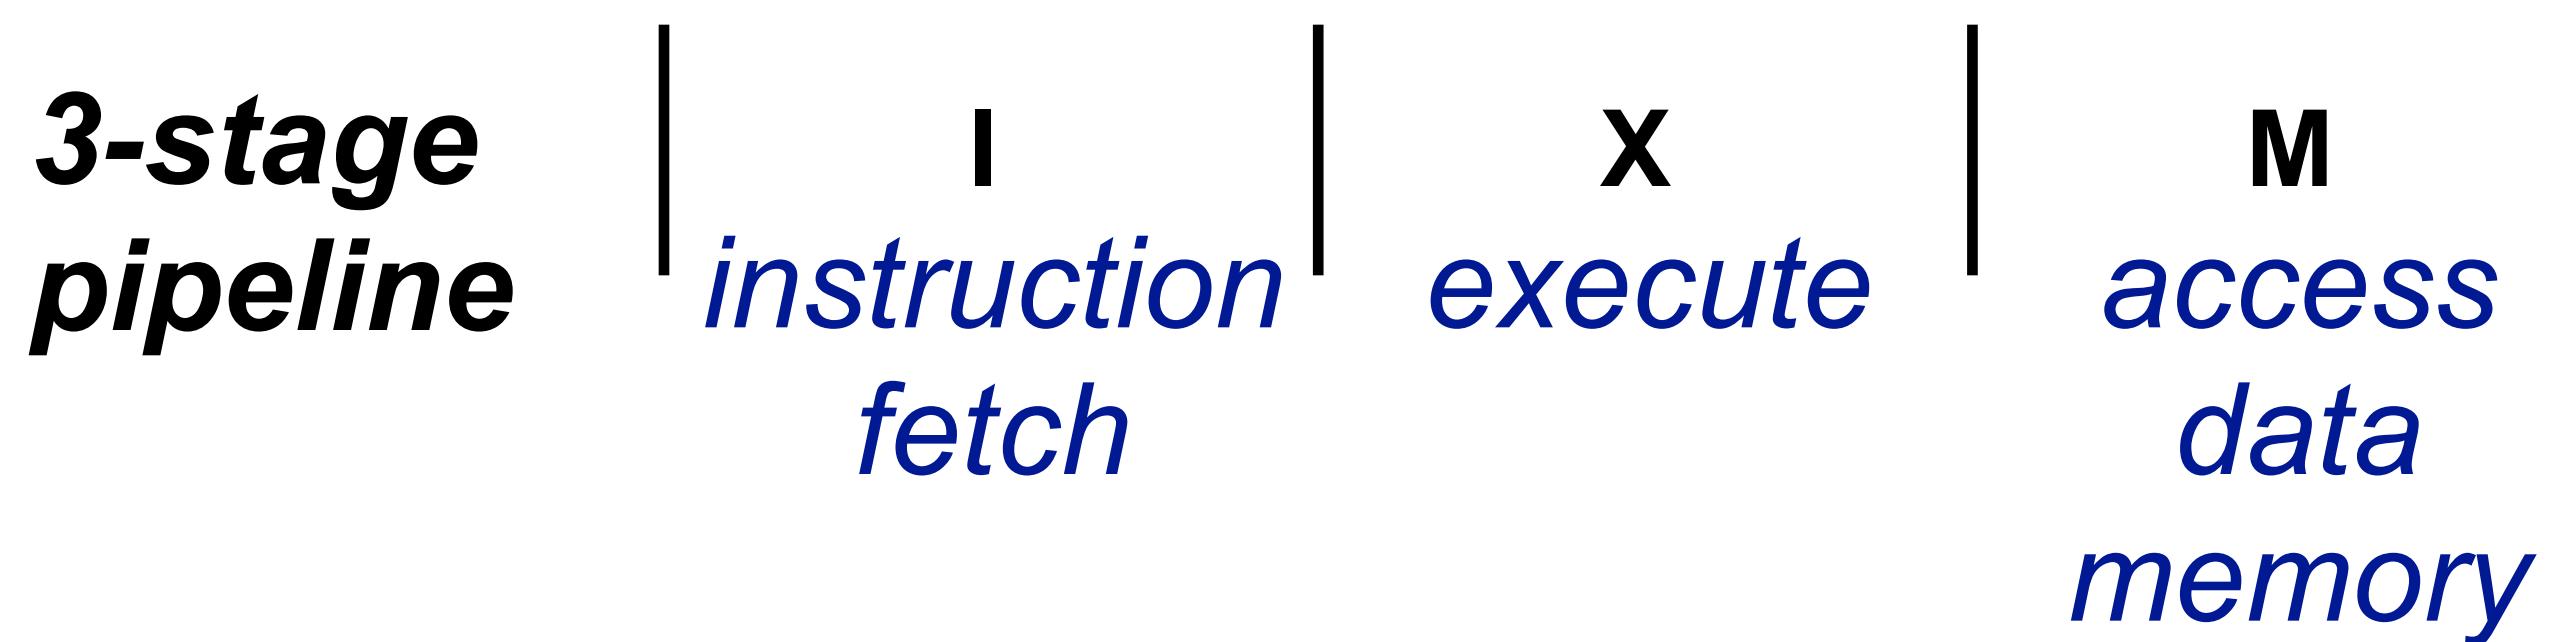
bneq x1, x1, L1	I	X	M		
add x5, x3, x4		I	X	M	
add x6, x1, x2			I	X	M
L1: sub x7, x6, x5				I	X

Not taken

beq x1, x1, L1	I	X	M		
add x5, x3, x4		I	nop	nop	
L1: sub x7, x6, x5			I	X	M

Taken

EECS151 Project CPU Pipelining Summary



- Pipeline rules:
 - Writes/reads to/from DMem are clocked on the leading edge of the clock in the “M” stage
 - Writes to RegFile at the end of the “M” stage
 - Instruction Decode and Register File access is up to you.
- Branch: predict “not-taken”
- Load: 1 cycle delay/stall on *dependent* instruction
- Bypass ALU for data hazards
- More details in upcoming spec